

Progressive Retrieval Attention: A Position Paper on Demand-Driven Context Expansion

Jorge Simao

EInnovator R&D – Center for the Sciences of Computation, Cognition, and Complexity

August 15, 2026

Abstract

Large language models commonly expose non-parametric information as a flat prompt, even when only a small and changing subset is useful. This paper argues for a complementary interface: *Progressive Retrieval Attention* (PRA), an inference-time sparse approximation to long-context self-attention. URI references and an implicit bounded prompt head define logical memory. Long sources are processed as contiguous, model-safe encoding blocks and sliced into smaller routing chunks. Compact gists index those chunks, while selected full-detail, layer-native K/V joins direct K/V under the host decoder’s attention operation. The routing representation is therefore not the attention payload, and logical, requested, resident, and materialized memory are separate quantities. This article records the original position and updates its implementation gateway to the mature native-KV design. It labels the earlier adapted cross-attention prototype as historical and summarizes the evidence subsequently established by the companion series: fragmentation-aware native transport, bounded execution, positional continuity and RoPE re-binding, distinct semantic and attention geometries, and exact pretrained-family retrofits. A usable one-shot SDK now exists for sparse direct-memory workloads, but the work does not yet establish production serving economics or a general solution to multi-hop retrieval. Bounded associative discovery is now experimentally implemented, but reliable entry routing, evidence-centered materialization, and optimized off-device serving remain explicit next steps. The proposal is falsifiable by quality per materialized K/V token, latency, transfer, resident memory, and content-causal counterfactuals against strong long-context and retrieval baselines.

1 Introduction

The transformer attention mechanism made it possible for models to condition every token on every earlier token in a sequence [37]. The same mechanism also created a central systems problem: when a model needs access to many documents, code files, logs, prior decisions, or external knowledge fragments, the simplest interface is to concatenate more text into a longer sequence. Modern long-context systems have stretched this interface dramatically, but a larger window does not by itself solve the problem of *selective use*. If only a small fraction of the context is relevant at any generation step, eager inclusion wastes computation and forces the model to perform implicit search inside an expensive dense attention substrate.

PRA proposes a different working-set rule: **recent context remains directly visible at high resolution; older or external context is retained as addressable memory and disclosed in native-KV detail only when routing predicts it is useful.** Here, *URI-addressed memory* is external source state identified by a stable URI-like reference, summarized by compact routing state, and selectively materialized as layer-native K/V. Explicit handles provide external addresses;




Figure 1: Context interfaces differ in when and where information becomes local to the model. PRA is a proposed design point, not a claim that the other interfaces should disappear.

displaced prompt history and streaming rollover are implicit addressable-memory forms that use the same machinery.

Retrieval-augmented generation (RAG) addresses part of this issue by selecting external passages before generation [21, 15, 17]. However, standard RAG often makes a single retrieval decision at the boundary of the model call. The retrieved passages are still inserted as flat prompt text, and the model must consume them whether or not each passage becomes useful later. Agentic systems take the opposite approach: they inspect tool descriptions, issue searches, open files, follow references, and progressively disclose detail only when needed [43, 30]. This behavior is effective but expensive because the loop is mediated by external actions, repeated model invocations, and hand-designed orchestration.

This paper proposes Progressive Retrieval Attention as a research program for internalizing part of that progressive-disclosure pattern. Consider a model diagnosing a failing authentication test in a large repository. Its local prompt can advertise compact handles such as `<REF_1>`, whose local reference table maps to `repo://app/auth#module`. The runtime resolves that object, encodes it in bounded contiguous blocks, slices the resulting state into routing chunks, and exposes detailed layer-specific memory only for selected chunks. The released SDK default is one-shot. Paper 2.5 separately evaluates bounded memory-to-memory discovery; online recursive decode-resume behavior remains a research target rather than an implied property of one forward pass.

The central question is therefore: *can a model obtain better quality per unit of active context, latency, and memory by treating external information as a progressively materialized context graph rather than an eagerly flattened token sequence?* The paper develops the vocabulary needed to answer that question, separates architectural, runtime, optimization, and evaluation claims that are often collapsed under “retrieval,” and states experiments that could reject the proposal.

1.1 Scope and Claim Types

The paper makes an architectural proposal, reports implementation facts, summarizes historical evidence, and advances hypotheses. The proposal is the separation among naming, encoding, routing, admission, and native attention. The current standalone implementation establishes bounded source encoding, smaller routing chunks, multi-gist indexes, exact packed routing, reusable row-isolated caches, implicit-head slicing, one global materialization budget, and CPU-backed native K/V. Companion work subsequently established exact disabled-path pretrained integration and a cross-family research SDK. These are implementation and controlled-mechanism results. They do not establish production throughput, unrestricted benchmark accuracy, recursive retrieval benefit, or the economics of a serving-optimized deployment. This distinction is maintained throughout the paper so that a working mechanism is not mistaken for a general advantage.

Figure 1 locates the proposal among neighboring interfaces. The horizontal arrangement expresses increasing separation between the local token stream and the available information; it is not a historical sequence and does not imply that PRA should replace the earlier interfaces.

Because this position paper now also serves as the entry point to a multi-paper program, Table 1

Table 1: Map of the PRA series. “Established” means demonstrated under the scope and controls of the named paper, not proved for every model or workload.

Stage	Question	What the later work established or leaves open
Paper 1	Can logical scope grow while active native K/V remains bounded?	Fragmentation repair; separate encoding and routing granularity; fixed- k and fixed-fraction scaling; bounded native execution and an off-device lifecycle.
Paper 1.5	What position does retrieved native K/V own?	Source-relative positional continuity; exact RoPE rebinding at matched positions; and the finding that native attention geometry is not necessarily a good semantic search geometry.
Paper 2	Can PRA retrofit pretrained model families without changing ordinary behavior?	Exact PRA-off paths; host-native K/V replay; depth-dependent memory use; distinct discovery, integration, decision, and behavioral measurements; a positive one-shot QASPER boundary, a HotpotQA multi-hop boundary, and a cross-family SDK.
Paper 2.5	Does frozen pretrained memory expose useful associative topology?	Bounded memory-to-memory discovery; native-edge, path, layer, and granularity diagnostics; and frozen selected-set output validation. Entry competition and broad activation remain limiting.
Next	What remains beyond whole-chunk diagnostic materialization?	Evidence-centered interval gather and deduplication, then an adaptive controller, serving-optimized search, transfer, attention, and scale validation.

separates the original proposal from results established later. The later papers refine and test the architecture; they are not evidence that every detail was present in the initial formulation.

2 Motivation

2.1 The Cost of Flat Context

In a decoder-only transformer with dense self-attention, the attention matrix for a sequence of length T contains $O(T^2)$ pairwise interactions per layer. Efficient kernels and hardware-aware implementations reduce constants, and sparse-attention variants reduce asymptotic or effective cost [2, 44, 19, 29]. Yet a conceptual mismatch remains: many tasks contain large amounts of potentially relevant context but only a small amount of immediately relevant evidence.

Consider a coding task. A repository may contain thousands of files. A developer rarely reads every file before making a change; they inspect a project map, open candidate modules, follow definitions, and read details near the call sites that matter. The useful information is embedded in a hierarchy: repository, package, file, class, method, block. The flat prompt interface collapses that hierarchy into a token sequence. Once collapsed, the model must recover structure through positional and lexical cues that were never designed as stable memory addresses.

The same pattern appears in scientific QA, technical support, legal analysis, and book-length reasoning. The system often knows that a reference exists before it knows whether the reference must be expanded. A long context window encourages eager expansion because the only available memory interface is token inclusion.

2.2 The Shallowness of One-Shot Retrieval

RAG retrieves documents or passages and places them in the model context [21]. This often improves factuality and makes model behavior auditable because generated answers can be traced to retrieved evidence. The limitation is that retrieval is usually performed at one granularity and one time. A retriever may return a passage that is too broad, too narrow, or missing the second-hop information required by the question. The generator then has limited ability to request a more specific subdocument without leaving the forward pass and entering an external loop.

Hierarchical retrieval partially addresses this issue by retrieving coarse summaries before detailed chunks, but the retrieved evidence is still usually materialized as prompt text. The generator receives the results of a pipeline rather than participating in a continuous memory-selection process. PRA asks whether part of this selection can be made differentiable, layer-local, or runtime-local.

2.3 The Expense of Agentic Search

Agentic systems use tools to search, read, execute, and revise. Their strength is not simply tool use; it is progressive disclosure. A capable agent first consumes compact descriptions and only loads detailed instructions, files, or logs when the task demands them. This style is naturally suited to large workspaces and multi-document tasks. Its weakness is cost and latency. Each search or file read may require a model turn, an external call, and a new prompt construction step.

The motivating question for PRA is therefore not whether agents should disappear. Rather, the question is whether the model/runtime boundary can absorb a lower-level version of progressive disclosure: reference-aware memory expansion during inference, with external agents reserved for high-level planning, verification, and action.

3 From Flat Context to Context Graphs

The repository example makes the required abstraction concrete. Before the model can route to the authentication implementation, the system needs a stable name for the object, an explicit relation to any child objects, and a representation compatible with the layer that will consume it. This section introduces those three pieces in that order. A useful way to keep them straight is to distinguish three levels of detail. The handle is an address, the routing gist is a small representation used to decide whether a chunk looks useful, and the token-level K/V is the evidence the decoder can actually read. Moving between these levels is the “progressive” part of PRA.

3.1 Reference Handles

A *reference handle* is a lightweight token or latent marker that names an external memory object without expanding it immediately. It plays the role of a typed pointer in the prompt: the pointer says where the object can be found, but it is not a summary of the object and does not consume the object’s full token budget. In the simplest textual form:

`<REF_1> → repo://app/auth#module`

The handle is paired with a local reference table containing at least an identifier, token, and URI, with optional summary and metadata. The local table supplies the binding: `<REF_1>` has meaning only in the example or object that advertises it. Child edges exist only through this explicit table; URI-prefix similarity is not sufficient to create graph structure. The prompt can therefore advertise a possible source without pretending that the source has already been read.

3.2 URI-Addressed Memory

PRA assumes that external memory can be addressed by URIs. The URI is the stable identity of source state, not its routing gist or transported K/V. The scheme may denote local memory, files, repository objects, search results, database rows, or runtime objects:

```
mem://manual#installation.cuda.vllm
```

The anchor fragment after `#` is not merely a byte offset. It can encode semantic hierarchy: document, section, subsection, paragraph, function, class, or trace event. URI addressing gives the runtime a stable object to cache and gives the model a compact symbol for a potentially large subtree of context.

3.3 Prompt History as an Implicit Reference

Explicit handles are not the only way an object can enter the context graph. Once a prompt exceeds the direct-token budget D , its older prefix can be reclassified as a request-local object while the most recent suffix remains in causal self-attention. For token sequence $x_{1:N}$ with $N > D$, the implemented split is exact:

$$z_{\text{head}} = x_{1:N-D}, \quad x_{\text{direct}} = x_{N-D+1:N}. \quad (1)$$

The head is registered under `pra://implicit/prompt/head` (shown in diagnostics as `#_head`), encoded once as causal history, sliced into independently routable native-K/V chunks, and routed through the ordinary PRA memory path. No visible marker is inserted into the suffix, and no decode-re-encode round trip changes the token identities. The reserved URI is scoped to one batch row and one request; a new request replaces it rather than inheriting stale prompt history.

This mechanism turns chronological displacement into the same named-object interface used for external documents. It is not equivalent to keeping the entire prefix in dense self-attention: the head becomes non-causal memory after it has been observed, and only selected chunks are transported. The implementation also retains `truncate` and `error` overflow policies. During generation, expired routing-sized prefixes of the direct suffix are appended to the same implicit object and re-encoded through the bounded source path. This correctness-first rollover rebuilds the head cache; incremental append remains future systems work.

3.4 Logical Context Under a Native Bound

PRA separates the size of its logical address space from the maximum context processed by the underlying model in any one operation:

$$L_{\text{logical}} \gg L_{\text{model,max}} \geq L_{\text{encode}} > L_{\text{route}}. \quad (2)$$

Long sources are encoded in bounded, optionally overlapping blocks. Context-only overlap is discarded before each block’s native K/V is sliced into smaller routing chunks. After routing, all candidate sources share one score-ordered whole-chunk admission budget:

$$L_{\text{direct}} + L_{\text{materialized}} + L_{\text{reserve}} \leq L_{\text{model,max}}. \quad (3)$$

This does not make the base model natively support arbitrary context. It exposes a larger logical context while keeping each encoding and attention operation within a configured native limit.

Logical context may exceed native model context, but no underlying model operation may exceed the configured native limit. This is a runtime contract, not a claim that the base model’s positional or attention machinery has been enlarged.




Figure 2: Future recursive anchor expansion. A coarse handle exposes children; an iterative policy would materialize the highlighted path to greater depth. This is not the current one-shot algorithm.

3.5 Recursive Anchors

This subsection describes the next-stage architecture, not the current one-shot algorithm. Many references are best exposed recursively. A root object may advertise child anchors, and those child objects may advertise more specific anchors:

```
mem://manual#root → mem://manual#installation → mem://manual#installation.cuda
```

Recursive anchors would turn context into a traversed graph rather than a bag of chunks. Expansion must be child-first and bounded by shared depth, child, reference, token, latency, and byte budgets. Build states must prevent partially encoded cyclic entries from becoming searchable, while content/model/tokenizer/routing fingerprints protect cache identity. The current runtime can resolve bounded child tables during cache construction, but it does not yet use generated evidence to perform online decode-resume expansion. Paper 2.5 evaluates bounded associative discovery over a frozen memory graph, but it does not turn recursive page faults into a production decoding policy.

3.6 Layer-Specific Memory Cache

If an external fragment is encoded into memory for attention, the representation should respect the transformer’s layer structure. A key/value pair computed for one layer is not generally interchangeable with another. In practical terms, PRA does not encode a chunk once and hand the same vector to every layer. It runs the chunk through the same model stack used by the prompt and captures the K/V representation appropriate to each PRA-enabled depth. PRA therefore uses a layer-specific memory cache:

$$\text{cache}[u, c, \ell] = (K_{u,c,\ell}, V_{u,c,\ell}, G_{u,c,\ell}^K, G_{u,c,\ell}^V), \quad (4)$$

where u is a URI, c a deterministic chunk, ℓ a transformer layer, (K, V) the full token memory, and $G^K, G^V \in \mathbb{R}^{G_c \times d}$ paired contextual routing gists derived from that layer’s projected chunk keys and values. Mean, last-token, reference-end, and GRU modes produce $G_c = 1$; k-means, prototype, self-organizing-map, and hybrid modes may retain several local regions. A gist is an index entry, not a compressed replacement for the chunk: it is cheap to compare with a query, while the full K/V preserves token-level detail after selection. This distinction matters because lower layers may encode lexical and syntactic information while higher layers may encode task-specific abstractions. Raw token embeddings or document summaries are not substitutes for layer-specific routing and attention memory.




Figure 3: A context graph separates compact handles, addressable objects, and the active layer-specific working set.

4 Progressive Retrieval Attention

4.1 Conceptual Computation

Return to the authentication example. Local self-attention relates the test tokens and the code already present in the prompt. The memory branch answers a different question: given the current diagnosis, which advertised module chunks should become additional evidence at this layer? The computation below separates the cheap ranking operation from the more expensive attention over detailed evidence. Operationally there are two phases. A preparation phase resolves each advertised URI, partitions its content, and builds layer-specific chunk memories. A use phase lets the live prompt query the small chunk gists, materializes only the selected K/V, and fuses the resulting memory read with ordinary causal self-attention.

Let $X \in \mathbb{R}^{T \times d}$ contain the d -dimensional hidden states for T local tokens at layer ℓ . Standard causal self-attention first projects these states into queries Q , keys K , and values V . Each head has width d_h , and M masks future positions:

$$\text{SelfAttn}(X) = \text{softmax}\left(\frac{QK^\top + M}{\sqrt{d_h}}\right)V, \quad (5)$$

The output mixes values from preceding local positions according to query–key compatibility. A PRA layer preserves that path and adds memory selected from a cache. For routing, it uses the projected query at the most recent valid position, merges its heads into $q_{t,\ell}$, and compares it with every projected-key gist $g_{u,c,\ell,j}^K \in G_{u,c,\ell}^K$ for chunk c of object u :

$$q_{t,\ell} = \text{mergeheads}(XW_Q^{(\ell)})_t, \quad (6)$$

$$a_{u,c,\ell,j} = \cos(q_{t,\ell}, g_{u,c,\ell,j}^K), \quad (7)$$

$$a_{u,c,\ell} = \text{Agg}_{j=1}^{G_c} a_{u,c,\ell,j}. \quad (8)$$

The default aggregation is maximum; mean and log-sum-exp are alternatives. URI-level gist sets can likewise select references before chunk scoring, so “reference-first” is a real two-stage

route rather than a relabeling of flattened chunk scores. The aggregate score $a_{u,c,\ell}$ ranks small routing representations rather than the full chunk. Hierarchical routing aggregates chunk scores to select at most k_r references, then independently selects at most k_c chunks inside each selected reference. Content-derived gists are the default. Summaries, when available, may replace, augment, or compete with content scores under an explicit policy, but missing summaries fall back to content. Let $\mathcal{A}_{t,\ell}$ be the selected reference–chunk pairs. This routing score answers “which drawers should be opened?”; it is not an attention weight over the evidence inside a drawer. Canonical transport prepends selected historical keys and values to the local keys and values:

$$\text{PRA}_\ell(X) = W_O^{(\ell)} \text{softmax} \left(\frac{Q^{(\ell)}[K_{\mathcal{A},\ell}; K_{\text{local}}^{(\ell)}]^\top}{\sqrt{d_h}} + M_{\text{causal}} \right) [V_{\mathcal{A},\ell}; V_{\text{local}}^{(\ell)}]. \quad (9)$$

Here $K_{\mathcal{A},\ell}$ and $V_{\mathcal{A},\ell}$ concatenate full layer- ℓ token memory from selected chunks only. Memory and local positions compete in one normalization, and the layer uses its existing W_O ; there is no memory scale or new transport projection. In the running example, the gist may select the password-validation chunk, but the decoder receives that chunk’s token-level keys and values and can assign different attention weights to its individual tokens. Routing representations are therefore cheap selectors, not compressed substitutes for transported evidence.

Routing gists are selectors; selected chunks materialize full native token-level K/V. Consequently gist count, stored source K/V, transferred K/V, and active attention K/V are different accounting quantities.

The central research problem is selection: how should $\mathcal{A}_{t,\ell}$ be chosen? Possible triggers include contextual gist similarity, attention to explicit reference-token positions, uncertainty in the next-token distribution, learned gates, or runtime policies that trade accuracy against latency. Reference ranking, chunk ranking, and materialization must remain separately measurable.

4.2 PRA as Architecture, Runtime, and Research Program

The term PRA is intentionally broader than one neural module. As an *architecture*, it denotes attention layers that consume external reference memory through native K/V injection. Optional adapted modes may use cross-attention or adapters, but they are secondary variants rather than the definition of PRA. As a *runtime*, it includes the resolvers, caches, URI stores, expansion policies, and schedules that decide when and how the memory is constructed. As a *research program*, it includes the benchmarks, ablations, scaling studies, and analyses needed to test demand-driven context expansion.

These scopes support different conclusions. A standalone prototype can validate cache construction and attention mechanics without showing readiness for a production-scale pretrained model. Runtime experiments can study scheduling and K/V reuse without settling how to train the selector. Scaling studies can ask which task distributions benefit from context graphs without committing to one resolver or cache implementation. The following commitments keep those conclusions separable.

5 Conceptual Commitments

PRA is most useful as a research frame when several concerns remain separate. Collapsing them produces systems that may work but cannot be interpreted.

5.1 Naming Is Not Materialization

A reference handle states that a memory object exists and can be addressed. It does not imply that the object’s text, embeddings, or K/V tensors have already been loaded. This distinction allows a runtime to expose a large logical context graph while maintaining a small active physical context. It also supports late binding: two deployments may resolve the same logical URI through different stores, permissions, or freshness policies.

The separation creates a useful sequence of representations:

$$\text{handle} \rightarrow \text{metadata} \rightarrow \text{optional summary} \rightarrow \text{text/chunks} \rightarrow \text{layer memory}. \quad (10)$$

Each transition has a cost and can be denied, cached, audited, or deferred. Progressive disclosure is therefore not merely truncation; it is movement through explicit levels of materialization.

For a reader following one reference end to end, the sequence is concrete. The prompt first carries a symbolic address. Resolution turns that address into an authenticated, versioned object. Encoding creates small layer-specific gists plus detailed token K/V. Routing scores the gists, and only then does memory attention consume selected K/V. A system may prepare some of these representations eagerly or retrieve them from a cache, but the logical boundaries remain the same.

5.2 Selection Is Not Transport

Selecting a URI and transporting its representation into a model are different operations. Selection may use lexical retrieval, vector similarity, explicit handle positions, learned gates, uncertainty, or an oracle. Canonical transport uses native layer K/V; prompt concatenation, cross-attention, and adapters are comparison transports. An experiment that changes both at once cannot determine whether a gain comes from finding better evidence or from presenting the same evidence through a better neural interface.

PRA evaluations should first bypass selection: restoring all historical native K/V must match full self-attention when positions and masks are equivalent. Oracle sparse selection then measures the cost of omitting K/V, and only afterward should routed, shuffled, and learned selection be evaluated. Cross-attention remains a separate adapted-transport comparison.

This sequence supports a quantitative error decomposition. The gap from full SelfAttention to native all-memory measures transport, partition, and position cost; the gap from all-memory to sparse oracle measures sparsification cost; and the gap from oracle to routed memory measures selection cost. For condition c , Recovered Context Benefit is

$$\text{RCB}_c = \frac{L_{\text{tail}} - L_c}{L_{\text{tail}} - L_{\text{full}}}. \quad (11)$$

RCB asks how much of the predictive advantage of dense context is recovered over a local tail. It is undefined when dense context offers negligible advantage and is not clipped when memory is harmful or outperforms the dense baseline. The companion systems quantity is active token-K/V fraction, $(T_{\text{local}} + T_{\text{retrieved}})/(T_{\text{local}} + T_{\text{displaced}})$; routing gists are reported separately. The scaling thesis is supported only if RCB remains high while this active fraction falls, not merely when a PRA loss is lower than an unrelated baseline.

The same separation applies inside selection. Reference ranking answers which addressable objects remain candidates; chunk ranking answers which portions of those objects are materialized. Independent budgets are necessary to prevent a prolific reference from consuming every chunk slot. Reported cache size, selected gist count, and transported token count are different quantities and should never be used interchangeably.

Active K/V is not total resident K/V. Sparse attention becomes a GPU-capacity benefit only when unselected payload can remain off device, be evicted, or be shared without materialization.

5.3 Implementation Invariants

The mature implementation is easier to reproduce when its invariants are stated before its optimizations. First, when PRA is disabled or has no selected memory, the host attention path is bypassed rather than approximately reconstructed. The standalone tests require the native history replay to match full causal attention when positions and masks are equivalent; pretrained-family adapters additionally require exact PRA-off logits, hidden states, generation, and tested cache behavior. Exactness applies to the named controls, not to sparse memory as an approximation of an omitted dense context.

Second, the compact index and attention payload have different jobs. A routing gist may be a pooled hidden state, projected key, or learned semantic vector. It identifies a URI and routing chunk. Once selected, the payload is the host layer’s full-detail native K/V, including its head layout, dtype, positional treatment, and cache convention. Canonical PRA introduces no separate memory output projection or residual scale; the historical cross-attention mode does.

Third, accounting follows the lifecycle rather than one ambiguous “context fraction.” Let T_{logical} be all addressable source tokens, $T_{\text{requested}}$ the tokens belonging to selected references or parents, and $T_{\text{materialized}}$ the K/V tokens admitted to an active operation. Then

$$0 \leq T_{\text{materialized}} \leq T_{\text{requested}} \leq T_{\text{logical}}, \quad (12)$$

but equality is not assumed. Routing gists, host-resident K/V bytes, transferred bytes, padded attention allocation, and attended physical K/V tokens are reported separately. This prevents a logical selection rate from being presented as a physical memory saving.

Finally, the lifecycle remains bounded. Every source encode, direct-plus-memory attention operation, resolver expansion, transfer, and cache admission is subject to explicit token, byte, depth, identity, and isolation constraints. CPU or off-device residency reduces accelerator capacity only when selective transfer keeps unselected payload off the active device; it does not erase encoding, storage, or communication cost.

5.4 Logical Isolation Must Survive Physical Batching

Each sequence has a logical memory address space even when kernels share physical tensors. Two rows may even use the same URI string for different content; the row-local reference table decides what that name means. An efficient implementation can still run the prompt as one batch by keeping a private cache namespace for each row, routing row i only against namespace i , and masking padded memory positions. Correct batching may evaluate each row independently, pad one masked rectangle, or group rows into deterministic length buckets. These execution plans are equivalent only if empty rows receive an exact zero memory residual, outputs return to original order, and one row’s cache cannot influence another row’s logits. Padding allocation and bucket count are systems metrics, not semantic changes to routing.

5.5 Language Ability Is Not Reference Adaptation

Adding a memory path can improve a reference-formatted task while damaging ordinary language modeling. Conversely, freezing the local model can guarantee preservation while limiting adaptation. The relevant object is therefore a two-dimensional outcome: reference-task improvement and

local-task retention. A useful method should report both, including the exact trainable parameter groups.

This distinction motivates staged optimization. First establish that each architecture learns ordinary language. Then initialize from the matching checkpoint and train only newly activated reference-specific modules. Finally, compare joint fine-tuning under a mixed plain/reference curriculum. Preservation obtained by freezing is a property of the optimization policy, not proof that an architecture is intrinsically immune to forgetting.

5.6 Memory Identity Requires Isolation

URI identity is part of the conditioning state. If a batched implementation combines references from unrelated examples without a corresponding mask, it changes the task and can leak evidence across samples. The standalone implementation exposed this issue directly: a batch-union cache made within-batch shuffling ineffective. Correctness required a cache per example. Efficient PRA runtimes must preserve the same logical isolation through cache indices or attention masks even when physical storage and kernels are shared.

6 Virtual Memory for Transformers

Virtual memory changed computing by separating the address space visible to a program from the finite physical memory resident at a given moment. Pages are named independently of their current location, loaded lazily, cached while useful, and evicted under pressure. Working-set theory connects recent access behavior to resource allocation rather than pretending that every address must remain resident [9]. Foundation-model context has an analogous tension: the logical information relevant to a task can be far larger than the token and K/V memory physically active during one decoding step.

PRA proposes a *virtual context space*. URIs and anchors play the role of stable logical addresses. Resolved text or structured objects are backing-store content. Encoded layer-specific K/V entries are resident pages. Selection or gating acts as an admission policy. A missing selected representation creates a context fault that triggers resolution and encoding. The active reference cache approximates a working set, while eviction and reuse determine whether the abstraction saves end-to-end cost.

6.1 PRA as Semantic Virtual Memory

The target is not literal infinite context. It is a systems regime in which logical context is much larger than physical active context:

Logical addressable context can grow far beyond the model’s native attention window and GPU-resident K/V capacity without requiring all logical context to be simultaneously materialized in attention.

We call this *effectively unbounded logical context*, *externally backed logical context*, or *demand-loaded context*. All three phrases remain conditional on finite backing storage, an executable lookup policy, bounded native operations, and measured systems cost. They do not imply mathematically infinite memory, constant lookup time, or quality that is independent of corpus scale.

The target GPU footprint is approximately

$$M_{\text{GPU}} \approx M_{\text{model}} + M_{\text{head}} + \sum_{\ell \in \mathcal{L}_{\text{PRA}}} M_{\text{active}}^{(\ell)} + M_{\text{runtime}}, \quad (13)$$

where $\mathcal{L}_{\text{PRA}}$ is the set of consumer layers. The intended signature is that M_{active} depends primarily on the task/query working set, not on every token in the logical store. Routing and index work may still grow with the store, CPU/reference memory grows with retained content, and a poor policy may broaden the working set. Scale-independent physical memory is therefore an experimental target, not a present claim.

Later results refine the early “chunk equals page” analogy into three granularities:

$$G_{\text{encode}} \neq G_{\text{search}} \neq G_{\text{materialize}}. \quad (14)$$

G_{encode} is the context used to create meaningful representations; G_{search} is the conceptual routing or native-graph node; and $G_{\text{materialize}}$ is the actual native-K/V interval loaded into active attention. Coarser encoding can contextualize finer search nodes, while physical materialization may later union or window several selected nodes. Paper 2.5 shows why the separation matters: finer search nodes can reduce conceptual payload while also removing useful associative edges.

Sparsity also has a layer axis. PRA may be sparse over

$$\text{memory regions} \times \text{token/K/V detail} \times \text{consumer layers}. \quad (15)$$

Paper 2 establishes selective pretrained-model consumer paths. Paper 2.5 Gate-3 output diagnostics test sparse and interleaved layer bands against all-layer injection, but those model- and dataset-specific findings remain preliminary placement evidence rather than a frozen deployment policy.

6.2 Living Architecture Status

Table 2 is a living contract between this position paper and the later PRA series. **Demonstrated** means implemented and measured in the stated scope; **Strong** means supported by multiple mechanisms or artifacts without a decisive end-to-end systems result; **Partial** means a working component has an important missing subproblem; **Emerging** marks an initial but not yet stable design; **Planned** marks specified next work; and **Open** marks an unresolved required property. These labels are not interchangeable with production readiness. The structured source `virtual_memory_status.json` records artifact paths, headline results, remaining gaps, and the update date for every row.

6.3 Paper-Series Roadmap

The table assigns a specific role to each manuscript rather than treating later papers as a single undifferentiated validation:

- **Paper 0, *Progressive Retrieval Attention: A Position Paper on Demand-Driven Context Expansion***: architectural promise and living virtual-memory contract.
- **Paper 1, *Progressive Retrieval Attention: Bounded Sparse Native-KV for Addressable Logical Memory***: bounded native-K/V transport, logical/native separation, and off-device lifecycle [33].
- **Paper 1.5, *Position, Attention Geometry, and Semantic Routing for Retrieved Native-KV Memory***: positional and representation correctness, including RoPE geometry [32].
- **Paper 2, *Progressive Retrieval Attention for Pretrained Transformers: Sparse Native-K/V Memory with Semantic Routing***: Hugging Face integration, externally backed pretrained-model memory, and one-shot SDK boundaries [34].

Table 2: Semantic virtual-memory status, part I. Evidence is bounded to the named paper scope.

Virtual-memory concept	PRA analogue	Status	Evidence / next validation
Virtual address space	URI references + logical positions	Strong	Papers 0–2; arbitrary-scale namespace validation remains open.
Backing store	CPU/external source + reference K/V	Strong	Papers 1–2; distributed freshness and cache economics remain open.
Page/index metadata	Gists, hierarchy, spans, reference metadata	Strong	Papers 1–2.5; large-corpus index maintenance is unmeasured.
Address/page lookup	Query facets + semantic root routing	Partial	Paper 2.5 oracle root ceiling is high; executable competition is weaker.
Locality between pages	Native Q/K successor graph	Strong	Paper 2.5 finds natural edges, with strong layer/granularity dependence.
Page-in	Selected CPU→GPU native K/V	Demonstrated	Papers 1–2; fused asynchronous serving transfer remains open.
Bounded physical working set	Bounded head + active reference K/V	Demonstrated	Paper 1 controlled scaling; multi-million-token task validation remains open.
Page granularity	Separate encode/search/materialize units	Partial	Paper 2.5 diagnoses the tradeoff; Paper 3 must choose physical windows.
Cross-page gather	Logical interval gather across chunks	Planned	Paper 3 target; no evidence-centered gather implementation yet.
Deduplication	Union overlapping logical K/V intervals	Planned	Paper 3 target; graph identity deduplication is not physical interval union.

Table 3: Semantic virtual-memory status, part II (continued).

Virtual-memory concept	PRA analogue	Status	Evidence / next validation
Demand paging policy	Route → discover → disclose K/V	Partial	Paper 2.5 executes the stages; entry and memory-use interference remain.
Replacement / eviction	Bounded head + adaptive working-set policy	Partial	Rollover is implemented; learned eviction and pressure tests are absent.
Page-fault handling	Discover missing memory, materialize, continue	Emerging	Stages execute separately; no calibrated online decode-resume loop.
Multi-level cache	External/CPU K/V → GPU active K/V	Demonstrated	Basic Papers 1–2 lifecycle; no distributed coherence or hit-rate study.
Locality-aware prefetch	Associative graph expansion	Demonstrated	Paper 2.5 mechanism and natural diagnostics; benefit depends on entry and topology.
Thrashing control	Adaptive routing, K/B, layers, materialization	Open	Broad activation is measured; no stable adaptive controller.
Physical page size	Native-K/V materialization window	Open	Whole selected chunks are diagnostic; Paper 3 targets evidence-centered windows.
Placement across depth	Selected PRA consumer layers	Emerging	Paper 2 plus preliminary Paper 2.5 Gate-3 diagnostics; no general policy.
Translation / controller	Facet controller + trained router	Partial	Direct routing works in regimes; Paper 2.5 retains a large oracle-facet gap.
Scale-independent memory	Bounded active K/V as logical memory grows	Open	Small fixed- k curves only; decisive 32K–8M+ scaling is required.




Figure 4: Semantic virtual-memory target. Logical memory and its backing store remain outside the GPU attention boundary. Root routing and associative search choose a conceptual working set; only the final native-K/V intervals and consumer layers determine active attention state.

- **Paper 2.5, *Iterative Progressive Retrieval Attention: Bounded Associative Discovery over Latent Gist Memory*:** associative discovery, graph topology, granularity/layer effects, and diagnostic selected-set output validation [31].
- **Paper 3 (planned):** evidence-centered native-K/V materialization, interval gather, overlap union, and physical page-size selection.
- **Paper 3.5 (planned):** trained facet/controller policy, adaptive topology/search, GPU kernels, serving metrics, and the decisive logical-memory scaling experiment.
- **Paper 4 (conditional):** PRA-aware or from-scratch training if frozen geometry and runtime control prove insufficient.

6.4 Address Spaces, Protection, and Lifetime

A useful virtual context address is more than a string. It should identify scheme, authority, object, semantic anchor, version, tenant, and access policy. Two identical texts with different provenance may not be interchangeable; one URI whose content changes over time should not silently reuse stale K/V. Cache keys may therefore require a tuple such as

$$(u, v, \ell, m, \theta, p), \quad (16)$$

where u is URI, v content version, ℓ layer, m model identity, θ relevant encoder parameters, and p policy domain. Cache lifetime can range from one sequence to a user session, repository revision, model deployment, or immutable corpus version. Longer reuse improves amortization but increases freshness and revocation risk.

Permissions must be enforced before resolution and again when resident memory is shared. A model should not infer a protected reference because another tenant populated the same physical cache. Logical address spaces and attention masks must survive batching, deduplication, paging, and distributed placement.

6.5 Page Faults and Progressive Disclosure

A conventional page fault is generated by an exact address access. A model’s “context fault” is less crisp: it may be inferred from explicit handle attention, uncertainty, a learned router, or a runtime policy. False negatives omit evidence; false positives waste resolution and encoding. This turns fault prediction into a joint learning-and-systems problem. Recursive anchors further resemble multi-level page tables only loosely: following a child changes semantic granularity rather than merely translating an address.

6.6 Where the Analogy Breaks

CPU pages are intended to preserve exact program semantics regardless of placement. Neural memory is approximate, content-dependent, and transformed by model parameters. Loading a document into K/V does not guarantee that the model will use it correctly. Attention can blend untrusted entries, positional semantics may matter, and encoding may be more expensive than transport. The virtual-memory analogy is therefore valuable for identifying address, isolation, admission, working-set, and eviction problems; it is not evidence that operating-system policies can be copied unchanged.

Classical translation starts from an exact virtual address and follows deterministic page tables. PRA translation is learned and content-addressed:

$$\text{question} \rightarrow \text{semantic candidates} \rightarrow \text{associative expansion} \rightarrow \text{materialized native K/V}. \quad (17)$$

It therefore has semantic failure modes with no direct analogue in ordinary virtual memory: a wrong root, distractor competition, a missing native edge, over-broad activation, or a miscalibrated controller. Conversely, PRA can propose content-related neighbors and perform associative prefetch rather than waiting for an exact address. That extra capability is useful only when its recall, false activation, and physical cost are measured together.

6.7 Decisive Virtual-Memory Scaling Experiment

The unresolved systems claim requires a direct scaling experiment, not extrapolation from small fixed- k curves. Construct matched logical memories of 32K, 128K, 512K, 2M, and at least 8M tokens while holding task families and evidence requirements constant. At every scale report: answer quality; root and evidence recall; active native-K/V tokens and bytes; peak HBM; CPU/reference-store memory; routing, index, and graph-search latency; interval-gather and H2D latency; TTFT; TPOT; and, when the runtime supports it, throughput and concurrency. Dense native context, sparse attention, RAG, and agentic search should receive matched task and hardware conditions.

The desired signature is stable matched-task output quality while logical memory grows from 32K to 8M+ tokens and the GPU native-K/V working set remains approximately bounded by a task-dependent budget. Search or index latency may grow, and backing-store memory should grow; the narrower hypothesis is that GPU attention working-set cost does not grow linearly with total logical context. Failure to preserve recall, quality, or bounded active K/V at scale would narrow or reject the virtual-context thesis. Until this experiment exists, PRA does not provide scale-independent physical memory, constant-cost page faults, or reliable routing at arbitrary memory size.

7 Current Architecture and Reference-Memory Lifecycle

A PRA runtime is a lifecycle rather than a single retrieval call. The current one-shot path is:

1. **Advertise:** expose handles, summaries, types, provenance, and access policy.
2. **Select references:** rank candidate handles for a sequence, position, layer, or decoding interval.
3. **Resolve:** map selected URIs to immutable or versioned content.
4. **Encode:** process source text as bounded contiguous blocks and retain host-native, layer-specific K/V; context-only overlap is not a separately addressable payload.
5. **Slice and index:** divide encoded K/V into smaller routing chunks and derive one or more compact gists without replacing the detailed payload.
6. **Route:** rank reference and chunk identities against the compact index.
7. **Admit and transfer:** choose whole chunks under one native-operation budget and move only selected K/V from CPU or other backing storage when needed.
8. **Attend:** allow authorized sequence positions and layers to consume the memory.
9. **Reuse or evict:** retain useful entries across tokens or requests, or release them under pressure.
10. **Audit:** record which version, URI, layer, and policy influenced the output.

This lifecycle suggests a cost objective. For task quality Q , prompt tokens T_p , encoded reference tokens T_r , latency L , and active cache bytes B , a system should be compared on a frontier such as

$$\max Q \quad \text{subject to} \quad c_p T_p + c_r T_r + c_l L + c_b B \leq C. \quad (18)$$

The constants depend on hardware, cache reuse, and service-level objectives. PRA is not successful merely because it shortens the visible prompt. Reference encoding that is repeated for every request may cost more than eager context. The economic case depends on selective admission, amortized reuse, or better quality under a fixed resource budget.




Figure 5: Canonical PRA architecture. Solid arrows are the one-shot whole-chunk SDK path. The compact index selects identities; attention consumes selected full-detail, layer-native K/V. The dashed discovery loop is exercised by Paper 2.5 experiments but is not the SDK default; partial materialization remains planned.

8 Architectural Placement

PRA need not occupy every transformer layer. Lower layers often build local lexical and syntactic representations, while later layers may provide more useful queries for semantic memory. A hybrid architecture can therefore place ordinary SelfAttention below PRA-enabled layers. The opposite hypothesis is also plausible: early access to external tokens may be necessary for composition across all depths. Full PRA, upper-layer PRA, mixed blocks, and sparse PRA placement should be treated as competing designs.

The standalone study compares four local layers, four PRA layers, and two local plus two PRA layers. In the five-seed corrected follow-up, frozen-path reference loss is 4.6720 for full PRA and 4.7296 for hybrid PRA. Scratch-to-frozen improvement is also larger for full PRA (1.6837 versus 1.6346), so the tested hybrid does not benefit more from staging. Hybrid PRA does use half as many trainable memory output projections and trains faster in this regime, but one tiny scale and an unlearned selector cannot resolve placement. A proper comparison still needs matched compute budgets, layer-wise reference attention diagnostics, and tasks stratified by whether required evidence is lexical, syntactic, or semantic.

9 Relationship to Existing Work

PRA draws mechanisms from several mature lines of work. The comparison is organized by *which resource changes*: prompt text, token connectivity, positional range, retained K/V, persistent state, retrieved neural state, or the external action boundary. This framing narrows the novelty claim to an access contract that combines stable names, a compact semantic index, selective native-K/V materialization, and a bounded lifecycle. None of those ingredients is claimed as individually new.

9.1 Retrieved Text and Latent Retrieval

RAG, REALM, and dense passage retrieval select external text, usually before generation, and let the generator encode the selected passages in its current prompt [21, 15, 17]. Their changed resource is the input text. This preserves a simple model interface and lets retrieval scale independently, but selected text still incurs prompt encoding and dense participation once inserted. PRA can use the same retriever to choose candidate URIs, then route smaller chunks and consume cached layer-native state rather than append every selected passage.

Latent retrieval moves selection closer to model representations. The k NN-LM interpolates the model distribution with neighbors from a hidden-state datastore; Memorizing Transformers retrieve prior token memories; and RETRO retrieves text chunks for learned generator-side conditioning [18, 38, 3]. These methods and PRA all expose non-parametric information during inference. Their memory values and fusion rules differ: PRA retains URI identity and provenance, uses compact vectors only for search, and presents selected full-detail native K/V to ordinary attention. It does not interpolate a neighbor token distribution or require RETRO’s trained cross-attention architecture.

9.2 Sparse, Local, and Block Attention

Longformer, BigBird, Reformer, and Routing Transformer change token-to-token connectivity inside an admitted sequence through local windows, global tokens, hashing, or content routing [2, 44, 19, 29]. These methods answer how a long sequence should interact after it is present. PRA first asks which external state should enter the active operation. The approaches are complementary: a sparse or block kernel could execute direct-plus-materialized K/V after PRA admission.

Landmark Attention similarly gives a long sequence coarse random-access landmarks [26]. A PRA gist has a related indexing role, but it belongs to a separate addressable cache and is not the evidence consumed by attention. The selected routing chunk materializes its detailed K/V payload. This routing-payload distinction is the central difference.

9.3 Long-Context Position and Sequence Range

Context-extension methods change the length over which one model invocation can preserve useful token interaction. Transformer-XL adds segment recurrence and relative positions; Position Interpolation and LongRoPE extend the operating range of positional geometry [6, 4, 10]. Their strength is continuous access to a large admitted sequence. PRA does not enlarge the host model’s native limit. It constructs logical scope from model-safe encoding blocks and admits a bounded subset to each operation. A longer native window can therefore improve PRA’s encoding and direct context while leaving the logical-versus-active separation useful.

9.4 KV Retention, Compression, Eviction, and Quantization

H₂O and SnapKV evict or retain decoding-cache tokens, ClusterKV merges cache state in a semantic space, and KIVI reduces bytes through asymmetric quantization [45, 22, 23, 24]. Their changed resource is an already-created KV cache. PRA instead chooses provenance-bearing reference and chunk identities before their payload enters active attention. These techniques are compatible: eviction, merging, quantization, or future token-level materialization can operate inside PRA’s selected working set. Their errors must then be measured separately from routing errors.

9.5 Recurrent, Compressive, and Persistent State

Compressive Transformers, Infini-attention, TransformerFAM, Titans, and selective state-space models carry forward compressed, recurrent, feedback, learned, or fixed-width state [28, 27, 16, 1, 14]. StreamingLLM preserves a chronological attention working set through sinks [39]. These methods change the state propagated through time and can be more economical than retaining addressable token detail. PRA preserves separately named source objects and can revisit selected full-detail K/V out of chronological order. Its implicit prompt head is the closest overlap, but beyond the native horizon its bounded-block encoding is an explicit approximation rather than a sufficient recurrent state.

9.6 Prompt and Context Caching

PagedAttention manages serving-time KV allocation, while Prompt Cache reuses schema-named prompt modules and CacheBlend selectively recomputes cached chunks to restore missing cross-chunk context [20, 11, 42]. These systems change where prompt state resides and how much prefill is repeated. PRA shares their concern with positional validity, reuse, transfer, and backing storage, but adds semantic routing over addressable objects and a hard active-K/V admission budget. Conversely, PRA does not yet provide their serving kernels or throughput evidence; an optimized runtime should build on these systems rather than treat CPU residency alone as a speed result.

9.7 Memory-Augmented Networks

Neural Turing Machines and end-to-end Memory Networks couple a controller to differentiable external memory and demonstrate content-addressed reads, writes, and multiple hops [13, 35]. Hierarchical Attention Networks and adaptive computation likewise motivate multi-level representation and variable computation [41, 8, 12]. PRA inherits the general principle that access can be learned and iterative. It differs in preserving host-transformer-native K/V, stable URI identity, cache provenance, and an exact disabled path. Paper 2.5 tests parameter-free bounded associative discovery, but it is not equivalent to the trained multi-hop controllers in those systems and does not establish reliable online fault handling.

9.8 Agentic and Protocol-Level Access

ReAct and Toolformer coordinate reasoning with external actions [43, 30]. They change the action boundary and can search, verify, or modify the world across model calls. PRA targets a narrower read path for stable or versioned resources. It does not replace explicit tools, side effects, planning, or verification.

9.9 MCP and IDE Assistants

The Model Context Protocol (MCP) defines a client-server interface through which AI applications discover resources, prompts, and tools [25]. IDE assistants similarly expose repository search, symbol navigation, diagnostics, terminals, and version-control state. These systems make external capability discoverable, but the host application still decides how retrieved resource content enters a model context. PRA is complementary: MCP can supply the address space and resolver boundary, while PRA can define how selected resources become internal attention memory.

This combination suggests a layered interface. MCP or an IDE host handles authentication, capability discovery, transport, and side effects. A PRA runtime handles content identity, versioned




Figure 6: RAG selects before a model call, agents and MCP iterate across an application boundary, and PRA proposes addressable memory consumption within a runtime/model boundary.

encoding, cache admission, and model-local use. High-level tool calls should remain explicit and auditable; repeated read-only access to stable resources is the more plausible target for latent memory.

10 Embodied, Situated, and Referential Memory

Embodied and enactive approaches argue that cognition is shaped by recurrent interaction among brain, body, and environment rather than by isolated symbol manipulation alone [36]. The extended-mind thesis further asks when reliably available external artifacts participate in a cognitive process rather than merely serving as passive background [5]. PRA does not imply that a transformer is embodied or that neural attention reproduces human memory. These traditions instead offer a useful design question: what changes when external structure is not compressed entirely into internal parameters, but remains available through stable, action-conditioned references?

Human activity routinely relies on referential scaffolding. A person remembers that evidence is in a notebook, a clause is in a contract section, a procedure is in a manual, or an explanation belongs to a colleague. The remembered address, cue, and access routine can be more useful than a lossy internal copy. Retrieval is often situated: the object sought depends on the current task, physical or social environment, and sequence of prior actions. Hierarchical cues support progressive refinement from place or document to section and detail.

PRA translates only the computational aspects of this observation. Reference handles are cues; URIs stabilize identity; resolvers implement access routines; routing gists and optional summaries support selection; and layer-specific caches provide a temporary active representation. Recursive anchors make access path-dependent. The analogy predicts that a useful model memory system should preserve provenance and affordances, not only vector similarity. It also predicts failure when the environment is unreliable, addresses become stale, or access policy is ignored.

This perspective broadens evaluation beyond factual recall. A referential memory architecture should be tested on navigation, coordination, revision, and action: whether it can locate the right object, notice when it changed, follow a dependency, respect a permission, and update a plan after new evidence. Those are situated-memory behaviors, not merely next-token improvements.

Table 4: Conceptual foundation-model memory stack. Entries are architectural roles, not claims about current product implementations.

Layer	Typical content	Lifetime	Access mode
Parametric memory	linguistic and learned regularities	model version	dense forward pass
Working memory	current goals, dialogue, active evidence	sequence/session	local attention
Episodic memory	prior interactions and outcomes	user/project history	event retrieval
Semantic memory	documents, schemas, concepts, policies	corpus version	addressed retrieval
Procedural memory	tools, workflows, APIs, permissions	deployment	explicit action
World-state memory	environment, simulations, live systems	dynamic	observation/model update

11 A Future Foundation-Model Memory Stack

Foundation models are likely to require a hierarchy rather than one universal memory mechanism. Table 4 sketches distinct roles. The boundaries are conceptual and may be implemented by multiple coupled systems.

PRA is best viewed as an addressing and materialization fabric across this stack. It does not replace weights, local context, databases, tool protocols, or world models. It could allow compact working memory to refer to episodic events, semantic objects, and procedural capabilities; runtime policy would decide which representations can cross into model-local memory. Different classes require different consistency. Immutable literature can use content-addressed caches, while live operational state may require short leases and mandatory revalidation.

A complete stack also needs a control plane. It must assign identity, authenticate principals, describe provenance, negotiate capabilities, track versions, account for cost, and revoke access. The neural data plane then consumes authorized representations. MCP provides one possible application-level discovery and transport interface [25]; PRA addresses the separate question of selective internal materialization.

12 Enterprise and Professional Applications

Enterprise information is unusually compatible with context graphs because it already has identifiers, ownership, revisions, links, and permissions. The opportunity is not to pour every enterprise object into a prompt, but to expose a governed address space whose active working set follows the task.

12.1 Software Engineering

A software task can begin from repository, branch, commit, package, symbol, test, issue, and build-log handles. GitHub objects and IDE symbol graphs provide stable navigation edges; PRA could keep interfaces and summaries active while materializing implementation bodies or historical diffs only when needed. Required controls include commit-aware cache keys, generated-file labeling,

repository permission isolation, and tests that detect stale symbol memory. SWE-bench-style tasks and repository navigation benchmarks can measure whether progressive access reduces loaded tokens without hiding required dependencies.

12.2 Analytics, BI, and SQL

Data work spans warehouses, catalogs, SQL schemas, semantic layers, dashboards, lineage graphs, and business glossaries. A Snowflake table, BI metric, or transformation can be a versioned reference with ownership and freshness metadata. The model may first inspect a semantic metric definition, then expand the underlying model, columns, joins, and query history. Evaluation must penalize stale schemas, unauthorized columns, invalid SQL, and semantically wrong but executable queries. Live execution should remain an explicit tool action, while read-only schema memory is a candidate for caching.

12.3 Organizational Knowledge and Operations

Confluence pages, Jira issues, runbooks, incidents, and service catalogs form a cross-linked operational memory. Handles can preserve the distinction between approved policy, draft discussion, resolved incident, and current ticket state. Freshness and authority are central: an old runbook should not silently outrank a current service notice. Useful tasks include incident diagnosis, change-impact analysis, status reporting, and tracing a decision from ticket to documentation and code.

12.4 Scientific Literature

Scientific assistants need identity at article, version, section, figure, dataset, and claim level. Recursive anchors can move from abstract to method, result, appendix, and cited source. Provenance traces should distinguish quoted evidence from model inference. Benchmarks should test multi-paper synthesis, contradictory findings, retractions, and whether the model can expose the exact evidence path rather than only produce plausible prose.

12.5 Legal and Medical Assistance

Legal and medical domains make the risks clearest. Jurisdiction, effective date, patient identity, consent, and source authority are part of meaning, not optional metadata. PRA could help preserve these constraints through versioned, permissioned references and auditable access paths. It must not turn latent memory into an unreviewable decision channel. High-stakes outputs require authoritative-source validation, explicit uncertainty, human review, and strict separation between informational retrieval and actions affecting rights, treatment, or records. Evaluation should emphasize provenance completeness, stale-evidence detection, privacy leakage, and calibrated abstention rather than fluency.

13 Illustrative Design Space

Table 5 gives a conceptual comparison. The entries are hypotheses and design claims, not measured results.

Table 5: Conceptual comparison of context-access strategies. This table is illustrative and does not report experimental measurements.

Strategy	Context interface	Strength	Failure mode
Flat long context	Eager token sequence	Simple and general	High cost when relevance is sparse
One-shot RAG	Retrieved prompt passages	Good factual grounding	Retrieval granularity fixed before generation
Agentic search	External tool loop	Flexible progressive disclosure	High latency and orchestration cost
PRA	URI-addressed memory cache	Demand-driven latent expansion	Requires learned/runtime selection policy

14 Historical Feasibility Evidence

This section preserves the experimental path by which the proposal was corrected. The opening pilot and adaptation study use the optional adapted cross-attention transport, with a separately normalized memory branch, $W_{O,\text{mem}}$, and a residual scale. They are labeled historical and must not be read as measurements of canonical native-KV PRA. Beginning with “First Native-KV Quality Gate,” the experiments instead use selected host-native K/V and progressively repair fragmentation, routing, bounded execution, and off-device residency. Those studies remain small controlled probes; the companion papers provide the complete protocols and later pretrained-model evidence.

A standalone PyTorch pilot first provided limited implementation evidence for parts of this position. At 524,288 processed tokens and seeds 1 and 7, mean plain WikiText-2 losses were 4.7691 for SelfAttention, 4.7457 for full PRA, and 4.7574 for a two-SelfAttention/two-PRA hybrid. Its version-1 reference task suggested a memory-path contribution, but regenerated data per model seed and labeled all candidate parts as relevant. More importantly, the old implementation routed one raw summary embedding per reference, collapsed a batch query in an early path, and materialized whole references. We therefore retain all reference-conditioned measurements as engineering history rather than use them for a causal claim about the corrected architecture.

A corrected controlled follow-up uses paired seeds {1, 7, 21, 42, 87}; a 1,048,576-token budget; a fixed tokenizer; a fixed 512-example reference corpus and split; and SelfAttention controls parameter matched through FFN width. Table 6 summarizes ordinary-language results. Because external memory is absent in this phase, the comparison tests backbone compatibility. Both PRA variants finish below their matched controls in every paired seed, but five pairs permit only a minimum two-sided exact sign-flip p -value of 0.0625. This is encouraging directional evidence, not proof of superiority or equivalence.

The corrected reference follow-up compares training from scratch, frozen-reference-path adaptation, and direct joint fine-tuning. For full PRA, respective valid-reference losses are 6.3557, 4.6720, and 4.9129; for hybrid PRA they are 6.3642, 4.7296, and 4.9196. Frozen adaptation improves over scratch by 1.6837 and 1.6346, with the same effect direction in every paired seed. It exactly preserves the plain WikiText losses 4.5034 and 4.5500 because the base is frozen. Direct joint tuning increases plain loss by 0.1775 for full PRA and 0.1562 for hybrid PRA and ends 0.2408 and 0.1899 worse than frozen adaptation. It is therefore not sufficient under this budget. Full PRA benefits at least as much from staging as the hybrid. Extending full-PRA ordinary-language pretraining from the historical 524,288-token pilot to 1,048,576 tokens lowers mean loss by 0.2423, indicating benefit

Table 6: Five-seed corrected-router ordinary-language follow-up. Mean $\pm$ population standard deviation; lower is better.

Architecture	Parameters	Test loss
SelfAttention, same width	4,216,320	$4.5704 \pm .0076$
Full PRA	4,479,488	$4.5034 \pm .0163$
Hybrid PRA	4,347,904	$4.5500 \pm .0112$
SelfAttention, matched to full	4,478,976	$4.5855 \pm .0347$
SelfAttention, matched to hybrid	4,347,648	$4.5715 \pm .0082$

from the longer tested schedule without establishing how long is enough.




Figure 7: Paired ordinary-language losses for historical cross-attention-capable backbones. Each line connects one seed to its nearly parameter-matched SelfAttention control; diamonds mark means. The same direction across five seeds gives an exact two-sided sign-flip $p = 0.0625$, so the figure motivates replication rather than a superiority claim. It does not measure native-KV quality.

The corrected causal result remains negative. After frozen adaptation, full-PRA loss is 4.6720 with valid references, 4.6895 with disabled memory, 4.6769 with shuffled references, 4.6723 with irrelevant references, and 4.6702 with oracle content. Hybrid values are 4.7296, 4.7413, 4.7304, 4.7304, and 4.7279. The memory branch changes predictions, but shuffled penalties are at most 0.0048, irrelevant penalties are below 0.001, and oracle gains are below 0.002. Full and hybrid top-1 are 0.4221 ± 0.0156 and 0.3981 ± 0.0397 against a candidate-count-weighted chance rate of 0.4202. Pretraining improves language competence and final loss, not generalized selection speed. The likely next intervention is a differentiable or explicitly supervised router rather than a larger claim.

A post-refactor smoke rerun now exercises the corrected batched router at fixed total split counts 2 and 5. Twenty tiny CUDA runs share 32 source entries, one BPE tokenizer, context length 128, a fixed data-generation and split seed, and an eight-step budget while pairing model seeds {1, 7, 21, 42, 87} across architecture and split count. For full PRA, mean valid-reference loss improves over




Figure 8: Historical cross-attention adaptation and counterfactual results. Frozen-path adaptation gives the lowest reference-task loss while preserving the frozen backbone. Positive counterfactual deltas indicate worse loss than valid references. The much larger disabled effect than shuffled, irrelevant, or oracle effects exposes branch use without demonstrating content-causal retrieval.

disabled-reference loss by 0.0189 ± 0.0030 at split 2 and 0.0312 ± 0.0048 at split 5. Every seed has a positive gap, establishing reproducible branch activation. Mean shuffled-reference changes are only 0.0002 ± 0.0006 and -0.0006 ± 0.0009 , however, and split-5 reference top-1 is 0.150 ± 0.102 against a 0.25 uniform-candidate rate. Correct content is therefore not causally established. SelfAttention is exactly invariant to the reference interventions, as required of the control. Because split count also changes the target span and processed tokens, these values remain end-to-end regression and cost diagnostics rather than an isolated split-count effect.

These measurements do not establish the central efficiency thesis. Reference encoding was performed separately for every example on a small GPU, selective triggering was disabled, and cache reuse was not measured. They instead show why the conceptual separation matters: naming and a reference-specific residual can be engineered while selection can still fail independently. The corrected implementation adds the instrumentation and controls needed for a cleaner test, but software completion is not empirical confirmation.

14.1 First Native-KV Quality Gate

A subsequent five-seed synthetic experiment isolates native transport from routing and natural-language competence. A two-layer, width-64 SelfAttention model learns to predict a random binary value stored at a fixed historical marker among 63 source units. A disjoint 64-example evaluation set holds the complete source, one-token tail, answer, model weights, and example order fixed while repartitioning source units into 1, 2, 4, 7, 15, 31, or 63 references. Full-context SelfAttention reaches mean loss 0.0105 and 100% answer-token accuracy; the tail alone reaches loss 1.6999 and 50.6%, confirming that the target depends on displaced history.

Oracle native memory remains 100% accurate from 2 through 32 splits. Its losses at 3, 5, 8, 16, and 32 splits are 0.0624, 0.0653, 0.0711, 0.0900, and 0.1194, compared with shuffled losses 2.7645, 2.6853, 2.5565, 2.4730, and 2.7038. Content-causality gaps therefore remain between 2.38 and 2.70

synthetic native-KV benchmark (five paired seeds)




Figure 9: Historical independently encoded split scaling over five paired seeds. Oracle memory remains content-causal and accurate through 32 splits, but all-memory composition and the untrained router degrade under fragmentation. Later contextual/native slicing removes the 64-split failure; this figure is a diagnostic, not the corrected PRA scaling limit.

loss units. This demonstrates that weight-preserving native transport can recover causally required content when given the correct reference, without a learned memory projection or PRA fine-tuning.

The negative controls remain equally important. In this historical diagnostic, oracle accuracy falls to 70.6% at 64 splits. Native-all accuracy falls to 71.9% at split 3, 59.1% at split 5, and approximately chance by split 16; unsupervised routing is also near chance at fine partitions. Independent reference encoding restarts positions and does not reproduce a causally encoded full prefix, so practical native-all need not match the exact historical-K/V reinsertion identity. The result supports sparse oracle transport under moderate independent fragmentation. The later contextual-slicing study changes its interpretation: it is evidence for separating encoding and routing granularity, not a 64-way limit of corrected PRA.

14.2 Natural-Text Transport Gate

A five-seed HotpotQA-derived probe [40] replaces synthetic distractors with supporting-fact and distractor text while retaining a controlled **AnswerCode** `yes|no` anchor. This is not a HotpotQA benchmark score: the original question and evidence remain in metadata, but the local target isolates transport from the language competence of a tiny from-scratch decoder. Training and validation examples are disjoint, and the same 63 source units and target are repartitioned across all split counts.

Dense SelfAttention loss is 0.0027 versus 1.7038 for the tail. At split 32, oracle loss is 0.0060, oracle RCB is 0.996, and active token-K/V fraction is 0.092. Routed loss is 0.6812 and routed RCB is 0.634, so selection rather than sparse oracle transport is the dominant remaining limitation. Shuffled loss is 2.2251, yielding a 2.219 oracle content-causality gap. At split 64, mean oracle RCB falls to 0.546 with large seed variance. The result supports the narrower scaling claim that active K/V can fall while benefit is retained over a measured range; it rejects an unrestricted fragmentation claim and does not establish genuine multi-hop QA.




Figure 10: HotpotQA-derived transport probe over five paired seeds. Oracle native K/V preserves nearly all dense-context benefit through split 32 with a falling active fraction, while routed selection leaves a substantial error gap. The historical independently encoded split-64 condition exposes contextual fragmentation later repaired by native-K/V slicing.

The five-seed QASPER-derived replication [7] uses scientific-paper text and human evidence annotations under the same answer-code protocol. Dense loss is 0.0026 versus 1.8545 for the tail. At split 32, oracle loss is 0.0039, RCB is 0.999, and active fraction is 0.088; routed loss is 0.5550 and routed RCB is 0.602, while shuffled loss is 2.7351. At split 64, oracle RCB falls to 0.748. Thus two natural-text domains reproduce both the useful sparse-oracle frontier and the routing/fragmentation limits.




Figure 11: QASPER-derived replication over five paired seeds. Sparse oracle K/V preserves dense-context benefit through split 32, while all-memory composition and routed selection remain limiting; the historical 64-way result uses independent fragment encoding.

14.3 Sensitivity, Contextual Slicing, and Scale

A subsequent staged sensitivity study changes the interpretation of the 64-way failure. Increasing top- k helps the 32-way router but does not repair independently encoded 64-way microchunks, and multi-gist modes do not produce a robust cross-domain winner. The decisive intervention is to separate encoding context from routing granularity. Encoding eight consecutive regions together

restores five-seed annotated-oracle RCB to 0.996 on HotpotQA-derived text and 0.999 on QASPER-derived text; preserving global positions raises routed RCB to 0.865 and 0.993. Encoding the complete historical source once and slicing its native K/V yields 0.876 and 0.994 at approximately 10% active K/V. Thus the earlier 64-way result diagnoses contextual fragmentation rather than a hard limit on the number of named regions.

The scale follow-up then grows source length with addressability: nested sources contain 63, 127, and 255 units rather than repeatedly splitting the same short source. An ordinary SelfAttention model is trained on full 255-unit context and converted without PRA training. Historical source K/V is computed once, global positions are retained, and the tail continues after the displaced prefix. Native-all RCB remains 1.000 through 256. At $k = 8$, a 2.2–2.4M-parameter backbone retains routed RCB 1.000 on HotpotQA-derived text and 0.908 on QASPER-derived text at 256 units; a 6.6–6.9M-parameter backbone obtains 1.000 on both. Active K/V is only 4.4–5.1%. QASPER-derived MRR at 256 rises from 0.240 to 0.443 with capacity, showing that base-model representation quality affects inference-only routing geometry.

The original resource result was conditional. Retrieved K/V stayed nearly constant at about 11–13 tokens, so active fraction approximately halved as source size doubled, but the pre-tensorization Python router grew to 0.51 seconds per tiny example and 0.84–1.03 seconds per small example at 256 units. Figure 14 preserves that historical baseline; it is no longer the current routing implementation.

14.4 Exact Routing, Reuse, and Prompt-Head Evidence

The current exact router packs normalized gists, scores them in matrix operations, and applies hierarchical `torch.topk`. It changes neither selected chunk semantics nor the subsequent full native-K/V materialization. A five-seed comparison on the same real encoded caches separates scalar routing, cold tensorized routing, and repeated warm queries. At 256 units, scalar latency is 563–581 ms for tiny and 1,164–1,183 ms for small; cold tensorized latency is 14.8–18.9 and 29.6–33.7 ms; persistent-index latency is 5.5–6.3 and 11.5–11.6 ms. Index construction contributes roughly 45–54% of cold routing. URI, chunk, rank, score, and loss parity hold in every paired condition. This is an exact routing result, not an end-to-end serving result.

The implicit-head experiment uses a trained fixed-target answer-code model, a 24-token direct tail, total prompts of 24/48/96/192 tokens, four target locations, eight examples, and five seeds. Historical `#_head` routing matches dense mean loss at 48 and 96 tokens and improves over both dense and truncation loss at 192 while activating 45.8% of K/V. Target-chunk recall is 1.000, 0.969, and 0.894 at 48, 96, and 192 tokens. At 192, routed loss is 1.365 versus 1.772 for truncation and 1.723 for wrong head memory; independently encoded head chunks degrade to 1.929. This establishes recovery outside the direct window on the controlled task, not general long-context language understanding.

14.5 Bounded Logical Context and Rollover

A five-seed model-scaled follow-up constrains every native operation to 32 tokens while retaining an 8-token direct tail and four reserve tokens. Sixteen-token encoding cores with four context-only overlap tokens produce four-token routing chunks. Heads of 24/56/120/184 tokens require 2/4/8/12 encoding calls; the maximum observed call is 20 tokens, with no limit violation. Thus 192 logical tokens are six times the native limit, the 184-token head is 5.75 times that limit, and the largest observed bounded encoding call uses only 62.5% of the ceiling. At 192 total prompt tokens, historical routed loss is 1.057 versus 1.229 for truncation and 1.709 for independent encoding

PRA fragmentation sensitivity, successive halving (seed count shown)




Figure 12: Fragmentation decomposition at 64 units. Larger causal encoding context and global positions restore the transport ceiling while retaining independently addressable regions.




Figure 13: Five-seed context/addressability scaling at routed $k = 8$. Historical slicing preserves the transport ceiling; capacity repairs the tiny QASPER-derived routing decline. These are answer-code transport probes, not unrestricted QA.




Figure 14: Pre-tensorization prototype baseline. Fixed- k K/V transfer remains nearly constant, while candidate-wise Python/CUDA synchronization dominates routing.




Figure 15: Exact routing on real encoded caches. A persistent packed index makes warm selection nearly flat through 256 units and approximately 94–103 $\times$ faster than the legacy scalar path at the largest scale.




Figure 16: Controlled implicit-head results. Historical encode-once slicing preserves displaced evidence; truncation, wrong-memory, and independent-encoding controls separate access, content causality, and contextual fragmentation.

blocks. Routed target recall is 0.525; an approximate oracle obtains 0.874 loss, while wrong chunks obtain 1.120. The remaining wrong-memory benefit makes this supporting evidence, not a complete content-causality result.




Figure 17: Five-seed implicit-head execution under a 32-token native bound. Logical prompt length reaches 192 while encoding and selected attention remain bounded.

Score-ordered whole-chunk admission exposes a quality–active-memory curve. Increasing the memory budget from 4 to 20 tokens raises target recall from 0.325 to 0.525 and lowers loss from 1.242 to 1.057. Mean forward latency remains 18.5–19.3 ms on this tiny GPU, so the sweep supports bounded allocation but no speed claim. In a separate 48-token generation smoke test, 44 expired tokens enter `#_head` through 11 rollovers; 40 subsequent steps route memory, the direct suffix never exceeds eight tokens, and the largest native operation is 20 tokens. Rebuilding the complete head after each rollover is correct but not yet efficient.

14.6 Prototype GPU Capacity Reduction

An experimental residency policy keeps routing gists and their packed index on GPU while storing complete native token K/V in pinned CPU memory. Selection transfers only chosen chunks. At 256 units this removes 0.73–0.85 MiB of tiny and 2.94–3.42 MiB of small source K/V from GPU, lowering measured peak allocation by approximately the same amount with zero loss or selected-chunk delta. Selected transfer remains 24.6–25.4 KiB for tiny and 85.7–87.9 KiB for small, costing 4.0–4.7 and 8.4–8.9 ms. Stable-cache warm requests rise from 18.2–18.5 to 22.5–22.6 ms for tiny and from 32.7–35.6 to 45.9–46.2 ms for small. Cold construction is substantially slower because the prototype performs many small synchronous offloads; batched offload is required before a serving claim.




Figure 18: Whole-chunk admission at 192 logical tokens. More active K/V improves recall and loss on the controlled probe; production latency remains unestablished.




Figure 19: GPU- versus CPU-resident native K/V at fixed $k = 8$. CPU backing removes complete source K/V from GPU while selected transfer stays bounded; peak allocated memory falls, but transfer and unbatched cold offload add latency.

15 Evaluation Doctrine

PRA requires counterfactual evaluation because a single task score cannot reveal how memory was used. At minimum, a promising checkpoint should be evaluated with valid references, a disabled reference path, shuffled references, irrelevant references, empty references, oracle references, and all evidence visible in local context. Targets and prompt structure must remain fixed across these conditions. Shuffling must occur across examples while preserving reference count and approximate length, and batched caches must prevent cross-example leakage.

Four families of outcomes should be reported together:

1. **Task quality:** loss, perplexity, calibrated accuracy, or task-specific utility.
2. **Reference and chunk causality:** valid-versus-disabled, valid-versus-shuffled, valid-versus-irrelevant, oracle-chunk, full-reference, and gist-only deltas.
3. **Preservation:** ordinary language and non-reference task quality before and after adaptation.
4. **Resource use:** processed tokens, available versus selected memory tokens, synchronized cache/routing/attention latency, bucket padding ratio, throughput, peak memory, cache bytes, and reuse rate.

Evidence of genuine reference use requires more than a disabled-path gap. A branch can act as a generic domain adapter regardless of supplied content. Valid references should outperform shuffled and irrelevant content, ideally by larger margins on examples proven to be locally insufficient. Reference routing and chunk routing should separately rank known relevant objects and evidence spans above controls. Claims should be repeated across seeds and reported with matched token, parameter, and compute budgets.

16 Research Agenda

The original agenda proceeded through controlled standalone models, positional validation, and pretrained integration. Those stages now have dedicated companion papers. The remaining agenda begins at the measured one-shot boundary rather than repeating completed integration work.

Standalone controlled models. Continue to use small decoder-only transformers for mechanism isolation. The corrected studies now cover paired counterfactuals, native-KV transport, contextual slicing, fixed- k and fixed-fraction scaling, learned routing geometry, and model-bounded prompts. Their role is to falsify mechanisms cheaply, not to stand in for pretrained task evaluation.

Pretrained-model integration. Paper 2 wraps pretrained Hugging Face families while preserving rotary conventions, grouped-query layout, cache formats, and exact disabled behavior. The released SDK makes one-shot PRA usable for research and sparse direct-memory workloads. The next integration question is not whether wrapping is possible, but how routing and conditional memory use generalize across domains, depths, model scales, and realistic retrieval errors.

Runtime and serving integration. Study how reference memories are resolved, encoded, cached, paged, shared, and evicted. Add asynchronous host-to-device transfer, fused route-and-attend kernels, batching, approximate indexes, and matched dense/RAG baselines. The central metric is no longer only accuracy, but quality per materialized token, millisecond, transferred byte,

and resident byte. The current SDK is usable; a generally compelling production architecture and its serving economics are not yet established.

Scaling and theory. Run the decisive 32K–8M+ logical-memory protocol defined above on task families where relevant context occupies a controllable fraction of the available corpus. Separate active GPU native-K/V from CPU/reference storage and search/index cost. PRA should be evaluated against full-context, sparse-attention, RAG, summary-first RAG, and agentic-search baselines.

Selection and gating. Compare projected content gists, explicit reference-token attention, optional summaries, learned routers, uncertainty triggers, and hybrid policies while holding selected native-KV transport fixed. Paper 1.5 shows that reproducing native Q/K geometry does not guarantee a useful semantic index, and Paper 2 shows domain-dependent learned routing. Reference and chunk selection should be evaluated through recall–sparsity curves, MRR, NDCG, span IoU, counterfactual task loss, and learning speed, not only whether an expected URI entered a cache.

Hierarchical expansion. Paper 2.5 implements bounded associative discovery for tasks in which the initial query identifies an intermediate relation but not the final evidence. It establishes synthetic mechanism controls and natural native-edge topology, while also exposing weak executable entry routing, granularity-dependent missing edges, and broad activation. The next step is not another uncontrolled depth sweep: it is calibrated online entry/search control followed by evidence-centered materialization. Measure expansion depth, branching factor, missed-evidence rate, repeated-route cost, and quality per resolved and materialized token. Flat single-hop tasks cannot test the progressive part of PRA.

Training and preservation. Compare isolated reference-path adaptation, joint fine-tuning, adapters, low-rank updates, and curriculum mixtures. Report both reference improvement and local-task retention. Determine whether zero-initialized memory residuals remain stable at larger scale and whether learned gates avoid permanent saturation.

Safety, provenance, and access control. Treat reference resolution as a privileged operation. Study signed or versioned URIs, tenant isolation, policy-aware cache keys, provenance traces, revocation, poisoning detection, and whether the model can infer protected content through shared memory representations.

16.1 Falsification Criteria

The research program should be abandoned or substantially revised if repeated, well-controlled studies show any of the following:

- valid references do not outperform shuffled or irrelevant references once formatting and capacity are controlled;
- equivalent prompt concatenation or RAG matches quality at lower end-to-end cost;
- reference encoding and cache management erase any savings from shorter local context, even with reuse;
- learned expansion policies fail to outperform simple oracle-budgeted retrieval or cannot maintain recall;

- local language ability degrades unacceptably under the adaptation needed to use memory;
- URI and cache isolation cannot be enforced under realistic batching and multi-tenant serving.

These criteria make PRA a testable systems-and-model hypothesis rather than a label that can absorb any retrieval mechanism after the fact.

17 Risks and Limitations

PRA introduces several risks. First, reference selection can fail silently: if the system does not select the right chunk, the model may answer from incomplete context. Second, URI-addressed memory creates provenance and security questions. A reference table is a powerful interface; poisoning or misaddressing entries can mislead the model. Third, cache construction may move cost rather than eliminate it. Encoding references separately is only beneficial when cache reuse, selective transfer, or reduced prompt length outweighs encoding overhead. Fourth, pretrained integration is family-specific because external K/V must respect layer semantics, position encodings, attention variants, and batching constraints. Paper 2 demonstrates exact adapters for multiple tested families, but not universal compatibility.

Fifth, content-independent branch effects can be mistaken for retrieval. A memory residual may improve a domain simply because it contributes additional activations, even when references are shuffled. Sixth, batching can violate semantic isolation: a physically shared cache must not expose one sequence’s references to another, and empty rows must not inherit projection bias. Seventh, preservation claims can be artifacts of freezing rather than architecture. Eighth, hierarchical expansion can amplify latency variance and create denial-of-service surfaces when references recursively fan out. Ninth, raw summaries can introduce a representation mismatch when compared with projected layer queries; optional summary evidence must inhabit the same contextual layer space. Tenth, implicit prompt heads move historical encoding into a separate cold pass; bounded blocks approximate full historical context, and streaming currently rebuilds the complete head after rollover rather than appending K/V incrementally. Eleventh, CPU residency reduces measured GPU capacity but the current per-chunk offload makes cold construction slower. Finally, stale cache entries create a consistency problem: URI, content version, model/tokenizer/routing fingerprints, and encoded layer memory must remain aligned.

These limitations motivate careful experimental design. PRA should not be judged by a single accuracy number. It should be evaluated through controlled ablations that measure answer accuracy, reference and chunk ranking, expansion depth and budget exhaustion, logical, requested, and materialized fractions, timing decomposition, padding allocation, cache hit ratio, content-specific counterfactual deltas, ordinary-path retention, and cache isolation. The series now addresses tiny-scale capacity, positional continuity, learned semantic routing, exact pretrained retrofit, long-prompt displacement, selective transfer, measured GPU residency, and bounded associative discovery. Sparse end-to-end iterative benefit, evidence-centered materialization, batched asynchronous offload, multi-million-token scaling, and efficiency against strong production baselines remain open.

18 Conclusion

Progressive Retrieval Attention proposes a shift in how foundation models interact with memory. Rather than treating context as a flat token budget to be filled eagerly, PRA treats references as first-class computational objects that can be named, selected, materialized, attended to, expanded,

reused, revoked, and audited. The model’s logical context may then be much larger than its active working set, just as a program’s address space can exceed resident memory. Unlike conventional paging, however, neural materialization is approximate and learned; success must be demonstrated, not presumed.

The long-term vision is a foundation-model memory stack with explicit boundaries among parametric knowledge, working context, episodic history, semantic resources, procedures and tools, and changing world state. URI identity, provenance, freshness, hierarchy, permissions, and cache lifetime become architectural concerns rather than application-side annotations. Protocols such as MCP can expose resources and actions, agents can manage high-level plans, and PRA-like runtimes can test whether repeated read-oriented access can move closer to attention without losing governance.

The historical cross-attention follow-up supports preservation controls but fails its content-causality gate. Native-KV experiments provide stronger mechanism evidence: contextual slicing preserves transport through 256 addressable units, fixed- k active fraction falls to roughly 5%, and model-bounded encoding demonstrates a logical head 5.75 times the native limit. Paper 1.5 adds positional continuity and separates semantic search from native attention geometry. Paper 2 then establishes exact disabled pretrained paths, cross-family adapters, and a useful one-shot QASPER regime while identifying HotpotQA as the multi-hop boundary. Paper 2.5 then demonstrates bounded associative discovery and native graph structure, while showing that entry competition, granularity, broad activation, and memory-use interference still limit sparse end-to-end gains. These results make PRA a usable research SDK and support the narrow claim that named logical context can grow faster than the K/V materialized in one operation. They do not yet show that PRA is generally better than optimized long context or RAG in production. The decisive next tests are evidence-centered materialization, adaptive serving control, and the matched 32K–8M+ scaling experiment; failure there should narrow or reject the broader virtual-context thesis.

References

- [1] Ali Behrouz, Peilin Zhong, and Vahab Mirrokni. Titans: Learning to memorize at test time. *arXiv preprint arXiv:2501.00663*, 2025.
- [2] Iz Beltagy, Matthew E. Peters, and Arman Cohan. Longformer: The long-document transformer. *arXiv preprint arXiv:2004.05150*, 2020.
- [3] Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George B. M. van den Driessche, Jean-Baptiste Lespiau, Bogdan Damoc, Aidan Clark, et al. Improving language models by retrieving from trillions of tokens. In *International Conference on Machine Learning*, 2022.
- [4] Shouyuan Chen, Sherman Wong, Liangjian Chen, and Yuandong Tian. Extending context window of large language models via positional interpolation. *arXiv preprint arXiv:2306.15595*, 2023.
- [5] Andy Clark and David J. Chalmers. The extended mind. *Analysis*, 58(1):7–19, 1998.
- [6] Zihang Dai, Zhilin Yang, Yiming Yang, Jaime Carbonell, Quoc V. Le, and Ruslan Salakhutdinov. Transformer-XL: Attentive language models beyond a fixed-length context. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 2978–2988, 2019.

- [7] Pradeep Dasigi, Kyle Lo, Iz Beltagy, Arman Cohan, Noah A. Smith, and Matt Gardner. A dataset of information-seeking questions and answers anchored in research papers. In *Proceedings of NAACL*, 2021. Preprint: <https://arxiv.org/abs/2105.03011>.
- [8] Mostafa Dehghani, Stephan Gouws, Oriol Vinyals, Jakob Uszkoreit, and Lukasz Kaiser. Universal transformers. In *International Conference on Learning Representations*, 2019.
- [9] Peter J. Denning. The working set model for program behavior. *Communications of the ACM*, 11(5):323–333, 1968.
- [10] Yiran Ding, Li Lina Zhang, Chengruidong Zhang, Yuanyuan Xu, Ning Shang, Jiahang Xu, Fan Yang, and Mao Yang. LongRoPE: Extending LLM context window beyond 2 million tokens. *arXiv preprint arXiv:2402.13753*, 2024.
- [11] In Gim, Guojun Chen, Seung-seob Lee, Nikhil Sarda, Anurag Khandelwal, and Lin Zhong. Prompt cache: Modular attention reuse for low-latency inference. *arXiv preprint arXiv:2311.04934*, 2023.
- [12] Alex Graves. Adaptive computation time for recurrent neural networks. *arXiv preprint arXiv:1603.08983*, 2016.
- [13] Alex Graves, Greg Wayne, and Ivo Danihelka. Neural turing machines. *arXiv preprint arXiv:1410.5401*, 2014.
- [14] Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. *arXiv preprint arXiv:2312.00752*, 2024.
- [15] Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Ming-Wei Chang. Retrieval augmented language model pre-training. In *International Conference on Machine Learning*, 2020.
- [16] Dongseong Hwang, Weiran Wang, Zhuoyuan Huo, Khe Chai Sim, and Pedro Moreno Mengibar. TransformerFAM: Feedback attention is working memory. *arXiv preprint arXiv:2404.09173*, 2024.
- [17] Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. Dense passage retrieval for open-domain question answering. In *Proceedings of EMNLP*, 2020.
- [18] Urvashi Khandelwal, Omer Levy, Dan Jurafsky, Luke Zettlemoyer, and Mike Lewis. Generalization through memorization: Nearest neighbor language models. In *International Conference on Learning Representations*, 2020.
- [19] Nikita Kitaev, Lukasz Kaiser, and Anselm Levskaya. Reformer: The efficient transformer. In *International Conference on Learning Representations*, 2020.
- [20] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention. In *Proceedings of the ACM SIGOPS Symposium on Operating Systems Principles*, 2023. Preprint: <https://arxiv.org/abs/2309.06180>.

- [21] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Kuttler, Mike Lewis, Wen-tau Yih, Tim Rocktaschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Advances in Neural Information Processing Systems*, 2020.
- [22] Yuhong Li, Yingbing Huang, Bowen Yang, Bharat Venkitesh, Acyr Locatelli, Hanchen Ye, Tianle Cai, Patrick Lewis, and Deming Chen. SnapKV: LLM knows what you are looking for before generation. In *Advances in Neural Information Processing Systems*, 2024. Preprint: <https://arxiv.org/abs/2404.14469>.
- [23] Guangda Liu, Chengwei Li, Jieru Zhao, Chenqi Zhang, and Minyi Guo. ClusterKV: Manipulating LLM KV cache in semantic space for recallable compression. *arXiv preprint arXiv:2412.03213*, 2024.
- [24] Zirui Liu, Jiayi Yuan, Hongye Jin, Shaochen Zhong, Zhaozhuo Xu, Vladimir Braverman, Beidi Chen, and Xia Hu. KIVI: A tuning-free asymmetric 2bit quantization for KV cache. In *Proceedings of the 41st International Conference on Machine Learning*, pages 32332–32344, 2024.
- [25] Model Context Protocol Contributors. Model context protocol specification. <https://modelcontextprotocol.io/specification/2025-06-18/>, 2025. Protocol revision 2025-06-18.
- [26] Amirkeivan Mohtashami and Martin Jaggi. Landmark attention: Random-access infinite context length for transformers. In *Advances in Neural Information Processing Systems*, 2023.
- [27] Tsendsuren Munkhdalai, Manaal Faruqui, and Siddharth Gopal. Leave no context behind: Efficient infinite context transformers with infini-attention. *arXiv preprint arXiv:2404.07143*, 2024.
- [28] Jack W. Rae, Anna Potapenko, Siddhant M. Jayakumar, and Timothy P. Lillicrap. Compressive transformers for long-range sequence modelling. In *International Conference on Learning Representations*, 2020.
- [29] Aurko Roy, Mohammad Saffar, Ashish Vaswani, and David Grangier. Efficient content-based sparse attention with routing transformers. *Transactions of the Association for Computational Linguistics*, 9:53–68, 2021. Preprint: <https://arxiv.org/abs/2003.05997>.
- [30] Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Eric Hambro, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems*, 2023.
- [31] Jorge Simao. Iterative progressive retrieval attention: Bounded associative discovery over latent gist memory. Companion Paper 2.5 manuscript, 2026.
- [32] Jorge Simao. Position, attention geometry, and semantic routing for retrieved native-KV memory. Companion Paper 1.5 manuscript, 2026.
- [33] Jorge Simao. Progressive retrieval attention: Bounded sparse native-KV for addressable logical memory. Companion Paper 1 manuscript, 2026.
- [34] Jorge Simao. Progressive retrieval attention for pretrained transformers: Sparse native-K/V memory with semantic routing. Companion Paper 2 manuscript, 2026.

- [35] Sainbayar Sukhbaatar, Arthur Szlam, Jason Weston, and Rob Fergus. End-to-end memory networks. In *Advances in Neural Information Processing Systems*, 2015.
- [36] Francisco J. Varela, Evan Thompson, and Eleanor Rosch. *The Embodied Mind: Cognitive Science and Human Experience*. MIT Press, 1991.
- [37] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, 2017. Preprint: <https://arxiv.org/abs/1706.03762>.
- [38] Yuhuai Wu, Markus N. Rabe, DeLesley Hutchins, and Christian Szegedy. Memorizing transformers. In *International Conference on Learning Representations*, 2022.
- [39] Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. StreamingLLM: Efficient streaming language models with attention sinks. In *International Conference on Learning Representations*, 2024.
- [40] Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William W. Cohen, Ruslan Salakhutdinov, and Christopher D. Manning. HotpotQA: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of EMNLP*, 2018. Preprint: <https://arxiv.org/abs/1809.09600>.
- [41] Zichao Yang, Diyi Yang, Chris Dyer, Xiaodong He, Alex Smola, and Eduard Hovy. Hierarchical attention networks for document classification. In *Proceedings of NAACL-HLT*, 2016.
- [42] Jiayi Yao, Hanchen Li, Yuhua Liu, Siddhant Ray, Yihua Cheng, Qizheng Zhang, Kuntai Du, Shan Lu, and Junchen Jiang. CacheBlend: Fast large language model serving for RAG with cached knowledge fusion. *arXiv preprint arXiv:2405.16444*, 2024.
- [43] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. React: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations*, 2023.
- [44] Manzil Zaheer, Guru Guruganesh, Avinava Dubey, Joshua Ainslie, Chris Alberti, Santiago Ontanon, Philip Pham, Anirudh Ravula, Qifan Wang, Li Yang, and Amr Ahmed. Big bird: Transformers for longer sequences. In *Advances in Neural Information Processing Systems*, 2020.
- [45] Zhenyu Zhang, Ying Sheng, Tianyi Zhou, Tianlong Chen, Lianmin Zheng, Ruisi Cai, Zhao Song, Yuandong Tian, Christopher Ré, Clark Barrett, Zhangyang Wang, and Beidi Chen. H2O: Heavy-hitter oracle for efficient generative inference of large language models. In *Advances in Neural Information Processing Systems*, 2023. OpenReview version: <https://openreview.net/forum?id=RkRrPp7GK0>.