

Positional Continuity for Retrieved Native-KV Memory: Logical Offsets, RoPE Rebinding, and the Limits of Geometric Repair

Jorge Simao

EInnovator R&D – Center for the Sciences of Computation, Cognition, and Complexity

August 12, 2026

Abstract

Chunked key/value (K/V) memory introduces a positional ambiguity that ordinary prefix caching largely avoids. A token can belong at position 100 in its source yet be encoded at position zero because its chunk was processed independently. The resulting tensor is valid, but it no longer represents the source coordinate or attention relation that retrieval later assumes. We study this failure in Progressive Retrieval Attention (PRA), which uses compact gists to select chunks and then attends to their full-detail native K/V.

A source-relative logical offset is the general repair. Across learned absolute, sinusoidal, and rotary position embedding (RoPE) models, offsets remove layer-0 reset error exactly and reduce final-layer K error in all 30 paired comparisons over two capacities and five seeds. RoPE adds a narrower capability: pre- and post-RoPE keys are exactly equivalent at matched effective positions, while pre-RoPE storage permits retrieval-time rebinding. On the measured unfused CUDA path, that flexibility increases warm latency from 0.179 to 0.649 ms.

The central result is a separation rather than a universal quality gain. WikiText-2 representation improves in 30/30 pairs, but next-token loss improves in only 5/30; controlled QA-derived loss improves in 35/60 pairs and reverses direction for some RoPE conditions. Changing RoPE distance alone yields no stable useful optimum. By contrast, a frozen 32-dimensional learned metric over position-independent hidden states reaches .908–.985 one-of-15 evidence recall, versus .369–.723 for post-RoPE K and .000–.015 under shuffled supervision. Increasing centered-RoPE gist resolution from one to eight raises native-Q/K rank fidelity from .794 to .963, but sparse semantic AUC remains .651–.652; the learned metric reaches .971 AUC despite only .178 Q/K fidelity. For fixed selected identities, post-RoPE K, V, and native attention output remain bit-exact. A 192-token logical prompt also runs with every measured native operation at or below 24 tokens under a 32-token limit. Positional continuity is necessary, but coordinate correctness, representation fidelity, attention geometry, semantic selection, and downstream use are distinct empirical properties.

1 Introduction

A key retrieved from external memory carries more than content. Its hidden state reflects the causal context present when it was encoded, and its attention score depends on a position. If a long source is divided into bounded blocks and every block starts again at zero, later tokens silently inherit the coordinates of earlier ones. The cache remains a well-formed tensor, but it misrepresents the source.

This problem is broader than RoPE. Learned absolute embeddings and sinusoidal embeddings also change when a block is reset. RoPE is special for a different reason: because position is an

explicit rotation of queries and keys, it provides an algebra for preserving or changing a cached key’s binding. The offset requirement is general; the rebinding operation is rotary.

We study both problems in PRA, a sparse native-KV memory mechanism. A long source or displaced prompt head is encoded in bounded blocks. Compact gists route a later query to selected chunks, but attention consumes the chunks’ original token-level K/V. The position carried by those keys must therefore be a cache contract rather than an implementation accident.

We organize the problem into five increasingly demanding levels. Retrieval selection and downstream use are separate because a router can find annotated evidence while the model still fails to use that composition:

1. **Coordinate correctness:** were source-relative positions preserved?
2. **Representation fidelity:** does chunked K/V resemble the state produced with the intended coordinates and context?
3. **Attention-geometry fidelity:** does the positional treatment reproduce the intended query–key relation?
4. **Semantic retrieval utility:** are relevant chunks ranked highly by a position-independent content representation?
5. **Downstream task utility:** does the selected state improve language-model loss or answers?

These levels form a dependency, not an equivalence. Correct coordinates do not recreate missing causal context. Faithful K/V can still be selected poorly. Even annotated evidence can be an inferior context composition for a particular model. This leads to the paper’s headline result:

Geometric correctness and functional usefulness are not the same thing.

The paper makes eight outcome-oriented contributions:

- We formalize source-relative logical position as cache provenance for independently encoded K/V.
- We show that logical offsets repair reset fragmentation for learned absolute, sinusoidal, and RoPE models, rather than only for rotary models.
- We verify exact pre-/post-RoPE equivalence and quantify the runtime cost of deferred rotation and retrieval-time rebinding.
- We separate coordinate repair from overlap-based contextualization and measure the quality–encoding-cost tradeoff.
- We demonstrate that systematic representation improvement need not produce monotonic language-model or QA-derived loss improvement.
- We demonstrate that semantic evidence routing can use a tiny learned hidden-state projection while selected native post-RoPE K/V remains unchanged for attention, separating semantic relevance geometry from positional attention geometry.
- We show that averaging post-RoPE keys mixes positional phases, and that increasingly faithful approximations of native Q/K geometry need not improve semantic evidence retrieval.
- We demonstrate bounded implicit-head execution in which logical prompt length exceeds every native model operation, while carefully separating this mechanical result from dense long-context equivalence.

2 What Position Does a Retrieved Token Own?

Suppose a source span occupies positions 100 through 131. Independent chunk encoding often does this:

View	First token	Last token
Source/physical position	100	131
Wrong chunk-local reset	0	31
Correct logical/effective position	100 + 0	100 + 31

The correction is simply


$$p_{\text{effective}} = p_{\text{chunk,start}} + p_{\text{inside}}, \quad (1)$$

where $p_{\text{chunk,start}}$ is stored once for a contiguous chunk and $p_{\text{inside}} \in [0, T - 1]$ is its local index. The formula restores the coordinate the token owns inside its source.

The word *source-relative* is important. Positions 100–131 are meaningful within one document, reference, or displaced prompt head. Two unrelated documents do not acquire a meaningful global ordering merely because both are registered in one request. Source identity and logical offset are therefore cache provenance. A continuously advancing `#_head`, by contrast, represents one prompt history and retains one coordinate system as direct tokens roll into memory.

3 What Does PRA Store and Retrieve?

PRA is self-contained in the following path:



The gist is routing metadata, not a compressed substitute for the selected token state. Once a chunk is selected, its full token-level native K/V joins direct-context K/V under the same attention softmax. Position errors introduced while building that K/V survive routing and materialization; the router cannot repair them afterward.

Three lengths remain distinct. L_{logical} is the addressable source span, $L_{\text{model-op}}$ is the maximum K/V participating in one native operation, and L_{position} is the supported coordinate range. PRA can make $L_{\text{logical}} > L_{\text{model-op}}$ by caching and selecting bounded regions, but it does not by itself guarantee that a model has learned useful behavior at every L_{position} .

Table 1 gives the empirical answers before the details. Each row addresses a different question; no row certifies the one below it.

Table 1: Headline results. Each row answers a different correctness question.

Question	Verified result	Meaning
Do offsets repair reset coordinates?	Layer-0 K RMSE becomes zero in 30/30 paired comparisons	Source-relative continuity is a general requirement
Do deeper representations improve?	Final-layer K RMSE falls in 30/30 pairs	The repair survives two capacities and three mechanisms
Does overlap improve representation?	Yes in the verified WikiText scope; 29/30 loss pairs also improve modestly	Contextualization is separate and costs $1.94\times$ encoding at 8L
Are pre/post RoPE keys equivalent?	Zero K, logit, and output error at matched effective positions	Storage state does not change the intended geometry
Does better geometry guarantee task quality?	No: WikiText loss improves in 5/30; QA routed loss in 35/60	Functional utility is a separate empirical property
Can deferred rotation rebind keys?	Yes; warm latency rises from 0.179 to 0.649 ms	Flexibility has a measured systems cost
Must semantic routing reuse native attention geometry?	No: a 32-D learned projection reaches .908-.985 one-of-15 recall, versus .369-.723 for post-RoPE K; shuffled labels reach .000-.015	Semantic selection and positional attention can use different geometries
Can logical prompt exceed a native operation?	192 logical tokens, maximum measured operation 24 under a limit of 32	Bounded positional and transport mechanics

The mechanisms can now be stated without conflating their roles:

Table 2: Five levels of memory correctness and utility. Success at one level does not certify the next.

Level	Question	Main intervention or result
Coordinate correctness	Are source-relative positions correct?	Logical offsets remove chunk-local reset error.
Representation fidelity	Are chunk states contextually faithful?	Overlap or grouped encoding reduces contextual fragmentation.
Attention-geometry fidelity	Is native positional Q/K geometry preserved or deliberately rebound?	RoPE placement controls the native attention relation.
Semantic retrieval utility	Are relevant chunks ranked highly?	A learned 32-D projection exposes an evidence-relevance geometry.
Downstream task utility	Does selected memory improve loss or answers?	Composition and model use remain empirical even under an evidence oracle.

4 Do Logical Offsets Repair More Than RoPE?

The general experiment compares independently encoded 32-token blocks with a dense 128-token representation. It covers two model capacities, learned absolute, sinusoidal, and RoPE positions, and five matched seeds per condition. Table 3 reports layer-0 and final-layer K RMSE.

Table 3: Reset versus source-relative offsets. Every final-layer paired difference favors offsets; every offset layer-0 error is exactly zero.

Tier	Position	Reset L0	Offset L0	Reset final	Offset final	Improved
Tiny	Learned absolute	.6673	.0000	.6769	.1129	5/5
Tiny	Sinusoidal	.3640	.0000	.4722	.1436	5/5
Tiny	RoPE	.6463	.0000	.6745	.1373	5/5
Small	Learned absolute	.6626	.0000	.7193	.3066	5/5
Small	Sinusoidal	.3345	.0000	.5507	.2746	5/5
Small	RoPE	.5586	.0000	.7498	.4235	5/5




Figure 1: Offsets reduce final-layer K error in all 30 paired runs. The residual after exact layer-0 repair reveals a second problem: missing causal context across encoding blocks.

The sinusoidal result rules out a RoPE-only explanation. Logical offsets repair coordinate reset because they restore source position, regardless of whether position enters through an embedding table, a formula-defined additive embedding, or rotary Q/K geometry.

The residual final-layer error is equally informative. At layer zero, matched token content and effective positions suffice for exact equality. At later layers, a block lacks the causal left context available in dense encoding. This is contextual fragmentation, not failed coordinate repair.

5 What Does Overlap Repair?

Encoding overlap gives each block additional native left context and commits only its new core. It does not change routing granularity and should not be described as a routing intervention. In the controlled representation test, 25% overlap lowers final-layer mean error in all six tier-mechanism conditions. At 50%, five of six improve; small learned absolute RMSE moves slightly from 0.3066 to 0.3103. The intervention is useful but not monotonic.

WikiText-2 gives a clear quality/cost point. At an 8:1 logical-to-native ratio, overlap lowers final-layer error in every condition and next-token loss in 29/30 paired runs relative to offset-only encoding. The loss reductions are modest, from 0.0007 to 0.0173 in the tiny tier and 0.0102 to 0.0148 in the small tier, while encoding processes 1.9375 times the unique source tokens. Overlap buys contextualization through recomputation.




Figure 2: On WikiText-2, offsets sharply improve representation fidelity without a matching loss gain. Overlap gives smaller, more consistent loss changes at almost $1.94\times$ encoding cost at 8L.

Across the QA-derived probes, overlap effects are mixed and mechanism-dependent. That does not make overlap a failed coordinate method; it confirms that it repairs a different input to the task pipeline.

6 What Does RoPE Uniquely Add?

6.1 Intuition Before Algebra

RoPE rotates each query and key according to its position. Attention sees the difference between those rotations, so translating both a query and key together preserves their relative geometry. Translating only the key changes its distance from a fixed query and is not an invariant operation.

This yields two storage choices:

Post-RoPE K	Pre-RoPE K
Position is bound when the cache is written. Retrieval is cheap, but the binding is fixed.	Position is applied when the cache is read. Retrieval-time rebinding is possible, with extra work.

6.2 The Minimal Algebra

For a head of width d_h , let R_p denote the RoPE rotation at position p . The rotated query and key are

$$\tilde{q}_i = R_i q_i, \quad \tilde{k}_j = R_j k_j. \quad (2)$$

Because $R_i^\top R_j = R_{j-i}$,

$$\tilde{q}_i^\top \tilde{k}_j = q_i^\top R_{j-i} k_j. \quad (3)$$

Equation 3 answers one question: why does common translation preserve the query-key relation? Adding c to both positions leaves $j - i$ unchanged. It does not say that $R_i q_i$ has the same relation to $R_{j+c} k_j$.

Pre- and post-RoPE storage are algebraically equivalent when they use the same effective position:

$$K_{\text{post}}(p) = R_p K_{\text{pre}}. \quad (4)$$

The four-cell model test confirms zero K, logit, and output RMSE, with top-token agreement 1.0, for matched offset positions across both tiers, all layers, and five seeds. Reset binding instead has mean K RMSE 0.648 and output RMSE 0.063. Deliberate rebinding by +512 changes the output, as it should.




Figure 3: Matched pre- and post-RoPE storage reproduces the same attention geometry exactly. Reset coordinates and deliberate K-only rebinding change it.

The pure tensor test reaches the same conclusion. Translating Q and K together preserves the top-attended token, with maximum logit RMSE 5.79×10^{-5} through the measured shifts. Moving K alone by up to 8,192 positions changes fixed-query scores. This is the correct behavior, not an error in RoPE.

6.3 Deferred Rotation Has a Cost

At 160 selected small-tier tokens, storing post-RoPE K gives a 0.1791 ms warm path. Storing pre-RoPE K, reconstructing positions, and rotating at materialization gives 0.6487 ms. The position reconstruction itself is small; the unfused RoPE transform dominates. The nearly flat timing across token counts indicates launch-dominated behavior on the measured GTX 950M, so this is a prototype cost measurement rather than a production throughput prediction.

Both modes retain the same number of K/V elements. For 192 stored small-tier tokens, K/V uses 196,608 bytes. A contiguous chunk needs one 8-byte logical start for execution; a full 192-token int64 audit vector would use 1,536 bytes. Deferred rotation changes when phase is assigned, not the size of the native K/V payload.

6.4 Retrieval-Time Placement Preserves Local Geometry

Deferred rotation exposes a narrower control than an arbitrary position reset. Let a selected chunk have source positions $s_0, \dots, s_{n-1}$ and let the current query be at q . We define the nearest-token distance D and translate the chunk as a unit:

$$s'_j = s_j + (q - D - s_{n-1}). \quad (5)$$

Then $q - s'_{n-1} = D$, while $s'_j - s'_\ell = s_j - s_\ell$ for every pair of tokens. The experiment therefore changes the global query-memory displacement but preserves token order and all within-chunk spacing. Exact placement leaves s_j unchanged; local placement sets $D = 1$; clipped placement uses $\min(D_{\text{source}}, D_{\text{max}})$.

This distinction matters because RoPE does not simply lose a fixed amount of resolution with distance. It represents relative displacement through many periodic frequency pairs: high-frequency pairs vary rapidly and cycle often, while lower-frequency pairs vary more slowly. A displacement can be mathematically valid yet fall outside the distance distribution for which the learned Q/K projections became useful. Equation 5 tests that functional question without changing memory content.

7 Why Does Better Geometry Not Guarantee Better Loss?

The central empirical discovery appears when representation and task metrics are placed on the same axis. On 8L WikiText-2 windows, offsets reduce final-layer K RMSE in 30/30 paired runs but reduce next-token loss in only 5/30. Learned-absolute loss rises by 0.0200 in the tiny tier and 0.0292 in the small tier; sinusoidal loss rises by 0.0040 and 0.0707; RoPE loss is unchanged to the reported precision.

The controlled HotpotQA- and QASPER-derived answer-code probes likewise resist a universal story. Routed loss improves in 35/60 paired comparisons. Learned-absolute and sinusoidal means improve in both capacities and datasets, whereas RoPE means worsen in all four dataset-capacity cells. These probes use natural source text but simplified targets; they measure positional and retrieved-memory behavior, not unrestricted QA.




Figure 4: Every point moves left because offsets improve representation fidelity, but task-loss changes cross zero. Correct geometry is not a sufficient predictor of usefulness.

This is not a failed universal-gain result. A universal improvement would have hidden the fact that these are different properties. A model trained on short reset blocks may not have learned to use the repaired distance distribution. Routing gists may rank different chunks after positions change. Selected evidence may lack surrounding support, and distractors can alter the model’s computation. Position repair makes the state better defined; it does not choose the state or teach the model how to use it.

8 Does RoPE Distance Explain Task Quality?

This section stays within level 3, attention-geometry fidelity. It asks a specific hypothesis: perhaps exact source-relative RoPE is geometrically correct but functionally suboptimal because retrieved memory is too distant from the query. Can virtual rebinding help? The task results also leave a contextual alternative. Very small memory fragments might lack enough contextual state, or exact source placement might expose a learned relative-distance mismatch. We separate encoding context from materialized-unit size and begin with oracle evidence so routing does not obscure the comparison.

8.1 Fixed Distance Changes Attention, but Has No Stable Optimum

The fixed- D sweep reuses the same raw K , V , query, evidence, and local chunk geometry for every policy. It evaluates $D \in \{32, 64, \dots, 4096\}$ on tiny RoPE checkpoints over HotpotQA- and QASPER-derived probes, using five matched seeds. Exact nearest-token source distances average 87.3 and 90.4, respectively.




Figure 5: The same retrieved evidence under different RoPE displacements. Dotted horizontal lines show exact-placement loss; the vertical line marks the pooled selected value $D = 64$. Bands are one standard deviation over five seeds. Remote placement can sharply damage loss and top-token agreement, but the useful near range has no stable winner.

The sweep establishes that placement is not inert. At $D = 1024$, HotpotQA-derived mean loss rises from .0655 under exact placement to 1.1855, and QASPER-derived loss rises from .0377 to .8192. Agreement with the exact policy’s top-attended token falls to .090 and .067. The response is not monotonic: $D = 2048$ is less harmful than $D = 1024$, consistent with RoPE’s multi-frequency periodic geometry rather than a scalar distance-decay story.

The near-range result resolves the hypothesis negatively. In the two-tier confirmation matrix, fixed $D = 64$ beats exact placement in 50 of 100 matched dataset–tier–seed–composition cells and loses in 50; mean loss changes by $+8.1 \times 10^{-5}$. Clipping only source distances above 64 wins 68/100 cells, but its mean change is merely -9.3×10^{-5} . Log compression wins 24/100 and changes mean loss by $+1.3 \times 10^{-4}$. These effects are far smaller than the remote-distance failures and do not establish a task-optimal virtual distance.

8.2 Context and Materialized Composition Matter More

The encoding-context ablation jointly encodes one, two, or four adjacent references while retaining five routing units. Its evidence-only oracle row is invariant: the evidence span is causally first, so later references cannot alter its hidden state. This is a useful design check, not evidence that context never matters. When all encoded neighboring K/V is retained, small-tier HotpotQA-derived loss falls from .0235 with one reference per block to .0026 with four; QASPER-derived loss falls from .0936 to .0026. Tiny QASPER improves modestly from .0388 to .0380, while tiny HotpotQA is essentially flat.

The materialized-unit ablation keeps the encoding regime fixed and changes source partitioning. Exact-position oracle loss generally worsens as the source is divided more finely. HotpotQA-derived tiny loss is .0607, .0655, and .0775 for 2, 5, and 16 units; small-tier loss is .0026, .0030, and .0040. QASPER-derived small-tier loss similarly moves from .0026 to .0027 to .0036; tiny QASPER has a shallow five-unit plateau before rising to .0426 at 16 units. A larger selected unit includes more neighboring content, so this probe identifies a composition benefit rather than a pure representation-size law.




Figure 6: Left: larger encoding groups help when the resulting neighboring K/V is retained. Center: finer source partitioning generally raises oracle loss. Right: larger deviation from exact RoPE phase does not map monotonically to task loss. Error bars show one standard deviation over five seeds.

The combined diagnosis is therefore an interaction dominated here by contextualization and composition. Their effects are larger and more stable than the small near-range distance changes. Exact source placement is substantially better than a full local reset in the tiny-tier oracle rows, and a small near-query rebase does not reliably improve on exact placement. Remote rebinding can be destructive, but the measured models do not support a universal replacement geometry. For continuous history, exact source coordinates remain the default. For independent references, where no natural serialization distance exists, a pretrained-model oracle gate should compare exact and a conservative clipped policy before general retrieval-time rebinding is enabled. The broad adaptive-geometry program is not justified by this fixed- D evidence.

9 Can Routing Use a Different Geometry?

The placement sweep closes one route to better selection: changing RoPE distance alone does not produce a stable optimum. We therefore test a positive alternative. Native RoPE remains responsible for attention, while a small position-independent metric ranks chunks. Let h_q be the final query token’s attention-input hidden state and g_c the mean attention-input hidden state of candidate chunk c . Two bias-free linear maps define

$$z_q = A_q h_q, \quad z_c = A_c g_c, \quad s_{\text{route}}(q, c) = \cos(z_q, z_c), \quad (6)$$

where $A_q, A_c \in \mathbb{R}^{32 \times d_{\text{model}}}$. The transformer, positional encoding, Q/K/V projections, and payload path remain frozen. Only A_q and A_c receive a within-example margin-ranking loss against the programmatic evidence identity.

A tiny learned semantic geometry sharply improves evidence ranking without altering native positional K/V.




Figure 7: Semantic memory selection and positional attention can use different geometries without changing the selected native K/V payload.

Protocol. We reuse the established 16-way HotpotQA- and QASPER-derived partitions, giving 15 candidate references per example, both model capacities, and five seeds. Each reference is encoded independently at its exact logical RoPE offset. We fit on the established 80% training split and evaluate the held-out complement. The four non-oracle comparisons are mean post-RoPE Q/K, mean pre-RoPE Q/K, untrained hidden-state cosine, and the learned hidden-state projection. A fifth router has identical initialization and optimization but within-example evidence labels are replaced by equally many non-evidence labels. The target is the benchmark’s explicit answer-code evidence anchor; this isolates evidence-identity ranking but is not unrestricted question-conditioned QA retrieval.

Because 15 candidates are indexed, the requested R@5% and R@10% cutoffs realize one and two references, or 6.7% and 13.3%. Each example has one target identity, so any-evidence, evidence-identity, and all-evidence recall coincide in this probe.

Table 4: Held-out evidence-identity recall at the nominal 5% point (one of 15 candidates), mean over five seeds. Parentheses give the learned projection’s population standard deviation. f_{80}/f_{90} are mean selected fractions required to reach .80/.90 recall.

Probe	Tier	Post-K	Pre-K	Hidden	Learned	Shuffled	f_{80}/f_{90}
HotpotQA	Tiny	.723	.446	.077	.908 (.031)	.015	.067/.080
HotpotQA	Small	.400	.446	.000	.985 (.031)	.000	.067/.067
QASPER	Tiny	.585	.369	.138	.938 (.090)	.000	.080/.080
QASPER	Small	.369	.369	.000	.969 (.038)	.000	.067/.067




Figure 8: The learned 32-D hidden-state metric dominates raw hidden cosine and the shuffled-label control throughout the sparse range. Post-RoPE and pre-RoPE Q/K remain useful controls but do not define the best evidence-ranking geometry. Curves average five seeds.

At the two-reference point, learned recall is 1.000, 1.000, .985, and 1.000 in table order. Learned sparse AUC_{0-30} is .962–.980 and MRR is .954–.992. Its AUC exceeds raw hidden cosine and shuffled supervision in all 5/5 seeds of every dataset–tier cell. It also exceeds post-RoPE Q/K in all small-tier seeds and all QASPER-tiny seeds; HotpotQA-tiny is less uniform at 3/5 despite a positive mean difference. Thus the result is not that every learned initialization beats every K-space baseline. It is that evidence supervision exposes a routing geometry absent from both raw hidden cosine and the matched label-noise control.

Raw hidden-state cosine is an important negative result: one-reference recall is .000–.138, despite the learned maps receiving exactly those hidden states. Semantic information can be present in contextual hidden states while cosine in the native hidden coordinates remains a poor relevance metric. The asymmetric projection exposes a more useful relevance geometry. The shuffled-label projection collapses to .000–.015 recall, showing that the gain is learned evidence relevance rather than a generic benefit from low dimension or reparameterization.

The router adds 4,096 parameters in the tiny tier and 8,192 in the small tier, or .31–.79% of the dataset-specific backbone. A 15-candidate float32 index falls from 3,840 to 1,920 bytes in the tiny tier and from 7,680 to 1,920 bytes in the small tier. The corresponding full native detail-K/V averages 52.9–55.3 KB and 105.9–110.6 KB. For a fixed selected identity list, post-RoPE K is identical, V is identical, and native attention output is identical: every comparison has zero maximum absolute error. The projection changes only the search representation, not the selected payload or its use by attention. In short, routing geometry $\neq$ attention payload geometry.

This closes the placement loop in the controlled scope. Exact, local, fixed, clipped, and compressed RoPE placements did not yield a stable beneficial fixed distance. Learned routing does not repair that geometry; it avoids asking positional attention geometry to serve as a semantic relevance metric. Larger contextual encoding improves the representation being summarized, learned

projection improves evidence ranking, and RoPE placement controls native positional attention. They are separate axes.

This contextualization result independently reinforces Paper 1’s separation between encoding and routing granularity: fine addressability does not require equally fine independent encoding.

9.1 Does an Evidence Oracle Guarantee Better Use?

It does not. An evidence-only oracle is not universally best. Across dataset, tier, and position mechanism cells, the annotated evidence-only condition is worse than the router-selected set for 3.3%–78.3% of examples. In the continuous-head tiny-RoPE condition, moving from offset-plus-overlap routing to the constructed evidence oracle raises mean loss from 0.8763 to 1.0124 even though evidence recall becomes one. Annotated evidence identity is therefore not necessarily the functionally optimal context composition for a model. The result is a controlled warning, not a universal composition law.

10 Does Pooling Preserve Native RoPE Geometry?

Token-level post-RoPE keys are the correct native attention representation. Compact routing, however, usually replaces several token keys with one arithmetic mean. This operation mixes phases. For pre-RoPE token keys K_j^{pre} at source positions p_j , native keys are

$$K_j^{\text{post}} = R(p_j)K_j^{\text{pre}}. \quad (7)$$

A mean formed after rotation and a mean formed before rotation then bound to one representative center p_c are, in general,

$$\boxed{\frac{1}{m} \sum_j R(p_j)K_j^{\text{pre}} \neq R(p_c) \left(\frac{1}{m} \sum_j K_j^{\text{pre}} \right)}. \quad (8)$$

The left side is a valid mean of native keys, but it has no single RoPE phase. The right side has one coherent representative phase, but compresses the token-specific rotations. Thus exact token-level RoPE correctness does not make post-RoPE arithmetic pooling positionally coherent, and no one-center construction can preserve every token relation.

Centered subgists. For a chunk spanning $[p_{\text{start}}, p_{\text{end}}]$, the one-gist construction averages pre-RoPE keys and rotates the mean once at $p_c = (p_{\text{start}} + p_{\text{end}})/2$. For $G > 1$, we split the chunk into G balanced contiguous nonempty subspans, repeat this operation at each exact subspan center, and retain all subgist identities. Half-token centers are evaluated directly by RoPE rather than rounded; they occur in 30.4% of evaluated subspans. When a short chunk contains fewer than G tokens, G is capped at one nonempty subgist per token.

Geometry target and protocol. For query i and candidate chunk c , the exact native target is

$$s_{ic}^{\text{native}} = \max_{j \in c} \left[\frac{1}{H} \sum_{h=1}^H \frac{q_{ih}^\top R(p_j)k_{icjh}^{\text{pre}}}{\sqrt{d_h}} \right]. \quad (9)$$

Every gist method uses the same mean-over-heads then maximum-over-gists reducer. We report the per-example Spearman correlation ρ_{QK} between its 15-chunk ranking and Eq. 9, plus top-1

and top-3 agreement. Position bias is the per-example Pearson correlation between routing score and normalized chunk center; the table reports its absolute magnitude $|\rho_{\text{pos}}|$, while signed values remain in the artifacts.

The experiment reuses the frozen RoPE checkpoints, 16-way HotpotQA- and QASPER-derived partitions, established 80/20 split, tiny and small tiers, and seeds 1, 7, 21, 42, and 87. We compare post-RoPE mean, position-neutral pre-RoPE mean, centered $G = 1, 2, 4, 8$, raw hidden cosine, and the same 32-D learned router from Section 9. Semantic metrics use the programmatic evidence identity. Reported means first aggregate the four dataset–tier cells within each seed; 95% intervals use the five resulting seed means. The attention payload is always native token-level post-RoPE K and native V.

Table 5: Native-Q/K geometry and semantic evidence retrieval. ρ_{QK} includes a 95% confidence interval over five seed-level means. R10/R20 denote evidence-identity recall at nominal 10%/20% candidate fractions; AUC is sparse AUC_{0-30} , and f_{80}/f_{90} are selected fractions needed for .80/.90 recall.

Method	ρ_{QK}	$ \rho_{\text{pos}} $	R10	R20	AUC	f_{80}	f_{90}
Post-RoPE mean	.802 $\pm$.049	.590	.631	.688	.653	.350	.523
Pre-RoPE mean	.408 $\pm$.071	.239	.527	.596	.548	.600	.683
Centered $G = 1$	.794 $\pm$.052	.589	.631	.692	.651	.373	.517
Centered $G = 2$	.815 $\pm$.050	.588	.627	.662	.639	.393	.537
Centered $G = 4$	.847 $\pm$.040	.597	.638	.692	.653	.373	.493
Centered $G = 8$	.963 $\pm$.011	.573	.627	.681	.652	.370	.483
Hidden cosine	−.078 $\pm$.147	.237	.100	.154	.131	.917	.963
Learned 32-D	.178 $\pm$.058	.338	.996	1.000	.971	.070	.073




Figure 9: Left: centered subgists increasingly reproduce exact native token-Q/K chunk rankings. Right: that geometric improvement does not produce a corresponding gain in sparse semantic AUC. Error bars are 95% confidence intervals over five seed-level means.

The G sweep supplies the cleanest controlled contrast. From $G = 1$ to $G = 8$, mean ρ_{QK} rises by .169, and it rises in all 20 dataset–tier–seed conditions. Top-1 agreement rises from .723 to .946 and top-3 overlap from .715 to .923. Yet semantic AUC moves only from .6510 to .6516: 10

conditions improve, 8 decline, and 2 tie. R10 changes from .631 to .627 and R20 from .692 to .681. Diagnostic R@3/R@8/MRR are .692/.931/.653 at $G = 1$ and .681/.923/.662 at $G = 8$. Better approximation of the host attention ranking therefore does not imply better evidence ranking in this probe.

The learned router makes the dissociation sharper. It has only .178 Q/K rank fidelity, yet reaches .971 sparse AUC, .996 R10, and 1.000 R20. It is not an approximation of native Q/K geometry; it is a supervised relevance geometry. Conversely, post-RoPE mean is a strong Q/K approximation (.802) but only a moderate semantic index (.653 AUC). These results strengthen the distinction between levels 3 and 4 in the paper’s framework:

$$\boxed{\text{better native-attention approximation} \not\Rightarrow \text{better semantic retrieval.}} \quad (10)$$




Figure 10: Absolute correlation between routing score and normalized source position. Centered pooling does not monotonically remove position dependence; its signed mean correlations are negative ($-.569$ to $-.595$). The learned relevance metric reduces, but does not eliminate, the positional association. Error bars show 95% confidence intervals over five seeds.

For every method, fixed selected identities yield zero maximum absolute difference in selected native K, selected native V, and attention output. Pooling changes only the compact search representation. The result does not make token-level post-RoPE K wrong: native post-RoPE K is the correct attention payload. It shows why a mixed-phase mean can be a useful approximation of attention ranking and still be the wrong objective for semantic indexing.

11 Can Logical Prompts Exceed Native Operations?

The implicit `#_head` reference provides an intuitive long-prompt test. As the direct tail advances, displaced prompt tokens become one request-local source. Their logical positions continue rather than restarting at each rollover. The experiment uses $L_{\text{logical}} = L_{\text{position}} = 192$ and enforces $L_{\text{model-op}} = 32$.

Every measured encoding and attention call contains at most 24 tokens; native-limit violations are zero in all six tier-mechanism conditions. Source-relative offsets lower mean answer loss in 27/30 paired seed comparisons. Overlap lowers it further in 23/30. The effect size remains mechanism-dependent, and the small RoPE mean stays high despite the mechanical correction.

Table 6: Continuous-head mean answer loss over four examples and five seeds. All rows use 192 logical positions, a 32-token operation limit, and a maximum observed operation of 24.

Tier	Position	Reset routed	Offset routed	Offset + overlap
Tiny	Learned absolute	2.3837	.7304	.6474
Tiny	Sinusoidal	2.9691	.9388	.8021
Tiny	RoPE	2.5910	1.5547	.8763
Small	Learned absolute	3.8306	.6681	.0030
Small	Sinusoidal	5.0421	3.3860	.9110
Small	RoPE	6.7887	6.2249	3.0396

This demonstrates bounded positional and transport mechanics: the logical prompt exceeds one native operation, source-relative coordinates remain continuous, and no measured operation breaks its limit. It does not establish full functional equivalence to dense processing of all 192 tokens through every layer.

12 Related Work and Positioning

RoPE makes attention depend on relative displacement through position-dependent rotations [10]. Position Interpolation and LongRoPE extend dense-context operating ranges through rescaling and progressive extension [1, 4]. ALiBi and Transformer-XL provide other relative or distance-aware treatments [9, 2]. This paper is not another RoPE scaling recipe: it studies the ownership and rebinding of positions when K/V is encoded in chunks, stored externally, and retrieved sparsely.

Shortformer separates absolute positional information from reusable values [8]. Prompt Cache uses position-aware schemas for modular cache reuse, and CacheBlend selectively recomputes independently cached chunks that lack preceding context [5, 14]. MiniPIC similarly delays RoPE binding, while SemPIC learns a cache writer to address residual context dependence [7, 12]. These systems reinforce the distinction between position assignment and contextual hidden-state construction.

StreamingLLM preserves selected chronological decode states [11]. PRA instead routes among URI- or request-addressed chunks and must define both source identity and position semantics. Sparse attention and KV compression concern which state remains active or resident; the present question is whether a retained state still represents the coordinate and context its consumer assumes.

13 Limitations and Scope

The models are tiny and small controlled transformers, not pretrained production LLMs. The HotpotQA- and QASPER-derived tasks retain natural source text but replace open-ended answering with answer codes [13, 3]. WikiText-2 tests next-token behavior in bounded windows [6]. These choices isolate mechanisms but limit external validity.

Five paired seeds provide direction counts but weak inferential resolution: the smallest exact two-sided sign-flip result with five pairs is $p = 0.0625$. Formula-defined RoPE and sinusoidal coordinates can be evaluated outside the training range, but mathematical availability is not evidence of learned extrapolation. Learned absolute positions remain bounded by their table.

Overlap approximates missing left context through additional encoding and does not recreate a fully dense history at every layer. The evidence oracle is a programmatic source-identity control,

not a guarantee of optimal context composition. Deferred-RoPE timing uses an unfused prototype path on one older GPU and should not be projected to production serving.

The fixed-distance and context-gate probes use scratch-trained tiny and small models with a 256-token training context. Their simplified answer-code loss reveals remote-placement failures, but small near-range deltas may not transfer to pretrained long-context models. Encoding-block size changes the causal context of later references, whereas materialized-unit size changes included neighboring content; neither is a pure scalar chunk-size intervention. The exact evidence-only context row is structurally invariant because evidence is encoded first.

The learned-router confirmation is deliberately narrow. Its positive identity is an explicit answer-code anchor inside natural-text-derived sources, and its train and held-out examples come from the same controlled dataset family. It demonstrates that a tiny supervised metric can separate the mapped target from distractors without positional K-space, not that it solves open-ended HotpotQA or QASPER retrieval. We do not retrain memory use or claim a downstream loss gain from replacing the selector. Hard negatives, cross-domain transfer, and pretrained semantic structure require separate experiments.

The pooling-geometry experiment shares those evidence-identity and scale limitations. Its native target uses a maximum over token scores after averaging heads; other reducers can define a different approximation problem. Increasing G also increases compact-index entries and score work, so $G = 8$ is a fidelity control rather than a free serving recommendation. The study measures ranking and fixed-selection payload parity, not end-to-end answer quality after replacing the selector.

The paper deliberately excludes pretrained Hugging Face adaptation, learned memory-use training, recursive multi-hop retrieval, hidden-state multi-gist search, semantic-distance policies beyond the reported mechanism probes, and serving integration. Those extensions must not be read back into the present evidence.

14 Conclusion

The evidence supports seven conclusions. First, a source-relative logical start offset repairs reset fragmentation exactly at layer zero for learned absolute, sinusoidal, and RoPE models, and improves deeper representation in every paired test. Second, contextual fragmentation remains separate: overlap or grouped encoding can improve representation, but costs additional encoding and does not follow from correct coordinates.

Third, RoPE provides explicit algebraic Q/K rebinding, and matched pre-, post-, and deferred rotation preserve native attention geometry. That flexibility has a measured runtime cost, and no stable beneficial fixed retrieval distance emerged in the tested regime. Very remote placement can be destructive; the near-range response is non-monotonic rather than a universal distance rule.

Fourth, native positional attention geometry is not necessarily semantic relevance geometry. A mixed-phase post-RoPE mean approximates native token ranking well, and centered subgists raise that fidelity further, but the improvement does not raise semantic evidence retrieval. A routing representation may therefore be a poor semantic index even when it closely approximates the host attention geometry. Fifth, a tiny asymmetric projection of contextual hidden states sharply improves controlled evidence ranking. Sixth, routing can use that learned geometry while selected native post-RoPE K, V, and attention output remain bit-exact. Seventh, semantic retrieval success still does not guarantee downstream utility: annotated evidence and the functionally best context composition are not always identical.

The practical hierarchy is therefore: repair source coordinates, preserve contextually faithful

representations, define native positional-attention geometry, choose a semantic relevance geometry, retain native K/V for attention, and measure downstream use independently. **Position, context, attention geometry, routing relevance, and use are separate architectural problems.**

A Detailed Task Results

Table 7: WikiText-2 at an 8:1 logical-to-native ratio. Deltas are intervention minus baseline; negative loss is better. Every row contains five paired seeds.

Tier	Position	Δ final-K	Δ offset loss	Better	Δ overlap loss	Better
Tiny	Learned absolute	-.4281	+.0200	0/5	-.0173	5/5
Tiny	Sinusoidal	-.1897	+.0040	1/5	-.0007	4/5
Tiny	RoPE	-.6292	-.0000	2/5	-.0048	5/5
Small	Learned absolute	-.4044	+.0292	0/5	-.0140	5/5
Small	Sinusoidal	-.1810	+.0707	0/5	-.0102	5/5
Small	RoPE	-.6314	-.0000	2/5	-.0148	5/5

Table 8: Split-5 routed answer-token loss. Better counts paired offset-versus-reset seeds.

Dataset	Tier	Position	Reset	Offset	Offset + overlap	Better
HotpotQA	Tiny	Absolute	.559	.482	.482	4/5
HotpotQA	Tiny	Sinusoidal	.359	.304	.358	3/5
HotpotQA	Tiny	RoPE	.948	1.056	1.080	2/5
HotpotQA	Small	Absolute	.245	.043	.068	3/5
HotpotQA	Small	Sinusoidal	.061	.017	.012	4/5
HotpotQA	Small	RoPE	.129	.170	.196	1/5
QASPER	Tiny	Absolute	.629	.309	.301	4/5
QASPER	Tiny	Sinusoidal	.375	.244	.244	4/5
QASPER	Tiny	RoPE	.683	.869	.871	0/5
QASPER	Small	Absolute	.460	.299	.298	4/5
QASPER	Small	Sinusoidal	.403	.003	.007	5/5
QASPER	Small	RoPE	.169	.324	.184	1/5

B RoPE Storage Timing

Table 9: Small-tier CUDA timing in ms, averaged over five seeded runs of 500 warm iterations.

Storage	Selected K/V	Gather	Reconstruct	RoPE	Attention	Warm
Post-RoPE K	32	.0312	.0000	.0000	.1343	.1842
Pre-RoPE K + deferred	32	.0337	.0191	.4021	.1343	.6511
Post-RoPE K	160	.0327	.0000	.0000	.1342	.1791
Pre-RoPE K + deferred	160	.0329	.0145	.3838	.1342	.6487

C Implementation and Reproducibility

The implementation separates the native operation limit from position capacity. Each cached layer records token K/V, source identity, logical start/end, and whether K is pre- or post-RoPE. Contiguous token IDs are reconstructed as `logical_start + arange(length)`. Values are never rotated. Conversion copies the trained SelfAttention Q/K/V and output projections into PRA; with memory disabled, converted logits match the source model to test tolerance.

The final validation artifacts use commit `64f7cdf89b55`; generated cross-domain summaries use `052ae05`; the frozen distance/context run records source commit `03c9566ebbd7`. The original decomposition and storage matrix use `d36f67752701`. The learned-routing measurements use experiment implementation `d96c87b`; the pooling-geometry matrix records `527a50170802`. Each distance row records the checkpoint tokenizer fingerprint; this prevents a process-local tokenizer re-training from changing token IDs during evaluation. Raw CSV/JSON, compact publication CSVs, plots, expectations, and manifests are stored in the Paper 1.5 results directory. The reduction script regenerates manuscript artifacts without executing a model. Tests cover large logical offsets, overlap and splitting, independent operation/position bounds, cache metadata, no-memory conversion parity, pre-/post-RoPE equivalence, deliberate rebinding, retrieval policies, rollover, and CPU/GPU parity. Learned-routing tests additionally cover projection shapes, backbone freezing, deterministic ranking, shuffled-label construction, fixed-selection K/V and output parity, and CUDA scores. Pooling tests cover balanced contiguous subspans, exact fractional centers, the $G = 1$ construction, short-chunk capping, deterministic correlations and top- k agreement, and routing-representation independence from native payload materialization.

References

- [1] Shouyuan Chen, Sherman Wong, Liangjian Chen, and Yuandong Tian. Extending context window of large language models via positional interpolation. *arXiv preprint arXiv:2306.15595*, 2023.
- [2] Zihang Dai, Zhilin Yang, Yiming Yang, Jaime Carbonell, Quoc V. Le, and Ruslan Salakhutdinov. Transformer-XL: Attentive language models beyond a fixed-length context. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 2978–2988, 2019.
- [3] Pradeep Dasigi, Kyle Lo, Iz Beltagy, Arman Cohan, Noah A. Smith, and Matt Gardner. A dataset of information-seeking questions and answers anchored in research papers. In *Proceedings of NAACL*, 2021. Preprint: <https://arxiv.org/abs/2105.03011>.
- [4] Yiran Ding, Li Lyna Zhang, Chengruidong Zhang, Yuanyuan Xu, Ning Shang, Jiahang Xu, Fan Yang, and Mao Yang. LongRoPE: Extending LLM context window beyond 2 million tokens. *arXiv preprint arXiv:2402.13753*, 2024.
- [5] In Gim, Guojun Chen, Seung-seob Lee, Nikhil Sarda, Anurag Khandelwal, and Lin Zhong. Prompt cache: Modular attention reuse for low-latency inference. *arXiv preprint arXiv:2311.04934*, 2023.
- [6] Stephen Merity, Caiming Xiong, James Bradbury, and Richard Socher. Pointer sentinel mixture models. *arXiv preprint arXiv:1609.07843*, 2016.
- [7] Nathan Ordonez and Thomas Parnell. MiniPIC: Flexible position-independent caching in less than 100 lines of code. *arXiv preprint arXiv:2606.13126*, 2026.
- [8] Ofir Press, Noah A. Smith, and Mike Lewis. Shortformer: Better language modeling using shorter inputs. In *Proceedings of ACL-IJCNLP*, pages 5493–5505, 2021.
- [9] Ofir Press, Noah A. Smith, and Mike Lewis. Train short, test long: Attention with linear biases enables input length extrapolation. In *International Conference on Learning Representations*, 2022.
- [10] Jianlin Su, Yu Lu, Shengfeng Pan, Ahmed Murtadha, Bo Wen, and Yunfeng Liu. RoFormer: Enhanced transformer with rotary position embedding. *arXiv preprint arXiv:2104.09864*, 2021.
- [11] Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. StreamingLLM: Efficient streaming language models with attention sinks. *arXiv preprint arXiv:2309.17453*, 2023.
- [12] Hui Xie, Peng Xiao, Yutong Deng, Shuoran Dou, Jian Yang, and Jinyang Guo. SemPIC: Learning semantic position-independent KV caches. *arXiv preprint arXiv:2607.28069*, 2026.

- [13] Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William W. Cohen, Ruslan Salakhutdinov, and Christopher D. Manning. HotpotQA: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of EMNLP*, 2018. Preprint: <https://arxiv.org/abs/1809.09600>.
- [14] Jiayi Yao, Hanchen Li, Yuhan Liu, Siddhant Ray, Yihua Cheng, Qizheng Zhang, Kuntai Du, Shan Lu, and Junchen Jiang. CacheBlend: Fast large language model serving for RAG with cached knowledge fusion. *arXiv preprint arXiv:2405.16444*, 2024.