

PRA-SGLang: Query-Addressed Non-Prefix Memory beside RadixAttention

Jorge Simão
EInnovator R&D

September 2026

Abstract

SGLang’s RadixAttention reuses exact token prefixes. Progressive Retrieval Attention (PRA) selects versioned non-prefix resources and intervals according to the current query. We implement a logical PRA object namespace and an SGLang adapter and an MLX runner-cache bridge that keeps PRA memory outside Radix prefix accounting while consuming it as selected native K/V. A 25-request Apple-M5 smoke study with patched SGLang MLX separates exact radix reuse from selected-memory recovery. Prefix plus selected memory and full context both recover the fixed answer in 5/5 requests, while the no-evidence controls are 0/5. Selection uses 365 prompt tokens instead of 1,577. The scheduler reports 364 cached tokens on warm selected requests and 1,576 on warm full-context requests; warm TTFT is 34.1 and 52.7 ms, respectively. In a separate five-seed native-cache experiment, selected K/V recovers 5/5 answers and is bit-exact to SGLang’s ordinary split cache. All 28 attention layers consume selected memory under one normalization, while the scheduler-visible local offset excludes every PRA token. An external PRA HiCache then places immutable detail in attention-ready L1 arrays, host-owned L2 arrays, and compressed NPZ L3 backing without entering the Radix namespace. Across five seeds, answer recovery and Radix separation are 5/5; warm L1, L2-to-L1, and L3-to-L1 resolution average 0.009, 4.0, and 405.6 ms. Forced capacity pressure performs both L1-to-L2 and L2-to-L3 demotion in every seed. The same logical-object cache is then connected to SGLang’s built-in `HiCacheStorage` file backend. Five fresh-adapter reads recover 5/5 answers; raw 14.35 MB blobs write in 38.7 ms on average, restore to L1 in 122.2 ms on average (36.2 ms median), and hit warm L1 in 0.013 ms. An oracle prefetch sweep then rejects Python-thread overlap as a production mechanism: promotion takes 25–28 ms, but requested 10–50 ms leads delay the caller by 193–235 ms rather than overlapping useful work. In a combined runner-level cohort, an 81-token shared-prefix hit and external HiCache selection coexist in one request. Selected A and reselected C recover 5/5, cleanup request B recovers 0/5, and all selected requests contain exactly one memory copy. The bridge also runs through SGLang’s real prefill and batched-decode lifecycle: native and full context recover 5/5, disabled memory recovers 0/5, and native throughput rises from 6.59 to 14.54 requests/s between concurrency one and eight. Shared immutable K/V avoids 108.0 MiB at concurrency eight. Across 40 QASPER and 40 HotpotQA natural-text transport controls per condition, native and full context are 80/80, disabled is 40/80, and shuffled memory is 32/80. On a separate frozen 60-example routed-QA cohort, E0 selected text and E2 native K/V produce exactly the same output in all 840 pairs spanning cold, warm, multi-query, and concurrency-eight schedules while E2 removes 90.6–93.0% of visible prompt tokens. Cold end-to-end cost is slightly lower for E2, but warm, multi-query, and concurrent cost ratios are 1.136, 1.140, and 1.122. Semantic transport is closed; the unfused selected-cache path remains an economic optimization target. An expanded 149-question confirmation preserves 2,086/2,086 pairs and reduces the reused-path penalties to 5.4–7.4%. An expanded lifecycle study routes logical keys through the shared manager and built-in HiCache storage. WARM is 80/80 exact and restart recovery succeeds; int8 COLD is exact in only 14/80 sequences and remains opt-in. A five-example Llama-3.2-1B replication also preserves every lossless WARM output. Gemma-3-1B is not runnable on the pinned SGLang-MLX build because its per-layer sliding-window topology is unsupported before PRA is attached. An event-loop-owned replacement for Python-thread prefetch preserves 35/35 outputs and, with 250 ms lead, reduces median demand stall to 0.013 ms. The native streaming gateway then covers concurrency one through sixteen, cancellation, and request-owned cleanup; all 31 concurrent outputs are exact. TTFT p95 reaches 5.03 s at concurrency sixteen, exposing the serialized model-runner queue rather than hidden parallel execution. Finally, a two-Mac off-node WARM control transfers five exact 14.35 MB native objects over SSH-tunneled HTTP. Five independent transfers per lead show that 0–1,000 ms prefetch remains late, whereas 2,000 ms makes 40/40 concurrent demands ready. At concurrency sixteen, stable affinity reduces shared-resource traffic from 54.8 to 13.7 MiB and mixed-resource traffic from 260.1 to 68.4 MiB. This closes off-node placement as a controlled mechanism, not SGLang-native distributed E3.

1 Qualification Snapshot

The engine-specific question is whether semantic-resource identity remains useful beside SGLang’s prefix-addressed Radix cache and hierarchical HiCache. The native runner is validated at E2: it consumes frozen selected K/V outside Radix ownership and preserves all 2,086/2,086 matched E0 outputs. Visible prompt tokens fall by 90.4–93.0%. Macro E2/E0 cost ratios are 0.980 cold, 1.074 warm, 1.071 multi-query, and 1.054 concurrent. Local HiCache and a two-host WARM bridge demonstrate E3 mechanism candidates, but SGLang-owned distributed E3 remains unvalidated.

Integration level	E2 validated; local E3 mechanism observed, not promoted
Model/hardware	Pinned Qwen-family MLX runner on Apple M5
Workload	QASPER, HotpotQA, and 2WikiMultiHopQA; frozen selections
Quality	2,086/2,086 exact E0/E2 pairs; unchanged paired F1
Context reduction	90.4–93.0% fewer visible prompt tokens
TTFT/ITL tails	Online gateway measured separately; common matched tails NOT MEASURED
Throughput	Concurrent E2 inverse-cost ratio 1.054, or 94.9% of E0 rate
Peak memory	Native objects measured; complete engine decomposition NOT MEASURED
Reuse	Radix/PRA identities separate; local and two-host WARM reuse measured
Main limitation	Engine-native distributed HiCache, scheduler prefetch, and continuous batching

PRA primer. A PRA resource is a stable, authorized semantic object with selectable token intervals. E0 presents a frozen selection as ordinary visible text. E1 carries logical identity and deltas but still uses ordinary model execution. E2 feeds the same selection to native non-prefix attention with matched geometry, single normalization, request-local cleanup, and isolation. E3 adds scheduler ownership of physical tiers, prefetch, eviction, sharing, and batching. SGLang’s Radix prefix and PRA resource namespaces therefore coexist rather than alias.

Deployment recommendation. **Best validated deployment:** E0 selected text for ordinary repeated and concurrent serving; E2 when hidden native context or independent resource identity is required. **Use when:** semantic resources outlive or cut across sequential prefixes. **Do not use when:** the controlled HTTP bridge is being treated as SGLang’s native distributed HiCache provider. **Main remaining gate:** scheduler-owned distributed HiCache with continuous-batching tails and native placement telemetry.

2 Introduction

A prefix cache and a retrieval memory answer different questions. A radix tree asks whether token histories share an exact path. PRA asks which retained resource or interval is useful for a query. A long-running session can use both:

$$\underbrace{\text{stable active dialogue}}_{\text{radix prefix}} + \underbrace{\text{query-selected archive}}_{\text{PRA resource}}.$$

SGLang combines language-model serving with a radix cache and optional hierarchical KV mechanisms [1]. The systems question is not whether selected text can be sent to SGLang; a gateway already does that. It is whether a stable non-prefix PRA identity can share SGLang’s physical cache hierarchy and scheduling substrate without being misrepresented as a prefix.

This iteration contributes an engine-independent logical object lifecycle, a capability-honest SGLang adapter, a measured RadixAttention baseline, and a native runner-cache bridge measured through live prefill and batched decode. External placement, the built-in SGLang file-storage bridge, and the combined shared-prefix/native path are now measured locally; off-node distributed placement remains open.

3 Two Address Spaces

The two logical keys are

$$a_{\text{radix}} = (t_1, \dots, t_n), \quad a_{\text{PRA}} = (\tau, s, r, v, [i, j], \ell, m, d, p),$$

where the PRA key includes tenant τ , optional session s , resource and version (r, v) , source interval $[i, j)$, layer ℓ , model revision m , layout/dtype d , and positional/materialization policy p . Equality in one space does not imply equality in the other.

The black-box baseline selects exact evidence outside SGLang and serializes it into the current turn. RadixAttention may then cache that resulting token prefix. An engine-aware design would instead retain the PRA logical identity while placing its physical detail in GPU, host, or distributed hierarchy. Selected Context session deduplication is owned by the shared PRA runtime; engine-native prefix caching is measured independently. RadixAttention is therefore reported as prefix-tree physical reuse, not as logical PRA visibility.

4 Implementation Boundary

`src/prahf/engine_memory.py` provides validated states `INDEXEDONLY`, `OFFDEVICE`, `PREFETCHING`, `RESIDENT`, `PINNED`, `EVICTABLE`, and `INVALID`. Resident transitions require physical handles. Selection enforces tenant and authorization scope, and snapshots keep address, logical detail, and resident detail bytes disjoint.

`src/prasglang/adapters.py` exposes the OpenAI-compatible server as E0. It advertises automatic prefix caching and selected-text fallback, but no logical references or native K/V. Supplying a `SGLangNativeExecutor` promotes the adapter to E2 and transfers generation, streaming, session close, resource deltas, and block-store access to the executor. This pattern prevents a sidecar namespace from being mistaken for attention integration.

The hierarchical bridge in `src/prasglang/hicache.py` maps PRA identities to an independent HiCache-style namespace. It does not insert resources into the radix tree as artificial prefixes. Attention-ready MLX arrays form L1, host-owned arrays form L2, and compressed NPZ plus a logical-identity manifest forms L3. LRU pressure demotes L1 to L2 and L2 to L3; lookup promotes in the opposite direction before attention. On unified memory, L1 and L2 describe ownership and readiness rather than physically separate GPU and CPU DRAM.

`src/prasglang/hicache_backend.py` adds an immutable blob protocol over SGLang’s actual `HiCacheStorage` interface. A fixed header carries the payload length required by the backend’s preallocated-read contract; a hashed namespace key hides the resource URI. The same adapter can wrap the built-in file backend or a backend returned by SGLang’s factory. The portable interface does not expose per-key physical deletion, so PRA invalidation revokes logical access and relies on versioned identities plus backend eviction for byte reclamation. Raw framing is the default: an initial compressed transport added 2.8–3.0 seconds of CPU work and was rejected before final measurement.

`src/prasglang/mlx_native.py` implements the first attention bridge. Each layer cache exposes a local `offset` to scheduler and Radix bookkeeping, a distinct `rope_offset` to attention, and concatenated selected-plus-local K/V only inside the attention view. Prefix-pool synchronization sees local K/V arrays, so PRA detail cannot be mistaken for an exact prefix. Hooks cover initial prefill, request cleanup, and the runner’s batched-decode cache protocol. Release-time unwrapping is essential: the first live run showed that returning a selected-memory wrapper to SGLang’s reusable cache pool lets a later request inherit memory and can double-wrap a later native request. The corrected release hook returns only clean request-local K/V to the pool, and explicit geometry checks reject contaminated caches. The external HiCache supplies selected memory to this request-local bridge while remaining disjoint from ordinary reusable cache pools.

5 Environment

The host was an Apple M5 MacBook Pro with 16 GB unified memory, Metal 4, and macOS 26.4.1. SGLang revision `ef20fab38a03490e2cdf1b7377145ca3a3f2bfc5` ran under Python 3.11.15, MLX 0.32.2, and `mlx-lm` 0.31.3. The checkout required a local compatibility patch that restores the loader property expected by `MlxModelRunnerStub`; the patch is part of provenance, not hidden setup.

The model was Qwen/Qwen3-0.6B at revision `c1899de289a04d12100db370d81485cdf75e47ca`. The server used one running request, 2,048 total tokens, disabled CUDA-graph backends, and the MLX runner. Startup identified `UnifiedRadixCache` with LRU eviction. Built-in distributed HiCache was disabled for the HTTP control. A separate experiment uses SGLang’s built-in file backend through `StorageBackendFactory`; it tests the engine storage contract but not off-node transport. The control plane reported `torch.native`, while model inference was delegated to MLX; both facts are retained in the artifact.

The live native-runner experiments use the 4-bit `mlx-community/Qwen3-0.6B-4bit` checkpoint at revision `73e3e38d981303bc594367cd910ea6eb48349da8`. This separates them from the earlier unquantized HTTP smoke and prevents cross-precision latency comparison.

Table 1: SGLang MLX smoke results. TTFT is milliseconds. Warm cached tokens come from the scheduler log rather than the SSE usage object.

Condition	Exact	Prompt	Cold TTFT	Warm TTFT	Warm cached
None	0%	24	511.8	44.3	23
Prefix	0%	333	97.3	31.6	332
Selected	100%	56	40.8	31.5	55
Prefix + selected	100%	365	43.5	34.1	364
Full context	100%	1577	287.4	52.7	1576




Figure 1: Shared calibration view. Prompt sizes are identical across the Qwen tokenizers. Absolute engine latency is not comparable because SGLang used an unquantized model while vLLM and MLX-LM used a 4-bit checkpoint.

6 Method

The product-facing projection uses three fixed labels: *No PRA*, *PRA-No Adaptor*, and *PRA-Adaptor Bundle*. Within the mechanism study, selected-text E0 supplies the original-model condition and native E2/E3 supply PRA without learned adaptor behavior. An immutable bundle result is reported only when bundle resolution is present in the run receipt; absent bundle measurements remain `NOT_MEASURED`. Deltas are candidate minus No PRA and are not sign-inverted to make improvements appear positive. We used the same synthetic codeword task and serialized prompts as the vLLM and MLX-LM calibration. Twelve distractor records surround one authoritative record. Conditions are no prefix/no PRA, prefix only, selected evidence only, prefix plus selected evidence, and full context. Each is repeated five times. SSE timing provides TTFT; SGLang scheduler logs provide exact new and cached token counts. Tail percentiles are not reported from five samples.

This HTTP matrix is SMOKE evidence. Separate five-seed controlled experiments exercise native prefill, decode, request cleanup, concurrency 1-8, natural-text causal controls, and external L1/L2/L3 placement. They do not measure off-node cache affinity.

7 Results

The controls behave correctly: no-evidence conditions are 0/5, and every selected or full-context request recovers the answer. Prefix plus selected memory uses 365 tokens, 76.9% fewer than full context. Its cold TTFT is 43.5 ms versus 287.4 ms; its warm mean is 34.1 versus 52.7 ms.

The cache trace is the strongest mechanism observation. After the first request in each condition, the scheduler reports one new token and 364 cached tokens for prefix plus selected memory. Full context similarly uses one new and 1,576 cached tokens. Selected context therefore does not interfere with exact radix reuse. It also creates a much smaller cached sequence. These numbers measure token-prefix cache reuse, not PRA-resource reuse.

Table 2: Five-seed SGLang native-cache mechanism result. Radix separation is the fraction for which scheduler-local length excludes selected PRA tokens.

Seeds	Exact	Argmax parity	Radix separation	Latency ms
5	100%	100%	100%	222.0

8 Native Runner Results

The native path recovers all five answers, preserves next-token logits exactly, and keeps PRA tokens out of scheduler prefix length in every case. Mean completion latency is 222.0 ms. The SGLang MLX control plane reports no ordinary GPU KV pool because MLX owns inference state.

The next experiment installs the bridge on the real `MlxModelRunner`. It executes source ingestion separately, then runs selected memory through `prefill_start`, `prefill_finalize`, and batched decode. Full visible context and native memory recover 5/5 answers, while query-only disabled memory recovers 0/5. Every native request keeps all 141 source tokens outside the scheduler-local length.

Table 3: Five-seed live-runner concurrency. One immutable selected memory is shared by every request in a batch.

Concurrency	Exact	Requests/s	Shared MB saved	Scheduler isolation
1	100%	6.59	0.0	yes
2	100%	9.91	15.4	yes
4	100%	13.25	46.3	yes
8	100%	14.54	108.0	yes

Table 3 shows exact recovery at every tested concurrency. Mean throughput rises from 6.59 requests/s at one request to 14.54 at eight, with diminishing returns after four. Physical object sharing avoids 108.0 MiB of duplicate selected K/V at concurrency eight. These are controlled in-process measurements, not queued HTTP tail latency.

8.1 Natural-text controls

We also use balanced QASPER and HotpotQA text to build 40 examples per dataset and condition. Each source contains an inserted `AnswerCode` `yes/no` anchor among natural text. Constrained greedy decoding ranks the two answer tokens through SGLang’s own logit-edit path. This measures native transport and causal dependence, not unrestricted benchmark QA.

Table 4: SGLang live-runner natural-text transport controls.

Dataset	Condition	Samples	Ranked exact
qasper	disabled	40	50%
qasper	native	40	100%
qasper	ordinary_full	40	100%
qasper	shuffled	40	40%
hotpotqa	disabled	40	50%
hotpotqa	native	40	100%
hotpotqa	ordinary_full	40	100%
hotpotqa	shuffled	40	40%

Native and ordinary full context are 40/40 on each dataset. Disabled memory is 20/40 and shuffled memory is 16/40 on each. The negative controls establish that successful native decisions depend on the registered source rather than the query alone.

8.2 Matched selected text versus native K/V

We next freeze the routed selections for 20 original-answer examples from each of QASPER, HotpotQA, and 2WikiMultiHopQA. E0 writes the selected source once to ordinary Radix-owned prefix slots. E2 encodes exactly the same source tokens as immutable native K/V and leaves them outside Radix accounting. Both paths use the

same query suffix, greedy decoder, and 24-token output budget. The shared schedule includes cold one-shot, two warm repeats, three query formulations over the same resource, and an eight-request wave with independent prefill and batched decode.

Table 5: Matched-selection quality on the shared three-engine cohort. All parity pools 840 request pairs per engine across four schedules.

Engine	Dataset	Visible E0/E2	Reduction	Cold F1 E0/E2	Cold parity	All parity
mlx-lm	2wikimultihopqa	375/33	91.1%	0.067/0.067	100.0%	100.0%
mlx-lm	hotpotqa	415/39	90.6%	0.087/0.087	100.0%	100.0%
mlx-lm	qasper	399/28	93.0%	0.110/0.110	100.0%	100.0%
sglang-mlx	2wikimultihopqa	375/33	91.1%	0.074/0.074	100.0%	100.0%
sglang-mlx	hotpotqa	415/39	90.6%	0.084/0.084	100.0%	100.0%
sglang-mlx	qasper	399/28	93.0%	0.105/0.105	100.0%	100.0%
vllm-metal	2wikimultihopqa	378/33	91.2%	0.079/0.079	100.0%	100.0%
vllm-metal	hotpotqa	417/39	90.6%	0.074/0.074	100.0%	100.0%
vllm-metal	qasper	400/28	93.0%	0.101/0.101	100.0%	100.0%

Table 6: Macro E2/E0 execution-cost ratios; lower is better. Cold includes ingestion, concurrent is inverse requests/s, and E2 usable is milliseconds.

Engine	Cold	Warm	Multi-query	Concurrent	E2 usable ms
mlx-lm	1.130	1.155	1.148	1.132	51.2
sglang-mlx	0.976	1.136	1.140	1.122	55.0
vllm-metal	0.985	1.004	0.993	0.981	86.0




Figure 2: Macro representation cost after freezing retrieval. Lower is better; concurrent cost is inverse throughput. SGLang wins the ingestion-inclusive cold regime but pays unfused reuse and decode overhead.

SGLang preserves exact output and F1 on all 840 E0/E2 pairs. E2 reduces visible prompt tokens by 90.6–93.0%. Its macro cold end-to-end cost ratio is 0.976 after 55.0 ms mean time to usable context: native ingestion offsets the slower online path for a one-shot request. Warm and multi-query ratios rise to 1.136 and 1.140. At concurrency eight, E2 reaches 89.2% of E0 throughput, an inverse cost ratio of 1.122. The wrapper concatenates selected

and local arrays at each consumer layer and decode step, so native transport is semantically closed while fused materialization and decode economics remain open.

An expanded confirmation increases the cohort to 149 unique questions (49 QASPER, 50 HotpotQA, and 50 2Wiki). SGLang preserves all 2,086 paired outputs, with an engine-level Wilson 95% lower bound of 0.9982. Visible-token reduction is 90.4–93.0%. Macro cold, warm, multi-query, and concurrent E2/E0 cost ratios are 0.980, 1.074, 1.071, and 1.054. The larger cohort therefore retains the one-shot ingestion advantage and narrows, but does not remove, unfused reuse and decode overhead.

8.3 Semantic tiers versus HiCache tiers

The shared runtime now owns HOT/WARM/COLD/SOURCE as semantic service levels. For this paper, HOT maps to the attention-ready PRA object, WARM can map to a lossless host or HiCache L2 representation, and COLD can map to the built-in file backend or another validated L3 provider. This does not merge PRA identity with Radix identity. HiCache supplies physical storage and transport; the shared manager supplies typed-record priors, task/dependency retention, persistence grace, and deterministic quota eviction.

A common 20-run policy control (four engine labels, five seeds) preserves exact lossless tier round-trip and SOURCE reconstruction, retains open-task and dependent records in every run, and avoids all 20 immediate writes for a transient closure trace. Those are policy and byte-transport checks, not SGLang generation or distributed-HiCache measurements. The live runner, combined Radix/HiCache lifecycle, and matched E0/E2 rows below remain the engine evidence.

A new live bridge makes that ownership executable end to end. The shared manager uses an attention-ready SGLangPRAHiCache object for HOT and SGLang’s built-in HiCacheStorage file API for WARM. Request registration supplies only logical keys; the bridge promotes, pins, attaches exactly once, unwraps PRA detail before returning local caches to the ordinary pool, and unpins on cleanup. Three natural QASPER selections reproduce HOT output exactly after WARM promotion, and a fresh manager recovers all persisted objects after restart.

Table 7: Shared three-example live storage probe. Ratios include promotion and generation but exclude background tier transitions.

Engine	WARM exact	int8 exact	int8 first	Prefix	WARM/HOT	COLD/HOT
mlx-lm	100%	0%	67%	1.3	1.97	1.84
sglang-mlx	100%	0%	67%	0.7	1.90	1.89
vllm-metal	100%	0%	67%	4.0	1.96	2.00

For SGLang, median lifecycle latency is 279 ms HOT and 529 ms WARM (1.90×); background WARM placement costs 158 ms median. Int8 COLD promotion plus generation is 527 ms, while background conversion costs 701 ms. Int8 changes all three sequences and preserves the first token in two, so it remains an explicit diagnostic rather than a default HiCache policy. This closes lifecycle-manager control over the local built-in backend, not distributed off-node HiCache or concurrent tail curves.

An expanded 80-example matrix covers all three datasets at Qwen3-0.6B and QASPER at 1.7B. Lossless WARM is 80/80 exact and every manager recovers after restart; int8 COLD is exact in 14/80 sequences. WARM/HOT p50 ratios improve from 1.75 to 1.64 at 1.7B but remain 1.64–1.75 overall. This closes local lifecycle fidelity over a broader cohort, not distributed placement or online tail latency. A separate five-example Llama-3.2-1B QASPER cohort is 5/5 WARM exact with restart recovery. The equivalent Gemma-3-1B run is blocked before PRA attachment by the pinned backend’s lack of a per-layer sliding-window map; we preserve that failure as an explicit engine compatibility boundary.

Table 8: Expanded SGLang-MLX lifecycle cohorts. Latencies include synchronous promotion and generation.

Engine	Model	Data	n	WARM exact	int8 exact	HOT p50	WARM p50	WARM p95
sglang-mlx	Llama-3.2-1B-Instruct-4bit	qasper	5	100.0%	40.0%	242	287	302
sglang-mlx	Qwen3-0.6B-4bit	2wikimultihopqa	20	100.0%	15.0%	310	525	565
sglang-mlx	Qwen3-0.6B-4bit	hotpotqa	20	100.0%	20.0%	320	543	592
sglang-mlx	Qwen3-0.6B-4bit	qasper	20	100.0%	15.0%	307	538	604
sglang-mlx	Qwen3-1.7B-4bit	qasper	20	100.0%	20.0%	374	614	632

8.4 Scheduler-owned promotion and online gateway

The earlier prefetch probe used a Python thread and delayed the caller even when nominal lead exceeded promotion time. The revised path gives promotion to the event loop, deduplicates in-flight logical keys, and admits completed objects to a bounded hot set. All 35 generations remain HOT-exact. With no lead, median demand stall is 210.3 ms and demand-path cost is 1.561 times HOT. At 250 ms, four of five objects are ready, median stall is 0.013 ms, and the ratio is 1.035; by 350 ms all five are ready. This closes local overlap as a mechanism, while distributed prefetch policy remains open.

Table 9: Event-loop-owned WARM promotion. Lead and stall are milliseconds.

Engine	Lead	Ready	Stall p50	Stall p95	Demand/HOT
mlx	0 ms	0%	383.7	397.6	2.123
mlx	25 ms	0%	332.7	358.1	2.054
mlx	50 ms	0%	304.3	328.3	1.962
mlx	100 ms	0%	256.4	305.5	1.836
mlx	250 ms	0%	90.3	122.4	1.310
mlx	350 ms	40%	4.2	40.8	1.048
mlx	500 ms	100%	0.1	0.1	0.983
sglang	0 ms	0%	210.3	280.1	1.561
sglang	25 ms	0%	202.4	240.7	1.495
sglang	50 ms	0%	169.0	195.7	1.475
sglang	100 ms	0%	111.9	141.4	1.326
sglang	250 ms	80%	0.0	5.7	1.035
sglang	350 ms	100%	0.0	0.1	1.063
sglang	500 ms	100%	0.0	0.0	1.023

8.5 Two-host WARM prefetch and affinity

We next separate the model and WARM store across two Apple hosts. The model host uses the SGLang PRA HiCache adapter with Qwen3-0.6B; the storage host writes immutable objects atomically to disk. Because macOS local-network policy denied the benchmark Python binary direct LAN sockets, HTTP crosses an SSH local forward. We therefore label this transport *SSH-tunneled HTTP* and include its encryption, process, and forwarding overhead. It is a controlled off-node mechanism result, not a Mooncake, NIXL, or native SGLang HiCache result.

Each object contains 125 source tokens and occupies 14.35 MB. The lead-time curve uses eight simultaneous demands for one shared object and five fresh-cache transfers per lead. At 0 ms, median/p95 demand stall is 1,514/1,889 ms. Leads through 1,000 ms remain late in all 280 requests across those seven conditions. At 2,000 and 3,000 ms, all 80 demands are ready and p95 stall is below 0.003 ms. The non-monotonic short-lead values reflect five network transfers per point; the defensible conclusion is a 1–2 s crossover on this tunneled path, not a smooth latency law. Every restored tensor has zero maximum absolute error.

The placement experiment models four worker-local HOT sets over five rounds. Affinity hashes each resource identity to one worker; random placement samples a worker per request. At concurrency sixteen, affinity transfers one shared object (13.7 MiB) versus four under random placement (54.8 MiB). For five mixed resources it transfers 68.4 versus 260.1 MiB. Shared p95 demand stall falls from 5.27 to 1.56 s, and mixed p95 falls from 18.12 to 5.54 s. Aggregate arrival-wave throughput rises from 12.24 to 28.38 requests/s for shared resources and from 2.25 to 11.77 requests/s for mixed resources. These values include prefetch lead and the controlled transfer path, but not model generation or an SGLang distributed scheduler queue.

The native executor also runs through the OpenAI-compatible streaming gateway. A runtime correction normalizes raw native token rows into portable delta and completion events; before this correction the HTTP layer emitted metadata-only chunks and did not commit streamed sessions. The corrected run returns 24 content

Table 10: Two-host WARM placement at five repeated arrival waves. Stall is ms; HOT counts objects resident when prefetch is scheduled.

Concurrency	Workload	Policy	HOT/requests	p95 stall	Remote MiB	Requests/s
8	shared	affinity	32/40	982.8	13.7	17.83
8	shared	random	32/40	5465.4	54.8	5.94
8	mixed	affinity	32/40	6635.7	68.4	5.06
8	mixed	random	21/40	9348.5	260.1	1.54
16	shared	affinity	64/80	1558.4	13.7	28.38
16	shared	random	64/80	5274.2	54.8	12.24
16	mixed	affinity	64/80	5536.1	68.4	11.77
16	mixed	random	57/80	18115.1	260.1	2.25




Figure 3: Off-node WARM behavior. Left: five-transfer lead-time curve at concurrency eight. Center and right: stable affinity bounds remote traffic and tail stall for mixed-resource arrivals.

chunks per request with exact output in all 31 concurrency requests. Throughput remains 2.99–3.06 requests/s from concurrency one to sixteen; TTFT p95 rises from 35.6 to 5,026.6 ms under queued load. Cancellation after the first streamed token succeeds, and every exercised session releases request-owned PRA state. These measurements include HTTP handling and queueing but remain a single-host SGLang-MLX realization rather than distributed HiCache evidence.

Table 11: Online native-gateway concurrency, including queueing. Times are ms.

Engine	c	req/s	TTFT p50	TTFT p95	request p50	request p95
mlx-lm	1	1.70	302	302	589	589
mlx-lm	2	1.52	355	1028	642	1316
mlx-lm	4	1.66	950	2138	1233	2415
mlx-lm	8	1.58	2247	4779	2538	5066
mlx-lm	16	1.56	4838	9988	5123	10275
sglang-mlx	1	3.06	36	36	325	325
sglang-mlx	2	2.94	43	381	353	678
sglang-mlx	4	3.06	369	1007	670	1305
sglang-mlx	8	3.00	1037	2364	1345	2667
sglang-mlx	16	2.99	2337	5027	2619	5326

8.6 External HiCache placement

Table 12 resolves the same immutable selected K/V from each tier before live generation. All five outputs recover the target and all five keep selected tokens outside Radix prefix identity. Warm attention-ready lookup is 0.009 ms

Table 12: Five-seed external PRA HiCache result. L1/L2 on Apple unified memory denote attention readiness and host-array ownership; L3 is compressed disk backing.

Path or invariant	Mean latency ms	Success
Warm L1 hit	0.009	—
L2 $\rightarrow$ L1	4.0	—
L3 $\rightarrow$ L1	405.6	—
Answer recovery	—	100%
L1/L2 pressure demotion	—	100%
Radix namespace separation	—	100%

on average. Host-to-attention-ready promotion averages 4.0 ms. Compressed L3 restoration averages 405.6 ms and includes one 984 ms outlier, so it is a backing path rather than a latency-neutral cache hit. With one-object L1 and L2 budgets, inserting two additional resources forces one L1-to-L2 and one L2-to-L3 demotion in every seed. This closes local hierarchy mechanics; it does not establish distributed HiCache or server tail latency.

8.7 Built-in SGLang storage backend

The next control replaces the private NPZ L3 path with SGLang’s registered `file` backend while retaining the independent PRA logical namespace. For each seed, one adapter writes an immutable object and a fresh adapter, without process-local placement metadata, checks the backend and promotes it through L3 to attention-ready L1 before generation.

Table 13: Five-seed SGLang `HiCacheStorage` file-backend bridge. Times are milliseconds; the cold-read mean includes one 476.6 ms outlier.

Measure	Mean	Median	Result
raw object write	38.7	39.0	5/5 stored
cold backend read $\rightarrow$ L1	122.2	36.2	5/5 hits
warm L1 lookup	0.013	0.012	5/5 hits
generation after restore	—	—	5/5 exact

Each source has 125 tokens and 14.336 MB of native K/V; framing adds 17.2 KB. The result closes compatibility with SGLang’s real storage abstraction and fresh-adapter discovery. It does not establish Mooncake, NIXL, HF3FS, AIBrix, shared-memory, or cross-host transfer. Those backends can now be selected without changing PRA serialization or logical identity, but each requires its own measured correctness, transfer, and cleanup gate.

8.8 Prefetch overlap diagnostic

The storage bridge exposes a caller-owned prefetch hook so an engine scheduler can initiate promotion before a selected object is demanded. We evaluate an oracle signal, which removes selection uncertainty and gives the mechanism its most favorable timing case. Each condition contains five seeds and promotes a 14.336 MB, 125-token native-K/V object from the built-in file backend to attention-ready L1.

Table 14: Caller-owned Python-thread prefetch. “Actual lead” is elapsed time before the caller regains control and issues the demand. All 20 restored tensors are exact. Times are milliseconds.

Requested lead	Promotion	Actual lead	Ready at demand	Demand stall
0	27.9	0.4	0/5	27.6
10	27.6	203.3	5/5	0.094
25	25.0	192.6	5/5	0.088
50	26.3	235.3	5/5	0.109

The near-zero demand stalls at nonzero requested leads are not an optimization result. Python deserialization and MLX array construction delay the caller, so the nominal 10–50 ms scheduling window expands to 193–235 ms. The experiment closes exact asynchronous API semantics but leaves useful overlap open. A production path must use SGLang’s native asynchronous HiCache controller, a GIL-independent worker, or an equivalent engine-owned transfer path before prefetch latency claims are warranted.

8.9 Combined shared-prefix and external-memory lifecycle

Table 15: Five-seed A/B/C lifecycle through SGLang’s shared-prefix pool and external PRA HiCache. Ordinary B should not recover the selected fact.

Request	Recovery	Radix separated	One copy	Mean ms
A: selected after L2 promotion	5/5	5/5	5/5	196.7
B: ordinary after cleanup	0/5	5/5	–	127.8
C: reselected from warm L1	5/5	5/5	5/5	147.2

Table 15 drives the runner interface used after an exact Radix match. A warm request writes 81 stable-prefix tokens into SGLang’s shared K/V pool. Request A gathers those prefix slots while the PRA bridge promotes its separately named memory from L2 to attention-ready L1. Request B uses the same shared prefix after A is removed, but receives no PRA identity. Request C selects that identity again from warm L1. A and C recover 5/5; B recovers 0/5. Scheduler-local length excludes PRA in all 15 requests, and each selected cache reports the source token count once rather than twice.

This closes the local composition and cleanup mechanism. It also exercises capacity pressure across the five resources: three L1-to-L2 and six L2-to-L3 demotions occur while preserving recovery. The experiment directly controls the prefix-slot handoff rather than entering through SGLang’s HTTP scheduler; distributed HiCache, scheduler affinity, and concurrent WARM/COLD tier tails remain open. The local streaming-gateway queue is measured separately.




Figure 4: Live native-runner throughput and natural-text causal controls.

Gate	Status	Evidence
environment	closed	patched pinned SGLang-MLX starts reproducibly
radix baseline	smoke closed	exact scheduler cache-token trace
selected-text parity	smoke closed	selected and full are 5/5
HiCache non-prefix object	controlled closed	five-seed L1/L2/L3 promotion and forced demotion
built-in HiCache storage	controlled closed	file backend, fresh adapter 5/5, restored generation 5/5
shared storage lifecycle	controlled closed	logical-key promotion; WARM 85/85 exact across Qwen/Llama; restart recovery
prefetch overlap	negative/open	exact 20/20; 10–50 ms requests became 193–235 ms caller delays
event-loop promotion	controlled closed	35/35 exact; 250 ms lead reduces median demand stall to 0.013 ms
off-node WARM/affinity	controlled closed	two Macs; five transfer repeats/lead; zero tensor error; affinity bounds copies
selected native attention	controlled closed	5/5 exact, zero logit error, all 28 layers
radix/native identity	controlled closed	PRA tokens excluded from local prefix length
Radix + HiCache request	controlled closed	A/C 10/10, cleanup B 0/5, exactly one selected copy
runner lifecycle	controlled closed	prefill, batched decode, and clean request release
concurrency/sharing	controlled closed	five seeds at 1–8 requests; 108 MiB avoided
natural transport	controlled closed	native/full 80/80; causal controls near chance
matched natural QA	controlled closed	840/840 exact E0/E2 outputs; 90.6–93.0% fewer visible tokens
robustness	partial	one model; synthetic and two natural-text transport fixtures

9 Next Engine-Native Experiment

The matched comparison, local event-loop overlap, and controlled two-host WARM mechanism are complete. A fresh environment audit still finds no installed Mooncake, NIXL, HF3FS, or AIBrix backend on the Apple hosts. The built-in file backend remains local L3 evidence, while the new SSH-tunneled HTTP bridge proves that off-node identity, transfer, coalescing, and affinity are measurable. It does not prove integration with SGLang’s distributed HiCache scheduler.

The next experiment should replace the controlled bridge with one supported distributed backend and preserve the same frozen object/arrival schedule. It must measure backend queue time, network and device bytes separately, wasted prefetches, eviction/reload amplification, and p50/p95/p99 under continuous batching. The controlled result predicts that resource affinity matters: at concurrency sixteen, random placement creates 3.8–4.0 times as much remote traffic and 3.4 times the shared-resource p95 stall. A native scheduler result should either reproduce that direction or explain why backend sharing removes it.

9.1 Observability contract

PRA wrapper spans preserve W3C request identity across gateway mediation and the SGLang boundary. Native scheduler, Radix-cache, and HiCache metrics are scraped directly when exposed by the pinned runtime, while PRA reports normalized selected context, native attach, semantic tier movement, and session-affinity outcomes. Telemetry remains explicit and detailed tracing is off for primary qualification runs.

10 Falsification and Limitations

Deep SGLang integration is unnecessary if RadixAttention plus the tuned black-box path matches native PRA on quality-adjusted memory, TTFT, and throughput. The current baseline is already strong, so a metadata-only object registry cannot satisfy the claim.

The principal model measurements use one small model family and a local compatibility patch. The off-node control adds a second Apple host but uses an SSH-tunneled HTTP store rather than a supported distributed HiCache backend. The HTTP and native generation runs use different model precisions. The broad concurrency and natural-text runner cohorts disable Radix pooling; the dedicated combined cohort instead drives the exact

prefix-slot boundary directly. The later HTTP gateway exercises native streaming, cancellation, and cleanup, but serializes one local model runner. Engine-owned distributed HiCache, continuous-batching tails, production affinity routing, and archived-head transition remain unmeasured. L3 timing includes only five loads and one large outlier. The natural controls contain real dataset text but an inserted answer-code task. The original matched cohort uses only 20 routed selections per dataset and a 24-token greedy decode; the expanded confirmation improves sample size while retaining the bounded decoder. Its exact transport parity is stronger than its low absolute F1. The two-host control measures storage transfer and cache placement without generation; no distributed SGLang scheduling result is implied.

11 Related Work

SGLang combines a structured runtime with RadixAttention for shared prefix reuse [1]. Paged serving systems manage physical KV, while hierarchical caches extend residency across tiers. Retrieval systems select text or documents before generation. PRA differs by assigning stable identity to native non-prefix memory and making its active subset query dependent. The central integration constraint is to preserve both prefix and resource identities rather than collapse one into the other.

12 Conclusion

On the pinned SGLang-MLX environment, selected context and RadixAttention are clearly complementary: the selected condition preserves 5/5 recovery, removes 76.9% of visible prompt tokens, and reuses 364 of 365 tokens on warm requests. The native SGLang-cache path additionally gives 5/5 recovery and exact logit parity without adding PRA tokens to Radix prefix length. The live runner keeps that result through batched decode, scales to 14.54 requests/s at concurrency eight, and passes 80/80 natural-text transport decisions while disabled and shuffled controls fall near chance. SGLang attention can therefore consume selected non-prefix K/V while preserving the two address spaces. Paper 6.1 also establishes local L1/L2/L3 promotion and forced demotion without entering the Radix namespace and uses SGLang’s built-in file storage through a fresh adapter with 5/5 recovery. A final A/B/C cohort composes the shared-prefix pool with external HiCache, preserves 10/10 selected recoveries, prevents 5/5 cleanup leaks, and attaches exactly one memory copy. A five-seed prefetch diagnostic recovers all 20 tensors exactly but finds no useful Python-thread overlap: the caller is delayed by 193–235 ms. Event-loop promotion corrects that local mechanism, and the two-host control then shows a 1–2 s tunneled-WARM crossover and $3.8\text{--}4.0\times$ less remote traffic under stable affinity at concurrency sixteen. All off-node tensors remain exact. It leaves a supported distributed HiCache backend, full-scheduler integration, and pressure-scale server validation as the focused next experiments. On matched routed QA, it preserves all 840 E0 outputs while moving 90.6–93.0% of prompt tokens to native memory; the measured decode overhead makes fused selected-cache consumption the immediate systems optimization.

References

- [1] L. Zheng et al. SGLang: Efficient Execution of Structured Language Model Programs. 2024.