

PRA-SGLang: Query-Addressed Non-Prefix Memory beside RadixAttention

Jorge Simão
EInnovator R&D

August 2026

Abstract

SGLang’s RadixAttention reuses exact token prefixes. Progressive Retrieval Attention (PRA) selects versioned non-prefix resources and intervals according to the current query. We implement a logical PRA object namespace and an SGLang adapter and an MLX runner-cache bridge that keeps PRA memory outside Radix prefix accounting while consuming it as selected native K/V. A 25-request Apple-M5 smoke study with patched SGLang MLX separates exact radix reuse from selected-memory recovery. Prefix plus selected memory and full context both recover the fixed answer in 5/5 requests, while the no-evidence controls are 0/5. Selection uses 365 prompt tokens instead of 1,577. The scheduler reports 364 cached tokens on warm selected requests and 1,576 on warm full-context requests; warm TTFT is 34.1 and 52.7 ms, respectively. In a separate five-seed native-cache experiment, selected K/V recovers 5/5 answers and is bit-exact to SGLang’s ordinary split cache. All 28 attention layers consume selected memory under one normalization, while the scheduler-visible local offset excludes every PRA token. An external PRA HiCache then places immutable detail in attention-ready L1 arrays, host-owned L2 arrays, and compressed NPZ L3 backing without entering the Radix namespace. Across five seeds, answer recovery and Radix separation are 5/5; warm L1, L2-to-L1, and L3-to-L1 resolution average 0.009, 4.0, and 405.6 ms. Forced capacity pressure performs both L1-to-L2 and L2-to-L3 demotion in every seed. The bridge now also runs through SGLang’s real prefill and batched-decode lifecycle: native and full context recover 5/5, disabled memory recovers 0/5, and native throughput rises from 6.59 to 14.54 requests/s between concurrency one and eight. Shared immutable K/V avoids 108.0 MiB at concurrency eight. Across 40 QASPER and 40 HotpotQA natural-text transport controls per condition, native and full context are 80/80, disabled is 40/80, and shuffled memory is 32/80.

1 Introduction

A prefix cache and a retrieval memory answer different questions. A radix tree asks whether token histories share an exact path. PRA asks which retained resource or interval is useful for a query. A long-running session can use both:

$$\underbrace{\text{stable active dialogue}}_{\text{radix prefix}} + \underbrace{\text{query-selected archive}}_{\text{PRA resource}}.$$

SGLang combines language-model serving with a radix cache and optional hierarchical KV mechanisms [1]. The systems question is not whether selected text can be sent to SGLang; a gateway already does that. It is whether a stable non-prefix PRA identity can share SGLang’s physical cache hierarchy and scheduling substrate without being misrepresented as a prefix.

This iteration contributes an engine-independent logical object lifecycle, a capability-honest SGLang adapter, a measured RadixAttention baseline, and a native runner-cache bridge measured through live prefill and batched decode. HiCache placement remains a separate open boundary.

2 Two Address Spaces

The two logical keys are

$$a_{\text{radix}} = (t_1, \dots, t_n), \quad a_{\text{PRA}} = (\tau, s, r, v, [i, j], \ell, m, d, p),$$

where the PRA key includes tenant τ , optional session s , resource and version (r, v) , source interval $[i, j)$, layer ℓ , model revision m , layout/dtype d , and positional/materialization policy p . Equality in one space does not imply equality in the other.

The black-box baseline selects exact evidence outside SGLang and serializes it into the current turn. RadixAttention may then cache that resulting token prefix. An engine-aware design would instead retain the PRA logical identity while placing its physical detail in GPU, host, or distributed hierarchy.

3 Implementation Boundary

`src/prahf/engine_memory.py` provides validated states INDEXEDONLY, OFFDEVICE, PREFETCHING, RESIDENT, PINNED, EVICTABLE, and INVALID. Resident transitions require physical handles. Selection enforces tenant and authorization scope, and snapshots keep address, logical detail, and resident detail bytes disjoint.

`src/prasglang/adapters.py` exposes the OpenAI-compatible server as E0. It advertises automatic prefix caching and selected-text fallback, but no logical references or native K/V. Supplying a `SGLangNativeExecutor` promotes the adapter to E2 and transfers generation, streaming, session close, resource deltas, and block-store access to the executor. This pattern prevents a sidecar namespace from being mistaken for attention integration.

The hierarchical bridge in `src/prasglang/hicache.py` maps PRA identities to an independent HiCache-style namespace. It does not insert resources into the radix tree as artificial prefixes. Attention-ready MLX arrays form L1, host-owned arrays form L2, and compressed NPZ plus a logical-identity manifest forms L3. LRU pressure demotes L1 to L2 and L2 to L3; lookup promotes in the opposite direction before attention. On unified memory, L1 and L2 describe ownership and readiness rather than physically separate GPU and CPU DRAM.

`src/prasglang/mlx_native.py` implements the first attention bridge. Each layer cache exposes a local `offset` to scheduler and Radix bookkeeping, a distinct `rope_offset` to attention, and concatenated selected-plus-local K/V only inside the attention view. Prefix-pool synchronization sees local K/V arrays, so PRA detail cannot be mistaken for an exact prefix. Hooks cover initial prefill, request cleanup, and the runner’s batched-decode cache protocol. Release-time unwrapping is essential: the first live run showed that returning a selected-memory wrapper to SGLang’s reusable cache pool lets a later request inherit memory and can double-wrap a later native request. The corrected release hook returns only clean request-local K/V to the pool, and explicit geometry checks reject contaminated caches. The external HiCache supplies selected memory to this request-local bridge while remaining disjoint from ordinary reusable cache pools.

4 Environment

The host was an Apple M5 MacBook Pro with 16 GB unified memory, Metal 4, and macOS 26.4.1. SGLang revision `ef20fab38a03490e2cdf1b7377145ca3a3f2bfc5` ran under Python 3.11.15, MLX 0.32.2, and `mlx-lm` 0.31.3. The checkout required a local compatibility patch that restores the loader property expected by `MlxModelRunnerStub`; the patch is part of provenance, not hidden setup.

The model was Qwen/Qwen3-0.6B at revision `c1899de289a04d12100db370d81485cdf75e47ca`. The server used one running request, 2,048 total tokens, disabled CUDA-graph backends, and the MLX runner. Startup identified `UnifiedRadixCache` with LRU eviction. Built-in distributed HiCache was disabled for the HTTP control; the separate experiment uses the external PRA hierarchy described above. The control plane reported `torch_native`, while model inference was delegated to MLX; both facts are retained in the artifact.

The live native-runner experiments use the 4-bit `mlx-community/Qwen3-0.6B-4bit` checkpoint at revision `73e3e38d981303bc594367cd910ea6eb48349da8`. This separates them from the earlier unquantized HTTP smoke and prevents cross-precision latency comparison.

5 Method

We used the same synthetic codeword task and serialized prompts as the vLLM and MLX-LM calibration. Twelve distractor records surround one authoritative record. Conditions are no prefix/no PRA, prefix only, selected evidence only, prefix plus selected evidence, and full context. Each is repeated five times. SSE timing provides TTFT; SGLang scheduler logs provide exact new and cached token counts. Tail percentiles are not reported from five samples.

Table 1: SGLang MLX smoke results. TTFT is milliseconds. Warm cached tokens come from the scheduler log rather than the SSE usage object.

Condition	Exact	Prompt	Cold TTFT	Warm TTFT	Warm cached
None	0%	24	511.8	44.3	23
Prefix	0%	333	97.3	31.6	332
Selected	100%	56	40.8	31.5	55
Prefix + selected	100%	365	43.5	34.1	364
Full context	100%	1577	287.4	52.7	1576




Figure 1: Shared calibration view. Prompt sizes are identical across the Qwen tokenizers. Absolute engine latency is not comparable because SGLang used an unquantized model while vLLM and MLX-LM used a 4-bit checkpoint.

This HTTP matrix is SMOKE evidence. Separate five-seed controlled experiments exercise native prefill, decode, request cleanup, concurrency 1–8, natural-text causal controls, and external L1/L2/L3 placement. They do not measure distributed cache affinity.

6 Results

The controls behave correctly: no-evidence conditions are 0/5, and every selected or full-context request recovers the answer. Prefix plus selected memory uses 365 tokens, 76.9% fewer than full context. Its cold TTFT is 43.5 ms versus 287.4 ms; its warm mean is 34.1 versus 52.7 ms.

The cache trace is the strongest mechanism observation. After the first request in each condition, the scheduler reports one new token and 364 cached tokens for prefix plus selected memory. Full context similarly uses one new and 1,576 cached tokens. Selected context therefore does not interfere with exact radix reuse. It also creates a much smaller cached sequence. These numbers measure token-prefix cache reuse, not PRA-resource reuse.

7 Native Runner Results

The native path recovers all five answers, preserves next-token logits exactly, and keeps PRA tokens out of scheduler prefix length in every case. Mean completion latency is 222.0 ms. The SGLang MLX control plane reports no ordinary GPU KV pool because MLX owns inference state.

The next experiment installs the bridge on the real `MlxModelRunner`. It executes source ingestion separately, then runs selected memory through `prefill_start`, `prefill_finalize`, and batched decode. Full visible context and native memory recover 5/5 answers, while query-only disabled memory recovers 0/5. Every native request keeps all 141 source tokens outside the scheduler-local length.

Table 3 shows exact recovery at every tested concurrency. Mean throughput rises from 6.59 requests/s at one request to 14.54 at eight, with diminishing returns after four. Physical object sharing avoids 108.0 MiB of duplicate

Table 2: Five-seed SGLang native-cache mechanism result. Radix separation is the fraction for which scheduler-local length excludes selected PRA tokens.

Seeds	Exact	Argmax parity	Radix separation	Latency ms
5	100%	100%	100%	222.0

Table 3: Five-seed live-runner concurrency. One immutable selected memory is shared by every request in a batch.

Concurrency	Exact	Requests/s	Shared MB saved	Scheduler isolation
1	100%	6.59	0.0	yes
2	100%	9.91	15.4	yes
4	100%	13.25	46.3	yes
8	100%	14.54	108.0	yes

selected K/V at concurrency eight. These are controlled in-process measurements, not queued HTTP tail latency.

7.1 Natural-text controls

We also use balanced QASPER and HotpotQA text to build 40 examples per dataset and condition. Each source contains an inserted **AnswerCode** **yes/no** anchor among natural text. Constrained greedy decoding ranks the two answer tokens through SGLang’s own logit-edit path. This measures native transport and causal dependence, not unrestricted benchmark QA.

Table 4: SGLang live-runner natural-text transport controls.

Dataset	Condition	Samples	Ranked exact
qasper	disabled	40	50%
qasper	native	40	100%
qasper	ordinary_full	40	100%
qasper	shuffled	40	40%
hotpotqa	disabled	40	50%
hotpotqa	native	40	100%
hotpotqa	ordinary_full	40	100%
hotpotqa	shuffled	40	40%

Native and ordinary full context are 40/40 on each dataset. Disabled memory is 20/40 and shuffled memory is 16/40 on each. The negative controls establish that successful native decisions depend on the registered source rather than the query alone.

7.2 External HiCache placement

Table 5: Five-seed external PRA HiCache result. L1/L2 on Apple unified memory denote attention readiness and host-array ownership; L3 is compressed disk backing.

Path or invariant	Mean latency ms	Success
Warm L1 hit	0.009	–
L2 → L1	4.0	–
L3 → L1	405.6	–
Answer recovery	–	100%
L1/L2 pressure demotion	–	100%
Radix namespace separation	–	100%

Table 5 resolves the same immutable selected K/V from each tier before live generation. All five outputs recover the target and all five keep selected tokens outside Radix prefix identity. Warm attention-ready lookup is 0.009 ms

on average. Host-to-attention-ready promotion averages 4.0 ms. Compressed L3 restoration averages 405.6 ms and includes one 984 ms outlier, so it is a backing path rather than a latency-neutral cache hit. With one-object L1 and L2 budgets, inserting two additional resources forces one L1-to-L2 and one L2-to-L3 demotion in every seed. This closes local hierarchy mechanics; it does not establish distributed HiCache or server tail latency.




Figure 2: Live native-runner throughput and natural-text causal controls.

Gate	Status	Evidence
environment	closed	patched pinned SGLang-MLX starts reproducibly
radix baseline	smoke closed	exact scheduler cache-token trace
selected-text parity	smoke closed	selected and full are 5/5
HiCache non-prefix object	controlled closed	five-seed L1/L2/L3 promotion and forced demotion
selected native attention	controlled closed	5/5 exact, zero logit error, all 28 layers
radix/native identity	controlled closed	PRA tokens excluded from local prefix length
runner lifecycle	controlled closed	prefill, batched decode, and clean request release
concurrency/sharing	controlled closed	five seeds at 1–8 requests; 108 MiB avoided
natural transport	controlled closed	native/full 80/80; causal controls near chance
robustness	partial	one model; synthetic and two natural-text transport fixtures

8 Next Engine-Native Experiment

The next useful patch connects the verified external hierarchy to distributed HiCache placement and scheduler affinity. The matched comparison is RadixAttention plus black-box selected text versus Radix plus native PRA at equal quality and selected-token budget, including cold/warm tier location, prefetch overlap, server tail latency, and multi-request pressure. Only then should learned placement or cache-aware scheduling be optimized.

9 Falsification and Limitations

Deep SGLang integration is unnecessary if RadixAttention plus the tuned black-box path matches native PRA on quality-adjusted memory, TTFT, and throughput. The current baseline is already strong, so a metadata-only object registry cannot satisfy the claim.

The measurements use one small model family, a local compatibility patch, and one Apple host. The HTTP and native runs use different model precisions. The live native runner disables Radix pooling so this iteration does not yet measure prefix and native memory inside the same scheduled request. Distributed HiCache, HTTP-level native concurrency, tail latency, resource-affinity routing, and archived-head transition remain unmeasured. L3 timing includes only five loads and one large outlier. The natural controls contain real dataset text but an inserted answer-code task; they are not QASPER or HotpotQA end-task scores.

10 Related Work

SGLang combines a structured runtime with RadixAttention for shared prefix reuse [1]. Paged serving systems manage physical KV, while hierarchical caches extend residency across tiers. Retrieval systems select text or documents before generation. PRA differs by assigning stable identity to native non-prefix memory and making its active subset query dependent. The central integration constraint is to preserve both prefix and resource identities rather than collapse one into the other.

11 Conclusion

On the pinned SGLang-MLX environment, selected context and RadixAttention are clearly complementary: the selected condition preserves 5/5 recovery, removes 76.9% of visible prompt tokens, and reuses 364 of 365 tokens on warm requests. The native SGLang-cache path additionally gives 5/5 recovery and exact logit parity without adding PRA tokens to Radix prefix length. The live runner keeps that result through batched decode, scales to 14.54 requests/s at concurrency eight, and passes 80/80 natural-text transport decisions while disabled and shuffled controls fall near chance. SGLang attention can therefore consume selected non-prefix K/V while preserving the two address spaces. Paper 6.1 also establishes local L1/L2/L3 promotion and forced demotion without entering the Radix namespace. It leaves distributed HiCache, combined Radix-plus-native scheduling, and pressure-scale server validation as the focused next experiments.

References

- [1] L. Zheng et al. SGLang: Efficient Execution of Structured Language Model Programs. 2024.