

PRA-MLX: Query-Addressed Memory beside Prompt and Rotating KV Caches

Jorge Simão
EInnovator R&D

August 2026

Abstract

We ask whether Progressive Retrieval Attention (PRA) Native Memory can remain exact and economically useful as dense model size grows on Apple Silicon. Unified memory reduces host-accelerator transfer boundaries but does not make all retained context free to index, store, or attend to. We implement Selected Context (E0) and detached native K/V (E2) beside MLX-LM prompt and rotating caches, with request isolation, persistent HOT/WARM/COLD lifecycle, segmented attention, and explicit consumer-layer profiles.

On frozen QASPER, HotpotQA, and 2Wiki selections, E2 removes 90.6–93.0% of visible prompt tokens and preserves all 840 paired outputs; an expanded 149-question cohort preserves 2,086/2,086. Its cold, warm, multi-query, and concurrent cost ratios are 1.043, 1.139, 1.133, and 1.072, so unfused native consumption remains slower on reused paths. A synchronized Qwen3 ladder from 4B to 32B preserves all 120 selected-text sequences. Warm E2/E0 converges from 1.054 at 4B to 1.011 at 32B, while full segmented overhead falls from 1.165 to 1.047. No tested reduced consumer suffix passes the preregistered natural-QA quality gate; model scale improves cost convergence but does not justify a smaller correctness profile.

Storage pressure exposes an eight-resource working-set threshold: fitting all eight resources removes reloads and roughly halves resolution time, while smaller budgets preserve quality but reload every revisit. Lossless WARM storage and restart recovery are exact; full int8 COLD K/V is exact in only 13/80 expanded sequences and remains opt-in. A matched ten-question Qwen3-4B study adds BF16, MLX 8-bit, and MLX 4-bit deployed pipelines. PRA task score is preserved in BF16 and 8-bit but loses one or two answers at 4-bit; the cohort is too small and its selectors are not frozen independently, so it is not a precision-invariance result. A separate 50-question, selector-frozen Qwen3-4B run preserves every contiguous E0/E2 output at 2K and 4K; retained E2 reduces warm total latency by 79.0% and 86.9%, respectively, against selected-text re-prefill. A matched Qwen3-8B replication preserves every contiguous output and reduces repeated-selection latency by 82.1%, but larger generation does not repair the generic selector. Chunk-resident source-local K/V is not exact and remains a candidate policy. MLX provides a working native PRA mechanism and a clear model-scale trend. A fused segmented kernel, broader natural cohorts, and learned consumer gates remain the decisive steps for an economic and profile-sparsity claim.

1 Qualification Snapshot

The engine-specific question is whether selected native K/V can improve the quality–memory–latency frontier in Apple unified memory beyond selected-text E0. The live native executor is validated at E2 and preserves 2,086/2,086 matched E0 outputs while removing 90.4–93.0% of visible prompt tokens. The current unfused path remains slower: macro E2/E0 cost ratios are 1.043 cold, 1.139 warm, 1.133 multi-query, and 1.072 concurrent. E0 is therefore the default when visible selected text is acceptable; E2 is a functional native mechanism whose memory and non-disclosure properties require workload-specific valuation.

Integration level	E2 validated in-process; no E3 scheduler claim
Model/hardware	Qwen, Llama, and Gemma-family checks on Apple M5; Qwen3 1.7B–32B and 30B-A3B on M4 Pro
Workload	QASPER, HotpotQA, and 2WikiMultiHopQA; frozen selections
Quality	2,086/2,086 matched E0/E2 pairs; six additional oracle-transport cohorts preserve aggregate F1 exactly
Context reduction	90.4–93.0% fewer visible prompt tokens
TTFT/ITL tails	Online gateway measured separately; common matched tails NOT MEASURED
Throughput	Concurrent E2 inverse-cost ratio 1.072, or 93.3% of E0 rate
Model scaling	Warm concat E2/E0 approaches parity from 1.054 (4B) to 1.011 (32B); cold ratio rises from 0.932 to 0.976
Peak memory	Unified-memory components and controlled pressure measured; full live frontier open
Reuse	HOT/WARM lifecycle and restart recovery measured
Main limitation	Fuse segmented attention; no reduced consumer-layer profile passes the natural-QA smoke gate

PRA primer. PRA gives each durable semantic resource stable identity, version, authorization, and selectable intervals. E0 renders one frozen selection as visible text. E1 preserves logical identity while executing an ordinary model path. E2 consumes the identical selection as detached native K/V with correct positions, masks, cache topology, one normalization, cleanup, and isolation. E3 adds a scheduler that owns placement, promotion, eviction, sharing, and batching. MLX prompt or rotating caches retain sequential local state; PRA holds independently selected archive state.

Deployment recommendation. **Best validated deployment:** E0 selected text for latency-sensitive serving; E2 for bounded native residency, non-disclosure, or persistent resource reuse after profiling. **Use when:** a selected archive should remain outside the visible prompt. **Do not use when:** a large speed advantage is assumed from model scale alone. **Main remaining gate:** a fused segmented kernel plus a joint local-window/PRA-budget/consumer-layer frontier.

2 Introduction

MLX provides an efficient array framework and model runtime for Apple Silicon [1]. MLX-LM adds explicit prompt caches, rotating K/V, optional cache quantization, batching, and an OpenAI-like server. These tools make sequential context reuse practical in unified memory. PRA addresses a different object:

large retained memory $\rightarrow$ query-selected region $\rightarrow$ active native detail.

Selected Context session deduplication is owned by the shared PRA runtime; engine-native prefix caching is measured independently. MLX prompt, rotating, and split caches report physical sequential reuse separately from logical visibility and native PRA memory. Unified memory may reduce copy overhead, but it does not make every retained token free to index, keep, or attend to.

The companion runtime paper defines the shared deployment, evidence, and fallback contract [2]. This paper asks the narrower MLX question: does native selected K/V stay exact and become more economical as model size and unified-memory pressure increase?

Secondary experiments test rotating archives, persistence and sharing, consumer-layer sparsity, quantized residency, and the boundary between routing error and native-K/V consumption error.

3 Cache Taxonomy

We keep three mechanisms distinct:

Prompt cache

exact reusable prompt-prefix K/V, retained by MLX-LM.

Rotating cache

bounded recent sequential K/V, with old positions evicted as the prompt advances.

PRA memory

query-addressed non-prefix resources whose address state, native detail, and active selection have separate lifecycles.

Prompt-cache bytes are physical retained state. PRA backing may be much larger than active detail. Rotating capacity bounds direct context but does not create an address for evicted evidence.

4 Implementation

The engine-neutral namespace in `src/pra_hf/engine_memory.py` identifies a block by tenant/session scope, resource/version, token interval, consumer layer, immutable model revision, dtype/layout, materialization profile, and position policy. Its residency state machine records indexed-only, off-device, prefetching, resident, pinned, evictable, and invalid states. In MLX, the tier name denotes array residency or backing policy within unified memory; we do not mislabel it as CUDA host-to-device transfer.

`src/pra_mlx/adapter.py` exposes the HTTP server as E0. It advertises automatic prefix caching and text fallback, but no native K/V. An in-process `MLXNativeExecutor` is required before the adapter advertises E2, selected interval materialization, logical resource deltas, or explicit prompt cache handles. `src/pra_mlx/native.py` now provides that executor.

The design choice is important. Precomputing a prompt cache from selected text still makes that text a sequential prefix. A native executor must instead preserve the selected region’s identity and positional policy while combining its K/V with local K/V under correct model attention.

The implementation encodes immutable resource text once and retains post-RoPE K/V in $[B, H_{kv}, T, D]$ form for every layer. A request-local cache returns *selected memory plus local sequential K/V* to each attention module and constructs a mask in which every query sees selected memory while local keys remain causal. Query positions advance from the source extent; selected text never appears in the visible prompt. The residency manager deduplicates materialization, pins blocks for request lifetime, and evicts only unpinned detail under a bounded byte budget.

For compact residency, each K and V head receives a symmetric int8 scale over its token and head-width axes. Compact arrays remain the retained object; selection dequantizes them to the model-native dtype before constructing the request-local attention cache. This version measures storage economy, not an int8 attention kernel: active attention still consumes full-precision K/V.

Consumer-layer profiles preserve the source-relative query position even on layers that omit selected detail. Such a layer receives only request-local K/V, but its RoPE offset still advances from the source extent; otherwise a profile sweep would confound memory removal with a positional-frame change. Persisted arrays carry a strict sidecar fingerprint over model and tokenizer revisions, dtype, position policy, consumer profile, and resource version. A mismatch is rejected before arrays enter attention.

5 Environment

For cards and cross-engine qualification, we project measurements into *No PRA*, *Selected Context*, *Native Memory*, and *Native Serving*, with bundle use represented independently. E0 and E2 remain internal mechanism labels: selected-text E0 is PRA Selected Context and native E2 is Native Memory without learned adaptor behavior. The model-size and precision curves therefore report E2 minus E0, not PRA minus No PRA. A bundle condition requires a recorded immutable bundle revision and otherwise keeps an explicit missing state. Every delta preserves its sign and declares source, target, unit, aggregation, and direction. The host was an Apple M5 MacBook Pro with 10 CPU and 10 GPU cores, 16 GB unified memory, Metal 4, and macOS 26.4.1. The environment used Python 3.12.7, MLX 0.32.0, and mlx-lm 0.31.3. The model was `mlx-community/Qwen3-0.6B-4bit` at revision `73e3e38d981303bc594367cd910ea6eb48349da8`. The HTTP server retained up to 16 prompt-cache entries. It emits an explicit warning that its basic security checks are not production-grade; we treat it as an experimental deployment baseline.

The end-task study uses `mlx-community/Qwen3-1.7B-4bit`, MLX-LM 0.31.3, five seeds, and four held-out examples per seed and dataset. Evidence is oracle-scoped from dataset annotations and truncated to 384 tokens; routing quality is deliberately outside the first engine-transport experiment. A replication uses an M4 Pro with 20 GPU cores and 48 GiB unified memory, MLX 0.32.0, and MLX-LM 0.31.3. It repeats the same five seeds for Qwen3-1.7B and Qwen3-4B. Thus hardware and model scale change while dataset selection and execution

Table 1: MLX-LM server smoke. TTFT is milliseconds; warm cached is the mean over repeats two through five.

Condition	Exact	Prompt	Cold TTFT	Warm TTFT	Warm cached
None	0%	24	523.6	505.2	23
Prefix	0%	333	446.4	355.2	332
Selected	100%	56	415.9	355.0	55
Prefix + selected	100%	365	432.0	352.3	364
Full context	100%	1577	565.4	351.8	1576

Table 2: Five-seed controlled rotating-cache experiment. Cache MB is the sum of layer-cache `nbytes`; latency is the mean completion time.

Condition	K/V tokens	Exact	Cache MB	Latency ms
Full sequential K/V	full	100%	112.0	468.9
Rotating only	64	0%	7.0	706.6
Rotating only	128	0%	14.0	730.2
Rotating only	256	0%	28.0	722.1
Rotating only	512	0%	56.0	721.2
Rotating + selected archive	64	40%	7.0	704.7
Rotating + selected archive	128	0%	14.0	736.2
Rotating + selected archive	256	0%	28.0	713.1
Rotating + selected archive	512	0%	56.0	732.2

conditions remain fixed. A second matched-budget cohort retains every HotpotQA and 2Wiki candidate document and every QASPER paper paragraph, then routes four candidates before native materialization.

6 Server Prompt-Cache Experiment

The shared fixed task contains a stable 309-token prefix, one authoritative codeword record, twelve distractor records, and a short query. Five conditions separate no memory, prefix only, selected memory only, prefix plus selected memory, and full context. Each is repeated five times. The server reports cached tokens in the streaming usage object, and SSE events determine TTFT.

Table 1 shows correct controls: selected and full conditions are 5/5, while conditions without evidence are 0/5. Prefix plus selected memory uses 365 tokens and full context uses 1,577, again a 76.9% visible reduction. Warm selected requests reuse 364 tokens; warm full requests reuse 1,576. The server’s nearest-prefix cache therefore composes cleanly with the selected-text baseline.

Warm TTFT is nearly flat, 352.3 versus 351.8 ms, because exact prompt caches remove most repeated prefix work in both conditions. This is a useful negative latency result: after a perfect repeat cache hit, reducing visible input does not improve warm TTFT on this setup. Physical retention still differs. The log shows five distinct prompt-cache sequences occupying about 0.28 GB at the end of the matrix.

7 Rotating Cache and Selected Archive

The Python-API control uses five deterministic filler permutations. One authoritative four-digit fact appears near the beginning of a roughly thousand-token prompt. We compare full sequential K/V with rotating capacities 64, 128, 256, and 512. For each rotating capacity, a second condition reintroduces the authoritative record immediately before the final question. The answer metric accepts harmless formatting but requires the correct digits.

Full sequential K/V is 5/5 with 112 MB of cache. Rotating-only is 0/5 at every capacity while reducing layer-cache storage from 112 MB to 7, 14, 28, and 56 MB. Selected archive text reaches 2/5 at 64 and 0/5 at larger capacities. Mean completion latency is 469 ms for full K/V and approximately 705–736 ms for bounded conditions. The generated outputs show chat-format and answer-path instability, not merely loss of the early fact.

This non-monotonic result should not be read as evidence that smaller cache is better. It indicates that the direct rotating-cache intervention changes more than evidence availability for this model and prompt. The selected-text fallback does not restore stable computation. A native PRA executor must retain local sequential semantics while adding selected memory separately; the next section evaluates that mechanism directly.




Figure 1: A bounded rotating cache is not a reliable archive in this control. Selected late text recovers two of five cases only at the smallest capacity.

Table 3: Five-seed native selected-K/V control. Latency excludes the separately reported one-time native encoding cost. Native MB is all-layer selected K/V.

Condition	Exact	Latency ms	Native MB	Max error
Ordinary split cache	100%	187.0	0.0	0
Native selected K/V	100%	166.3	15.4	0
Native + rotating local 64	100%	215.5	15.4	0

Peak process memory remains about 1.01–1.07 GB across these conditions despite the 7–112 MB cache range. Model and runtime allocations dominate on this small checkpoint, so layer-cache bytes are the interpretable economy metric. Larger models and memory pressure are required before process-level savings claims.

8 Native Selected-K/V Experiment

The source contains 141 tokens. Ordinary split-cache and native selected-K/V produce identical logits, identical argmax tokens, and 5/5 exact recovery. Native selected execution averages 166.3 ms after a separately measured 55.4 ms one-time encoding step and retains 15.4 MB of all-layer K/V. Most importantly, replacing the local sequential cache by a 64-token rotating cache still recovers 5/5. The selected archive is therefore independent of local rotation rather than late visible text competing inside that bounded cache.

Table 4 repeats the mechanism on Llama 3.2 1B and Gemma 3 1B. Llama achieves 5/5 native and rotating-local recovery with exact split-cache logits while retaining 4.34 MB. Gemma also has zero logit error and matching argmax decisions, but both its ordinary and native conditions are 0/5 on this particular answer-format task. Gemma therefore establishes mechanism parity, not answer recovery. Together the runs cover three attention topologies and tokenizers; they are not yet natural-workload results.

Table 4: Cross-model native selected-K/V replication. Native and rotating columns report exact recovery over five seeds; error is zero for every row.

Model	Seeds	Native exact	Rotating exact	Native MB
Qwen3 0.6B 4-bit	5	100%	100%	15.42
Llama 3.2 1B 4-bit	5	100%	100%	4.34
Gemma 3 1B 4-bit	5	0%	0%	3.58

9 Profiles, Persistence, and Concurrency

Table 5: Five-seed consumer-layer sweep. Active MB counts selected native K/V, not model weights or request-local K/V.

Consumer profile	Layers	Exact	Active MB	Max error
all	28	100%	15.42	0
disabled	0	0%	0.00	6.188
last_4	4	0%	2.20	4.844
last_half	14	100%	7.71	4.031
last_quarter	7	0%	3.86	4.188

Table 5 exposes a sharp controlled boundary. Full native memory and the last-half profile both recover 5/5 answers. Last quarter, last four, and disabled profiles recover 0/5. Last half reduces active selected K/V from 15.42 to 7.71 MiB. Its nonzero logit difference is expected because fourteen layers no longer see the archive; exact recovery, not bit parity, is the profile criterion. This is a candidate economy profile for this fixture, not a production profile calibrated from five examples.

Table 6: Five-seed persistence and concurrent request sharing. Persist load is the mean for a 15.43 MiB selected-K/V sidecar.

Concurrency	Exact	Requests/s	Mean latency ms	Shared MB saved	Persist load ms
1	100%	6.04	164.9	0.0	7.2
2	100%	10.80	180.1	15.4	7.2
4	100%	16.98	222.4	46.3	7.2
8	100%	16.01	383.9	108.0	7.2

Persisting and reloading native arrays yields zero maximum logit error. Mean save and load times are 28.7 and 7.2 ms. Separate request-local caches share one immutable memory object: at concurrency eight this avoids 108.0 MiB of duplicate selected K/V. Aggregate throughput increases through concurrency four and then plateaus, while mean per-request latency rises from 164.9 to 383.9 ms. These are in-process controlled measurements, not queued server tail latency.




Figure 2: Left: the controlled consumer-layer quality–memory boundary. Right: shared immutable K/V raises throughput until the small Apple host saturates.

10 Larger-Model Segmented-Attention Calibration

The 48 GiB M4 Pro permits a single-process Q4 ladder that was unavailable to the original small-host study. We load pinned Qwen3-8B, 14B, and 32B checkpoints and Qwen3-30B-A3B. Each model evaluates five unique annotated evidence examples from each of QASPER, HotpotQA, and 2Wiki. The selected source is bounded at 384 tokens and held fixed across selected-text E0, concatenated native E2, live segmented native E2, and six contiguous consumer profiles. Retrieval is absent from this experiment. It asks whether MLX transports the same selected evidence and how sensitivity to consumer depth scales.

Table 7: M4 Pro larger-model calibration over 15 natural-QA examples per model. Agreement is greedy sequence agreement with selected-text E0; $|\Delta|LP$ is the mean across dataset-level absolute gold-log-probability deltas. PRA MiB is active all-layer selected K/V, excluding weights and local K/V.

Model	Layers	E0 F1	Concat agr.	Seg. agr.	$ \Delta LP$	PRA MiB	Seg./E0
Qwen3-8B	36	0.236	1.000	1.000	0.189	39.6	1.08
Qwen3-14B	40	0.277	1.000	1.000	0.027	43.9	1.05
Qwen3-32B	64	0.231	1.000	0.867	0.039	70.3	1.03
Qwen3-30B-A3B	48	0.256	1.000	0.933	0.146	26.4	1.12




Figure 3: Left: likelihood sensitivity as contiguous late consumers are removed. Dense 32B is less sensitive in likelihood, but reduced profiles still change generated sequences. Right: selected text, concatenated native K/V, and the unfused segmented implementation on the same selected evidence.

Concatenated E2 is the numerical reference path: all 60 model-example pairs match E0 sequences. The segmented path performs the same mathematical normalization with separate memory and local score blocks, accumulating the shared denominator and value numerator in FP32. It matches every 8B and 14B sequence and 27/30 across dense 32B and MoE. The remaining changes are tie-sensitive numerical differences after 48–64 repeated attention layers, not evidence-selection differences; mean absolute answer-likelihood drift stays at or below .189 nats. The current Python/MLX composition is also not a speed optimization: it costs 1.08, 1.05, 1.03, and 1.12 times E0 completion latency. Its systems result is removal of the temporary K/V concatenation, not a fused-kernel throughput claim.

The placement result is more conditional. Across the dense ladder, mean absolute likelihood loss for the last-three-quarter profile falls from 2.545 to 1.665 to .356 nats. This supports the hypothesis that larger dense models can be less likelihood-sensitive to omitted early consumers. Yet sequence agreement for that profile is only .267, .467, and .133, respectively. The 30B-A3B point is fast and uses 26.4 MiB of active all-layer K/V versus 70.3 MiB for dense 32B, but its last-three-quarter loss is 2.662 nats. Parameter sparsity therefore does not imply semantic-context sparsity. All reduced-layer rows remain CALIBRATION PENDING; the correctness profile remains all eligible layers.

10.1 Matched lifecycle scaling law

The preceding profile experiment synchronized ordinary source encoding, but an earlier 1.7B/4B runner did not. MLX executes lazily: that runner called the ordinary source forward outside its timer without calling `mx.eval`, so the work completed inside generation. Native encoding did call `mx.eval`. Its apparent 44.5–54.5% E2/E0 completion ratio therefore mixed cold E0 with warm E2 and is retired.

The corrected runner makes lifecycle state part of the metric. *Cold* cost is synchronized representation encoding plus query generation; *warm* cost is generation from synchronized reusable ordinary or native state. It evaluates two independently shuffled examples per dataset under each of five seeds, or 30 examples per model, at 384 selected source tokens and 12 greedy output tokens. The 4B anchor and pinned 8B, 14B, and 32B models each run ordinary warm E0, concatenated warm E2, and segmented E2 at all, 7/8, 3/4, 2/3, and 1/2 of decoder layers. This produces 840 model-backed condition rows; retrieval and selected evidence are fixed within each pair.

Table 8: Matched lifecycle scaling on the M4 Pro. Ratios divide E2 by E0; lower is faster. Concat is the numerical reference. Segmented agreement is exact greedy-sequence agreement. “None” means no reduced profile passed the smoke quality gate.

Model	n	concat warm	concat cold	seg. warm	seg. agr.	balanced	PRA MiB
Qwen3-4B	30	1.054	0.932	1.165	0.833	none	52.6
Qwen3-8B	30	1.031	0.932	1.102	0.767	none	52.6
Qwen3-14B	30	1.018	0.950	1.055	0.933	none	58.5
Qwen3-32B	30	1.011	0.976	1.047	0.833	none	93.6




Figure 4: Left: matched cold and warm E2/E0 costs. Warm overhead approaches parity with model size, but cold E2’s advantage contracts. Right: paired task-F1 change as consumer layers are removed; the dotted line is the -0.03 smoke margin, not a production noninferiority bound.

Concatenated E2 preserves every one of the 120 E0 sequences and has zero gold-answer-likelihood drift. Its warm ratio falls monotonically from 1.054 at 4B to 1.011 at 32B. Paired bootstrap 95% intervals are $[1.052, 1.056]$ and $[1.009, 1.013]$ at the endpoints. This four-point trend is convergence toward parity rather than a native execution advantage. The cold ratio moves in the opposite direction, from 0.932 at 4B and 8B to 0.950 at 14B and 0.976 at 32B. Reference encoding and ordinary source-cache construction scale similarly, so larger models do not amplify the measured cold benefit.

The unfused segmented path also improves with scale: all-layer warm overhead falls from 1.165 to 1.047. Its exact agreement is 0.833, 0.767, 0.933, and 0.833, while mean absolute gold-log-probability drift remains 0.090–0.214 nats. These are small reduction-order effects accumulated across many layers; the concatenated reference shows that source K/V itself is exact.

Table 9: Consumer-depth quality and cost. $|\Delta|LP$ is mean absolute gold-answer-log-probability change from E0. Agreement is exact sequence agreement. Reduced profiles save active PRA K/V proportionally but none passes both the paired F1 and likelihood gate.

Model	profile	fraction	$\Delta F1$	$ \Delta LP$	agreement	warm ratio
Qwen3-4B	all	1.000	-0.011	0.214	0.833	1.165
Qwen3-4B	last 7/8	0.889	-0.037	2.131	0.300	1.177
Qwen3-4B	last 3/4	0.750	-0.010	3.412	0.233	1.155
Qwen3-4B	last 2/3	0.667	-0.040	7.029	0.167	1.140
Qwen3-4B	last 1/2	0.500	-0.085	14.086	0.167	1.115
Qwen3-8B	all	1.000	+0.018	0.167	0.767	1.102
Qwen3-8B	last 7/8	0.889	-0.037	1.758	0.267	1.092
Qwen3-8B	last 3/4	0.750	-0.052	3.825	0.200	1.078
Qwen3-8B	last 2/3	0.667	-0.078	7.019	0.233	1.070
Qwen3-8B	last 1/2	0.500	-0.135	12.587	0.133	1.053
Qwen3-14B	all	1.000	-0.012	0.114	0.933	1.055
Qwen3-14B	last 7/8	0.875	-0.026	1.298	0.333	1.053
Qwen3-14B	last 3/4	0.750	-0.090	2.148	0.333	1.043
Qwen3-14B	last 2/3	0.675	-0.098	2.719	0.333	1.035
Qwen3-14B	last 1/2	0.500	-0.186	11.460	0.133	1.023
Qwen3-32B	all	1.000	-0.001	0.090	0.833	1.047
Qwen3-32B	last 7/8	0.875	-0.030	1.009	0.167	1.032
Qwen3-32B	last 3/4	0.750	-0.111	1.495	0.167	1.025
Qwen3-32B	last 2/3	0.672	-0.058	1.791	0.100	1.020
Qwen3-32B	last 1/2	0.500	-0.078	2.173	0.133	1.011

The consumer-layer hypothesis is not supported under this evidence geometry. At 32B, the last-7/8 profile loses 0.030 mean F1, has a wide paired bootstrap interval $[-0.130, 0.058]$, and drifts by 1.009 nats in absolute answer likelihood. The last-half profile reaches near-E0 warm cost (1.011) but loses 0.078 F1 and drifts by 2.173 nats. Similar or larger likelihood losses occur at smaller sizes. The minimum supported consumer fraction therefore remains one. The SDK’s `BALANCED` profile consequently consumes native memory at all eligible layers. Reduced profiles remain `CALIBRATION PENDING`; parameter scale alone does not promote them.

Public exact-identity qualification. The release pipeline imports these measurements into model cards only after matching the model ID, immutable revision, 4-bit MLX quantization, `mlx-lm` version, `BALANCED` profile, and Native Memory execution mode. The public Qwen3-8B, 14B, and 32B cards use the concatenated all-consumer path, not the numerically drifting segmented variant. Across the 15 paired QASPER, HotpotQA, and 2Wiki examples at each size, token F1 is unchanged and all outputs agree exactly; visible prompt tokens fall by 89.1%. Mean completion overhead contracts from 3.0% to 1.1% and 0.5% with scale. The 32B median TTFT point estimate is 6.6% lower, but the cohort is too small for a tail- latency claim. We therefore label these rows `ENGINE QUALIFIED`, retain all eligible consumer layers in `BALANCED`, and leave `QUALITY` and `ECONOMY` as `CALIBRATION PENDING`.

10.2 M5 timing replication and held-out layer gate

We repeated the corrected runner on a 16 GiB M5 after one unmeasured graph warm-up per execution path. Each measured checkpoint uses two examples from each dataset under five seeds, or 30 paired examples per model. Qwen3-4B, 8B, and 14B complete; the 32B preflight requires 22.86 GB against 17.18 GB physical memory and is recorded as `NOT_RUN_CAPACITY_GATE`.

Table 10: Corrected M5 replication. Ratios compare matched warm or cold states; concat preserves every E0 output. No reduced profile passes the quality gate.

Model	n	concat warm	concat cold	seg. warm	seg. agr.	balanced	PRA MiB
Qwen3-4B	30	1.035	0.973	1.290	0.933	none	52.6
Qwen3-8B	30	1.015	0.968	1.156	0.867	none	52.6
Qwen3-14B	30	0.980	0.946	1.031	0.933	none	58.5

Warm concat ratios fall from 1.035 at 4B to 1.015 at 8B and 0.980 at 14B. The 14B paired-bootstrap interval, $[0.922, 1.012]$, includes parity, so the point estimate is promising rather than a confirmed native speedup. Cold usable-context ratios are 0.973, 0.968, and 0.946. Unfused segmented overhead also contracts from 1.290 to 1.156 and 1.031; agreement remains 0.933, 0.867, and 0.933. The M5 result therefore strengthens model-scale convergence and the exact concat reference. It does not justify a fused-kernel claim or a reduced consumer profile.

We also calibrated a parameter-free noncontiguous layer selector on Qwen3-4B. Seed 11 greedily tests eight contiguous layer groups under a 0.05-nat absolute-likelihood bound and 0.90 first-token agreement, then seeds 23/37/53/71 form the held-out cohort. Every candidate removal fails on the calibration split; even dropping the first five layers changes likelihood by 1.538 nats. The selected set is therefore all 36 layers.

Table 11: Held-out noncontiguous consumer-layer calibration. The calibrated gate is learned from the calibration split but resolves to all layers.

Condition	n	Layers	F1	Sequence agreement	Completion ms
Selected text E0	12	0/36	0.224	1.000	356.1
Segmented E2, all	12	36/36	0.224	1.000	469.3
Calibrated gate	12	36/36	0.224	1.000	470.5

The 12 held-out examples preserve every selected-text sequence, but because the calibrated action equals the all-layer action this is a negative sparsity result. A trained, query-conditioned gate remains open work. Fixed suffixes and calibration-only group elimination have both failed.

10.3 Cross-family replication at 8B

The concat transport result is not Qwen-specific. On a 48 GiB M4 Pro we run Llama 3.1-8B-Instruct-4bit on 60 natural-QA examples: 20 each from QASPER, HotpotQA, and 2Wiki, with five seeds. The selector output, source cap, and 12-token greedy decode are frozen between E0 and E2. Native E2 preserves all 60 output sequences, aggregate token F1 (0.3563), and gold-answer log probability (-6.0414). Warm request cost is 349.4 versus 338.7 ms, a 1.032 ratio; cold time to usable context is 1,076.1 versus 1,151.2 ms, a 0.935 ratio. Mean active native detail is 46.2 MiB. A current-code 45-example overlapping rerun reproduces exact output and likelihood parity with warm/cold ratios 1.032/0.939. This closes useful-scale transport replication for a second model family, but not reduced-layer calibration: the segmented consumer patch remains Qwen-specific.

Table 12: Llama 3.1 8B cross-family transport replication over 60 natural-QA examples. Times are means in milliseconds.

Mode	F1	Gold log-prob.	Warm	Cold usable	Agreement
E0 selected text	0.3563	-6.0414	338.7	1,151.2	60/60
E2 native concat	0.3563	-6.0414	349.4	1,076.1	60/60

11 Long-Context Segmented Execution

The preceding ladder fixes selected evidence at a short context. We next keep the same 384-token selection and add deterministic same-dataset distractors to 8K source tokens for Qwen3-8B, 14B, and 32B. Five unique examples per dataset give 15 independent questions per model. An 8B boundary run adds one example per dataset at 32K. The requested 64K rows are emitted as NOT_RUN_MODEL_CONTEXT_LIMIT: the pinned models expose a 40,960 position limit, and the harness does not silently truncate.

Four representations separate execution questions. FULL VISIBLE is the ordinary prefix. E0 SELECTED renders the frozen evidence as text. E2 SELECTED consumes that selection through segmented native K/V. E2 SOURCE RELATIVE consumes the full dense source through the same segmented path; QUERY RESTART reuses its K/V but incorrectly starts the query at position zero.

Table 13: Long-context MLX calibration. ΔLP is full visible minus selected text. Selected E2 agreement compares the two frozen selected-evidence representations. Source/restart columns compare full native K/V with the ordinary full prefix.

Model	Context	ΔLP full-selected	selected E2 agr.	source RMSE	restart RMSE	source agr.	restart agr.
Qwen3-8B	8K	+1.192	1.000	0.112	2.572	0.667	0.000
Qwen3-8B	32K	-0.603	1.000	0.128	2.530	0.333	0.000
Qwen3-14B	8K	-1.485	1.000	0.100	4.786	0.933	0.000
Qwen3-32B	8K	-0.999	0.867	0.121	3.671	0.667	0.000

Selected E2 preserves every selected-text sequence at 8B and 14B. At 32B it preserves 13/15 sequences; mean gold-answer log probability differs by only +0.030 nats and mean F1 by +0.017. The full-source segmented path is deliberately held to a stricter ordinary-prefix comparison: first-logit RMSE is 0.112 and 0.100 at 8K, whereas query restart gives 2.572 and 4.786. Source-relative sequence agreement is 0.667 and 0.933; restart agreement is zero. The segmented implementation uses one mathematical normalization but separate FP32 reductions, so repeated unfused numerical differences accumulate through depth. This is distinct from the concatenated-cache reference above, which preserves all 60 sequences.

The pressure result illustrates why full-source native K/V is not the product goal. At 8B, selected detail occupies 39.6 MiB at both 8K and 32K, while all-layer full-source K/V grows from 1,152 to 4,608 MiB. MLX allocator peak for the ordinary full condition rises from 9.10 to 22.55 GiB. At 14B and 8K, selected detail is 43.9 MiB versus 1,280 MiB full-source K/V. At 32B the corresponding pair is 70.3 versus 2,048 MiB, while ordinary full-prefix peak allocation reaches 25.62 GiB. The three-case 8B 32K boundary is useful for mechanism and pressure, not tail-latency inference.

Corrected occupied-context curve. An independent Qwen3-8B rerun uses six natural-QA examples at each of 2K, 8K, and 32K occupied source tokens. Selection remains fixed at the same bounded evidence, and each execution path is warmed before measurement.

Table 14: Corrected occupied-context pressure. Selected detail stays fixed as the distractor-bearing visible context grows. Times are in-process completion latencies and do not estimate queued tails.

Context	n	E0 selected ms	E2 selected ms	E2/E0	Selected MiB	Full peak GiB
2K	6	350.9	387.4	1.104	39.6	6.37
8K	6	350.9	387.4	1.104	39.6	11.44
32K	6	350.9	387.4	1.104	39.6	31.85

Selected E2 agrees with selected-text E0 on all 18 sequences and keeps active detail at 39.6 MiB. Its mean completion time is 387.4 ms at every occupancy, 1.104 times the 350.9 ms E0 selected path. In contrast, full-visible peak allocation grows from 6.37 to 31.85 GiB and completion time from 387.1 to 1,116.0 ms. Full-source native K/V grows from 288 to 4,608 MiB and is not the intended product mode. Query restart fails on every 8K and 32K sequence, confirming that source-relative position continuity is a correctness requirement rather than optional metadata.

12 Natural-Text Causal Controls

The QASPER and HotpotQA builders draw natural paper or article text and place a balanced **AnswerCode** **yes/no** anchor among 63 source units. They test whether an answer-bearing natural-text resource survives the native transport; they do not claim unrestricted question answering. Each dataset contributes eight examples under each of five seeds, for 40 examples per condition. We rank the next-token logits for **yes** and **no**; free-form generation is retained only as a diagnostic because this small completion model often ignores the requested answer format.

Ordinary split-cache and all-layer native execution agree on all 80 examples and have identical mean margins within each dataset. Last-half execution also ranks all 80 correctly while reducing the margin. Disabled memory is

Table 15: Natural-text transport controls. Margin is the mean gold-minus-other answer logit.

Dataset	Condition	Samples	Ranked exact	Margin
qasper	native_all	40	100%	2.953
qasper	native_disabled	40	50%	0.014
qasper	native_last_half	40	100%	1.009
qasper	native_shuffled	40	50%	-0.108
qasper	ordinary_split	40	100%	2.953
hotpotqa	native_all	40	100%	2.885
hotpotqa	native_disabled	40	52%	-0.002
hotpotqa	native_last_half	40	100%	1.113
hotpotqa	native_shuffled	40	50%	0.166
hotpotqa	ordinary_split	40	100%	2.885

50.0% on QASPER and 52.5% on HotpotQA; shuffled memory is 50.0% on both. The causal controls therefore rule out query-only label bias as the source of native success for this transport task.

13 Original-Answer QA and Quantized Residency

The next experiment removes the inserted answer code. It evaluates original QASPER, HotpotQA, and 2Wiki-MultiHopQA questions and short gold answers. For each example, annotated evidence is placed first and natural distractor or abstract text fills the 384-token source budget. This is an oracle-evidence materialization study: it asks whether the engine transports useful native K/V, not whether PRA routing found the evidence.

Table 16: Natural-QA generation with Qwen3-1.7B-4bit over five seeds. Strict EM is zero because the model often continues after the answer; token F1 and total gold-answer log-probability retain the useful quality signal.

Dataset	Condition	n	Token F1	Gold log-prob.	Resident MiB
qasper	Ordinary split K/V	20	0.136	-28.66	0.0
qasper	Native FP K/V	20	0.136	-28.66	37.6
qasper	Native int8 resident	20	0.147	-27.73	18.8
qasper	Shuffled native K/V	20	0.057	-52.18	37.6
qasper	No memory	20	0.042	-54.53	0.0
hotpotqa	Ordinary split K/V	20	0.129	-10.95	0.0
hotpotqa	Native FP K/V	20	0.129	-10.95	42.0
hotpotqa	Native int8 resident	20	0.140	-10.21	21.0
hotpotqa	Shuffled native K/V	20	0.045	-19.48	42.0
hotpotqa	No memory	20	0.058	-15.68	0.0
2wikimultihopqa	Ordinary split K/V	20	0.059	-10.30	0.0
2wikimultihopqa	Native FP K/V	20	0.059	-10.30	42.0
2wikimultihopqa	Native int8 resident	20	0.059	-9.92	21.0
2wikimultihopqa	Shuffled native K/V	20	0.060	-16.39	42.0
2wikimultihopqa	No memory	20	0.073	-13.88	0.0

Table 16 shows exact full-precision transport parity: ordinary split-cache and native K/V have identical aggregate token F1 and gold answer log-probability on every completed dataset. On QASPER the two paths both reach 0.136 F1 and -28.66 mean log-probability; on HotpotQA both reach 0.129 and -10.95 ; on 2Wiki both reach 0.059 and -10.30 . Shuffling memory lowers F1 to 0.057 and 0.045 on QASPER and HotpotQA and makes gold log-probability substantially worse on all three datasets. On 2Wiki, shuffled and no-memory F1 are slightly higher despite worse gold log-probability, exposing verbose generation as a noisy metric rather than supporting a quality gain. The exact transport parity and answer-likelihood controls establish source-dependent computation; broader decoding evaluation remains necessary.

Per-head int8 residency reduces mean selected storage from 37.6 to 18.8 MiB on QASPER and from 42.0 to 21.0 MiB on HotpotQA. Its aggregate F1 and gold log-probability are slightly higher in this small cohort, but this is not evidence that quantization improves quality: generation is discrete and the sample is small. The supported conclusion is non-degradation here plus a near-exact $2\times$ residency reduction. Dequantization restores model-native K/V before attention, so active materialization bytes are not reduced.

13.1 M4 Pro cross-model replication

The M4 Pro replication repeats all five execution conditions for 20 examples per dataset and model. It is independent of the earlier M5 run and adds Qwen3-4B. Table 17 reports the selector-frozen oracle-evidence comparison; Figure 5 keeps transport quality separate from the shuffled-memory causal control.

Table 17: Cross-model original-answer QA on the M4 Pro. E0 is ordinary split-cache execution; E2 is the same selected evidence as native K/V. FP/int8 is retained selected-K/V residency; int8 is dequantized before attention. The original mixed-state latency column is retired.

Model	Dataset	n	E0 F1	E2 F1	Shuf. F1	FP/int8 MiB
Qwen3-1.7B	2Wiki	20	0.059	0.059	0.060	42.0/21.0
Qwen3-1.7B	HotpotQA	20	0.129	0.129	0.045	42.0/21.0
Qwen3-1.7B	QASPER	20	0.136	0.136	0.057	37.6/18.8
Qwen3-4B	2Wiki	20	0.097	0.097	0.011	54.0/27.0
Qwen3-4B	HotpotQA	20	0.212	0.212	0.028	54.0/27.0
Qwen3-4B	QASPER	20	0.249	0.249	0.019	48.3/24.2




Figure 5: Cross-model transport and causal controls. Full-precision native E2 matches ordinary split-cache aggregate F1 in every cohort, while shuffled memory degrades quality and gold-answer likelihood.

Aggregate E0/E2 F1 is identical in all six cohorts. The 4B model improves E0/E2 F1 from 0.070 to 0.097 on 2Wiki, 0.143 to 0.212 on HotpotQA, and 0.153 to 0.249 on QASPER. Shuffled 4B memory falls to 0.011, 0.028, and 0.019, respectively, while its gold-answer log-probability drops by 6.81–38.58 nats relative to native E2. These interventions support a source-dependent native computation rather than query-only answer bias.

Full-precision residency is 37.6–42.0 MiB at 1.7B and 48.3–54.0 MiB at 4B; compact int8 residency halves each value. This is a retained-storage result: request-local dequantization still restores full-precision active K/V. The former speed comparison is invalid: the ordinary cache-producing forward was lazy and completed after its request timer started, whereas `encode_native_memory` calls `mx.eval` before native request timing. The corrected scaling benchmark synchronizes both representations and reports cold/cold separately from warm/warm. The selector-frozen online E0/E2 benchmark remains the serving-path economic reference.

13.2 Routed evidence rather than oracle evidence

We next remove the oracle ordering. Each HotpotQA and 2Wiki article, or each QASPER paragraph, becomes a typed `AgentResource`. The product SDK’s `PersistentResourceIndex` combines normalized token overlap, indexed BM25, and a deterministic signed-hash semantic control. It ranks the original question against all candidates; the top four are joined in rank order, truncated to the same 384-token budget, and encoded as native K/V. This is a parameter-free product-path baseline, not a learned retriever. Each dataset contains 20 distinct examples sampled under five seeds.

Table 18 exposes the discovery boundary. HotpotQA and 2Wiki have ten candidate documents per question and reach 0.725 and 0.775 recall@4. QASPER averages 46.6 paragraph candidates and reaches only 0.348. The selected sources consume 342–376 tokens. Mean scoring time is 2.85 ms on 2Wiki, 3.57 ms on HotpotQA, and

Table 18: Document routing before MLX materialization. Recall is the fraction of annotated evidence documents recovered in the first k ranks. Route time excludes one-time index construction.

Dataset	n	Candidates	R@1	R@2	R@4	Tokens	Route ms
qasper	20	46.6	0.119	0.140	0.348	370.6	11.92
hotpotqa	20	10.0	0.400	0.600	0.725	376.3	3.57
2wikimultihopqa	20	10.0	0.487	0.662	0.775	341.5	2.85

Table 19: Original-answer generation after routed materialization. Oracle uses annotation-scoped evidence; shuffled supplies another example’s routed K/V.

Dataset	Condition	n	Token F1	Gold log-prob.
qasper	Oracle	20	0.148	-27.25
qasper	Routed	20	0.116	-38.36
qasper	Shuffled	20	0.034	-54.13
qasper	No memory	20	0.026	-60.20
hotpotqa	Oracle	20	0.129	-10.95
hotpotqa	Routed	20	0.091	-15.98
hotpotqa	Shuffled	20	0.054	-20.30
hotpotqa	No memory	20	0.058	-15.68
2wikimultihopqa	Oracle	20	0.059	-10.30
2wikimultihopqa	Routed	20	0.082	-14.31
2wikimultihopqa	Shuffled	20	0.054	-17.00
2wikimultihopqa	No memory	20	0.073	-13.88




Figure 6: Left: evidence discovery improves as four candidates are admitted, but QASPER paragraph routing remains difficult. Right: routed memory carries more answer signal than shuffled memory, while the oracle gap remains.

11.92 ms on QASPER; corresponding one-time index construction means are 6.56, 8.20, and 32.98 ms. These small exhaustive Python cohorts establish cost accounting, not a production latency claim.

The routed ordinary-cache and native-K/V paths produce the same output and exactly equal gold-answer log-probability on all 60 pairs. MLX consumption therefore adds no measured quality error after selection. Routed memory is better than shuffled memory in both F1 and gold likelihood on every dataset. It remains below oracle likelihood, however: QASPER changes from -27.25 to -38.36 , HotpotQA from -10.95 to -15.98 , and 2Wiki from -10.30 to -14.31 . HotpotQA and 2Wiki routed likelihood also does not consistently beat the absent-memory control, despite better routed F1. Partial retrieval can be insufficient or misleading, while short free generation is noisy. The result closes the routed engine path and localizes the remaining deficit to evidence selection; it does not close a production QA-quality gate.

We then expand each dataset to 50 sampled examples (49 unique for QASPER and 50 for the other two). Table 20 confirms exact ordinary/native consumption parity and sharpens the dataset-specific result. On QASPER, routed native memory raises F1 from 0.038 without memory to 0.113, while shuffled memory reaches 0.042; on HotpotQA the corresponding values are 0.054, 0.113, and 0.034. Their gold-answer likelihoods move in the same causal direction. On 2Wiki, however, routed F1 is 0.076 versus 0.081 without memory and 0.089 shuffled, although routed likelihood remains better than shuffled. The larger cohort therefore confirms QASPER/Hotpot utility and a 2Wiki selection or answer-metric failure; it does not support a dataset-averaged quality claim.

Table 20: Expanded routed QA. Selection is frozen before ordinary or native consumption; oracle and shuffled conditions are causal bounds.

Dataset	Condition	Sampled/unique	Recall@4	F1	$\log p(y)$	Route ms
QASPER	no_memory	50/49	0.343	0.038	-51.37	14.27
QASPER	oracle_native	50/49	0.343	0.162	-22.83	14.27
QASPER	routed_native	50/49	0.343	0.113	-36.13	14.27
QASPER	routed_ordinary	50/49	0.343	0.113	-36.13	14.27
QASPER	routed_shuffled	50/49	0.343	0.042	-46.76	14.27
HotpotQA	no_memory	50/50	0.750	0.054	-17.99	4.17
HotpotQA	oracle_native	50/50	0.750	0.134	-11.09	4.17
HotpotQA	routed_native	50/50	0.750	0.113	-14.30	4.17
HotpotQA	routed_ordinary	50/50	0.750	0.113	-14.30	4.17
HotpotQA	routed_shuffled	50/50	0.750	0.034	-20.82	4.17
2Wiki	no_memory	50/50	0.725	0.081	-12.93	3.42
2Wiki	oracle_native	50/50	0.725	0.084	-11.53	3.42
2Wiki	routed_native	50/50	0.725	0.076	-14.12	3.42
2Wiki	routed_ordinary	50/50	0.725	0.076	-14.12	3.42
2Wiki	routed_shuffled	50/50	0.725	0.089	-15.22	3.42

13.3 Optional learned routing across MLX model families

The remaining discovery deficit motivates a controlled adapter comparison, but it does not justify replacing the parameter-free router by default. We freeze the normalized attention-input states produced by exact 4-bit MLX revisions of Qwen3-4B, Qwen3-14B, Llama-3.1-8B, and Gemma-3-1B. For each model, an asymmetric rank-128 linear router is trained on 48 QASPER/HotpotQA examples for 512 steps under five seeds. Selection uses combined validation AUC_{0-30} ; the 32 test identities remain held out. Generic cosine routing and learned routing then see the same chunks, queries, and model states.

The learned router improves QASPER R_{20} for every model, from 0.413 to 0.640 at Qwen3-4B, 0.318 to 0.679 at Qwen3-14B, 0.318 to 0.468 at Llama-3.1-8B, and 0.226 to 0.454 at Gemma-3-1B. It does not transfer uniformly: HotpotQA R_{20} is lower for all four selected routers. This result narrows rather than closes the discovery gap. Public bundles therefore retain generic cosine as the **balanced** profile and expose the adapter as **qasper-learned**. The comparison qualifies routing on exact MLX model identities; it does not alter the selector-frozen E0/E2 transport results or establish end-task generation quality.

Table 21: Held-out routing comparison for exact post-trained HF and quantized MLX model identities. Cohort sizes are retained in the source artifact. The learned router is selected from five seeds using validation data; R_{20} is evidence recall at a 20% chunk budget.

Model	QASPER R_{20}		HotpotQA R_{20}		Combined MRR	
	Generic	Learned	Generic	Learned	Generic	Learned
Qwen2.5-1.5B-Instruct	0.296	0.381	0.385	0.334	0.354	0.423
Qwen2.5-Coder-1.5B-Instruct	0.173	0.394	0.398	0.215	0.346	0.412
Qwen3-4B	0.413	0.640	0.386	0.342	0.333	0.549
Qwen3-4B 8-bit	0.446	0.566	0.407	0.180	0.404	0.499
Qwen3-14B	0.318	0.679	0.494	0.314	0.401	0.645
Qwen3.5-27B 4-bit	0.375	0.610	0.518	0.308	0.501	0.483
Llama-3.1-8B	0.318	0.468	0.616	0.421	0.404	0.537
Gemma-3-1B	0.226	0.454	0.336	0.318	0.316	0.494

13.4 Matched selected-text economics

The routed selections also support a direct E0/E2 comparison without changing the evidence. For 20 examples per dataset across five seeds, E0 restores an ordinary cache snapshot of the selected source and E2 restores its native PRA cache. The question, greedy decoder, and 24-token output budget are identical. The shared schedule includes cold one-shot, two warm repeats, three deterministic query formulations over the same resource, and eight threaded requests sharing one immutable cache object.

Table 22: Matched-selection quality on the shared three-engine cohort. All parity pools 840 request pairs per engine across four schedules.

Engine	Dataset	Visible E0/E2	Reduction	Cold F1 E0/E2	Cold parity	All parity
mlx-lm	2wikimultihopqa	375/33	91.1%	0.067/0.067	100.0%	100.0%
mlx-lm	hotpotqa	415/39	90.6%	0.087/0.087	100.0%	100.0%
mlx-lm	qasper	399/28	93.0%	0.110/0.110	100.0%	100.0%
sglang-mlx	2wikimultihopqa	375/33	91.1%	0.074/0.074	100.0%	100.0%
sglang-mlx	hotpotqa	415/39	90.6%	0.084/0.084	100.0%	100.0%
sglang-mlx	qasper	399/28	93.0%	0.105/0.105	100.0%	100.0%
vllm-metal	2wikimultihopqa	378/33	91.2%	0.079/0.079	100.0%	100.0%
vllm-metal	hotpotqa	417/39	90.6%	0.074/0.074	100.0%	100.0%
vllm-metal	qasper	400/28	93.0%	0.101/0.101	100.0%	100.0%

Table 23: Macro E2/E0 execution-cost ratios; lower is better. Cold includes ingestion, concurrent is inverse requests/s, and E2 usable is milliseconds.

Engine	Cold	Warm	Multi-query	Concurrent	E2 usable ms
mlx-lm	1.130	1.155	1.148	1.132	51.2
sglang-mlx	0.976	1.136	1.140	1.122	55.0
vllm-metal	0.985	1.004	0.993	0.981	86.0

MLX-LM produces exactly the same E0 and E2 output in all 840 pairs, with equal F1 and gold evidence by construction. E2 reduces visible prompt tokens by 90.6–93.0%. Its mean time to usable context is 51.2 ms, but cold end-to-end cost is 1.130 times E0 because the unfused selected cache slows generation. Warm and multi-query ratios are 1.155 and 1.148. Under eight threaded requests, E2 reaches 88.4% of E0 throughput, an inverse cost ratio of 1.132. This is exact semantic transport across all four regimes, but not yet an optimized native-attention kernel.

The expanded 149-question confirmation preserves all 2,086 MLX pairs, with an engine-level Wilson 95% lower bound of 0.9982. Visible-token reduction remains 90.4–93.0%. Macro cold, warm, multi-query, and concurrent cost ratios are 1.043, 1.139, 1.133, and 1.072. The larger cohort thus reduces the measured cold and concurrent penalty relative to the original 20-example-per-dataset run, but confirms that this unfused implementation remains slower




Figure 7: Macro representation cost after freezing retrieval. Lower is better; concurrent cost is inverse throughput. The current unfused MLX native cache is semantically exact but slower in every schedule.

than E0 in every regime.

We isolate the principal decode allocation with a one-normalization segmented attention primitive. Rather than concatenate selected and local K/V, it forms their score blocks separately, shares one maximum and denominator, and sums the two value numerators. We additionally graph-compile that primitive while keeping dynamic causal masks outside the captured graph. Table 24 reports paired M4 Pro kernel measurements with 384 selected tokens, eight heads, and 128-dimensional heads.

Table 24: Eager and graph-compiled segmented attention against concatenation. Ratios below one are faster. Avoided MiB is the temporary concatenated K/V per layer-call. Compilation is an MLX graph optimization, not a custom Metal kernel.

Local tokens	concat ms	eager/concat	compiled/concat	avoided MiB	max error
2,048	0.351	1.053	1.099	9.5	1.22e-04
8,192	0.712	1.040	1.016	33.5	3.42e-05
32,768	2.135	0.990	0.959	129.5	3.05e-05




Figure 8: Segmented-attention kernel cost and temporary-memory avoidance as the occupied local context grows. Graph compilation reaches 0.959 times concat at 32K while avoiding 129.5 MiB per layer-call, but does not remove the full-model request overhead.

The compiled kernel is slower at 2K (1.099 times concat), near parity at 8K (1.016), and faster at 32K (0.959), with maximum absolute error at most 1.22×10^{-4} . The end-to-end check is less favorable. On the same 15 Qwen3-32B natural-QA examples, compiled concatenated E2 is 1.009 times E0 and preserves all outputs; compiled segmented E2 is 1.052 times E0 and agrees on 13/15 outputs. The eager ratios on that subset are 1.012 and 1.048. Graph capture therefore removes neither the residual model-wide overhead nor the reduction-order drift. A genuinely fused Metal kernel remains an open systems optimization.




Figure 9: Expanded 149-question representation-cost confirmation across all three engines. Retrieval and selected evidence are held fixed.




Figure 10: Five-seed compact-K/V pressure over QASPER, HotpotQA, and 2Wiki. Each seed visits eight resources twice and then revisits the first. The budget $K = 8$ is the first condition that fits the working set.

Table 25: Natural-QA residency curve. Each row contains 85 requests: 17 requests per seed over five seeds. Reload fraction divides reloads by the nine accesses after each resource’s first load.

Dataset	K	n	Peak MiB	Reload frac.	Resolve ms	Token F1
QASPER	1	85	21.0	1.00	113.4	0.189
QASPER	2	85	41.9	1.00	113.4	0.189
QASPER	4	85	79.7	1.00	112.9	0.189
QASPER	8	85	151.4	0.00	53.2	0.189
HotpotQA	1	85	21.0	1.00	119.1	0.144
HotpotQA	2	85	42.0	1.00	118.6	0.144
HotpotQA	4	85	84.0	1.00	117.5	0.144
HotpotQA	8	85	168.0	0.00	55.2	0.144
2Wiki	1	85	21.0	1.00	117.9	0.049
2Wiki	2	85	42.0	1.00	117.9	0.049
2Wiki	4	85	84.0	1.00	118.5	0.049
2Wiki	8	85	168.0	0.00	57.2	0.049

Figure 10 and Table 25 validate physical lifecycle rather than theoretical byte sums. Each request pins one selected block, dequantizes it locally, generates, and releases it. With budgets $K \in \{1, 2, 4\}$, the cyclic eight-resource working set reloads all nine repeated accesses per seed. At $K = 8$, initial loads fall from 17 to eight, evictions and reloads fall to zero, and mean resolution falls from 113–119 ms to 53–57 ms. Peak compact residency rises from approximately 21 MiB at $K = 1$ to 151 MiB on QASPER and 168 MiB on HotpotQA and 2Wiki at $K = 8$.

Token F1 is unchanged across all four budgets within every dataset, and mean completion latency remains within roughly 5 ms because generation dominates the saved resolution time. The result is therefore a working-set threshold, not evidence that more residency generally improves answer quality or complete request latency. No pinned block is evicted and compact detail never enters an ordinary prompt or rotating cache.

13.5 Larger-model sustained pressure on M4 Pro

We repeat the bounded-residency experiment with Qwen3-4B on the M4 Pro. Each seed owns eight resources, visits the complete set three times, and revisits the first resource. Four resident budgets produce 25 requests per seed, or 125 requests per table row and 1,500 requests overall.

Table 26: Qwen3-4B sustained compact-K/V pressure on the M4 Pro. Budget is resident/logical resources. Resolve includes compact-detail load and dequantization; p95 retains unavoidable first-ingestion requests.

Dataset	Budget	n	F1	Reload	Resolve ms	p95 ms	Resident MiB
QASPER	1/8	125	0.218	0.680	417.1	460.6	24.4
QASPER	2/8	125	0.218	0.680	420.0	460.7	47.7
QASPER	4/8	125	0.218	0.680	419.3	460.5	91.5
QASPER	8/8	125	0.218	0.000	134.3	460.2	167.4
HotpotQA	1/8	125	0.231	0.680	457.2	460.2	27.0
HotpotQA	2/8	125	0.231	0.680	457.8	461.0	52.9
HotpotQA	4/8	125	0.231	0.680	457.6	461.1	101.5
HotpotQA	8/8	125	0.231	0.000	146.5	460.1	185.8
2Wiki	1/8	125	0.128	0.680	458.8	467.1	27.0
2Wiki	2/8	125	0.128	0.680	461.1	478.1	52.9
2Wiki	4/8	125	0.128	0.680	462.2	480.6	101.5
2Wiki	8/8	125	0.128	0.000	147.1	465.8	185.8




Figure 11: The eight-resource working-set threshold at Qwen3-4B. Fitting the complete set removes repeat reloads and lowers mean resolution; it does not remove each resource’s first cold ingestion.

For budgets 1, 2, and 4, 68.0% of all requests reload compact detail. At budget 8, reloads fall to zero. Mean resolution falls from 417.1 to 134.3 ms on QASPER, 457.2 to 146.5 ms on HotpotQA, and 458.8 to 147.1 ms on 2Wiki, a 67.8–68.0% reduction. Mean compact residency rises to 167.4 MiB on QASPER and 185.8 MiB on HotpotQA and 2Wiki. Token F1 and gold-answer log-probability are identical across all four budgets within each dataset.

The p95 resolve time remains 460–466 ms at budget 8 because eight of each 25-request sequence are compulsory first loads. Residency removes 17 repeat loads, not those cold ingestions. This result therefore supports working-set admission and reuse policy; it does not support a claim that larger budgets improve answer quality or cold-start tails.

13.6 Shared storage lifecycle and MLX realization

The product runtime now defines engine-neutral HOT/WARM/COLD/SOURCE semantics. MLX attention-ready arrays are HOT. Lossless retained arrays or strict fingerprinted files are WARM. Compressed persistence is COLD, with K/V quantization configured independently from compression. Apple unified memory can physically host both HOT and WARM, but readiness and lifecycle remain different service states.

Across the common four-label, five-seed policy workload, lossless tier recovery and SOURCE reconstruction are exact in 20/20 runs. Weighted retention keeps the shared document where plain LRU does not; active-task and dependency guards hold in 20/20, and a grace period avoids 20/20 transient writes. A generic per-array symmetric-int8 codec occupies 0.25 of FP32 bytes with 0.0103 RMSE. That codec result is not an MLX answer-quality claim. The model-native per-head experiment in this paper is the applicable quality evidence and obtains a roughly twofold compact-residency reduction because its metadata and scale layout differ. COLD still dequantizes before attention, so compact persistence does not reduce active attention bytes.

The live storage probe now replaces opaque WARM blobs with independently checksummed, mmap-readable K and V segments per layer. Three QASPER selections generate from MLX HOT arrays, lossless WARM reload, int8 COLD reload, and a fresh manager after restart. WARM reproduces all 3/3 HOT outputs and restart recovery succeeds. Median promotion-plus-generation latency is 327 ms HOT and 644 ms WARM (1.97 $\times$); the excluded background WARM transition costs 348 ms. COLD promotion plus generation is 600 ms and its background conversion costs 940 ms.

Table 27: Shared three-example live storage probe. Prefix is mean HOT/int8 common output prefix; ratios use lifecycle-inclusive request latency.

Engine	WARM exact	int8 exact	int8 first	Prefix	WARM/HOT	COLD/HOT
mlx-lm	100%	0%	67%	1.3	1.97	1.84
sglang-mlx	100%	0%	67%	0.7	1.90	1.89
vllm-metal	100%	0%	67%	4.0	1.96	2.00

Per-array int8 COLD changes all three 24-token sequences and preserves the first token in two. Its positive mean F1 change (+0.068) is an unstable three-example accident rather than evidence of improved quality: exact parity is zero and the mean common prefix is 1.3 tokens. Lossless persistence is the default; quantized COLD requires a larger model- and workload-specific quality gate. The separate per-head compact-residency experiment remains relevant but does not override this autoregressive sequence diagnostic.

We expand the protocol to 80 examples: all three datasets at Qwen3-0.6B and QASPER at Qwen3-1.7B. Lossless WARM remains 80/80 exact and all four managers recover after restart; int8 COLD is exact in 13/80 sequences. WARM/HOT p50 ratios are 2.01–2.15, with WARM p95 between 710 and 817 ms. The larger cohort therefore strengthens persistence correctness while making request-time promotion, rather than array storage, the immediate systems bottleneck. Two five-example QASPER replications add Llama-3.2-1B and Gemma-3-1B; both are 5/5 lossless-WARM exact with restart recovery. Llama int8 is 2/5 exact and Gemma is 5/5 in this small cohort, which is insufficient to override the broader Qwen-based opt-in gate.

13.7 Asynchronous promotion and online request behavior

The lifecycle ratios above charge a complete WARM promotion to the requesting thread. We therefore move promotion ownership to the event loop, deduplicate concurrent promotions by logical key, admit completed objects to a bounded hot set, and overlap promotion with query-side work. Table 29 and Figure 12 report five QASPER

Table 28: Expanded MLX lifecycle cohorts. Latencies include synchronous promotion and generation; int8 remains diagnostic.

Engine	Model	Data	n	WARM exact	int8 exact	HOT p50	WARM p50	WARM p95
mlx-lm	Llama-3.2-1B-Instruct-4bit	qasper	5	100.0%	40.0%	191	275	281
mlx-lm	Qwen3-0.6B-4bit	2wikimultihopqa	20	100.0%	15.0%	328	658	710
mlx-lm	Qwen3-0.6B-4bit	hotpotqa	20	100.0%	15.0%	329	707	817
mlx-lm	Qwen3-0.6B-4bit	qasper	20	100.0%	15.0%	326	672	714
mlx-lm	Qwen3-1.7B-4bit	qasper	20	100.0%	20.0%	337	723	744
mlx-lm	gemma-3-1b-it-4bit	qasper	5	100.0%	100.0%	391	480	502

selections at each lead. Every generation remains HOT-exact. With no lead, median demand stall is 383.7 ms and demand-path cost is 2.123 times HOT. A 250 ms lead lowers these to 90.3 ms and 1.310; 350 ms makes 40% ready at demand and reaches 1.048; 500 ms makes all five ready and reaches 0.983. Promotion is therefore hideable when the scheduler has enough useful work, but it has not become intrinsically cheaper.

Table 29: Event-loop-owned WARM promotion. Lead and stall are milliseconds.

Engine	Lead	Ready	Stall p50	Stall p95	Demand/HOT
mlx	0 ms	0%	383.7	397.6	2.123
mlx	25 ms	0%	332.7	358.1	2.054
mlx	50 ms	0%	304.3	328.3	1.962
mlx	100 ms	0%	256.4	305.5	1.836
mlx	250 ms	0%	90.3	122.4	1.310
mlx	350 ms	40%	4.2	40.8	1.048
mlx	500 ms	100%	0.1	0.1	0.983
sglang	0 ms	0%	210.3	280.1	1.561
sglang	25 ms	0%	202.4	240.7	1.495
sglang	50 ms	0%	169.0	195.7	1.475
sglang	100 ms	0%	111.9	141.4	1.326
sglang	250 ms	80%	0.0	5.7	1.035
sglang	350 ms	100%	0.0	0.1	1.063
sglang	500 ms	100%	0.0	0.0	1.023

The native executor is also exercised through the OpenAI-compatible gateway, including streaming, cancellation, cleanup, and session ownership. A shared gateway correction normalizes native token rows into portable deltas and a terminal completion event. At concurrency one through sixteen, throughput is 1.52–1.70 requests/s; TTFT p95 rises from 302 to 9,988 ms as the deliberately serialized model runner queues. Every one of the 31 concurrent outputs matches the prime output. The client cancels after the first streamed token (364 ms) and all 34 exercised sessions close without retaining request-owned native detail. This is a real HTTP-path measurement, but the single MLX model instance is intentionally serialized to avoid unsafe concurrent mutation of prompt-cache arrays.

13.8 Concurrent physical-tier curves

The lifecycle manager is next stressed at concurrency $c \in \{1, 2, 4, 8, 16\}$ with either one shared immutable resource or one resource per request. The model runner remains deliberately serialized, so these finite waves measure queueing, promotion, and physical sharing rather than continuous batching. All 124 HOT/WARM requests reproduce their HOT baselines exactly. At $c = 16$, shared HOT and WARM sustain 3.20 and 2.91 requests/s with request p95 of 4,996 and 5,491 ms. The shared WARM wave promotes one object and avoids 630 MiB of duplicate physical K/V. Independent WARM instead promotes 16 objects, reads 1.26 GiB, reaches 1.22 requests/s, and has 13,087 ms p95. The gap isolates resource sharing and promotion amplification; it is not evidence for parallel model execution. Int8 COLD remains outside the lossless gate.




Figure 12: Promotion lead trades demand-path stall for speculative work. The dashed line is HOT request latency.

Table 30: Online native-gateway concurrency, including queueing. Times are ms.

Engine	c	req/s	TTFT p50	TTFT p95	request p50	request p95
mlx-lm	1	1.70	302	302	589	589
mlx-lm	2	1.52	355	1028	642	1316
mlx-lm	4	1.66	950	2138	1233	2415
mlx-lm	8	1.58	2247	4779	2538	5066
mlx-lm	16	1.56	4838	9988	5123	10275
sglang-mlx	1	3.06	36	36	325	325
sglang-mlx	2	2.94	43	381	353	678
sglang-mlx	4	3.06	369	1007	670	1305
sglang-mlx	8	3.00	1037	2364	1345	2667
sglang-mlx	16	2.99	2337	5027	2619	5326




Figure 13: MLX physical-tier concurrency. Shared WARM amortizes one promotion; independent WARM scales storage work with the number of resources.

13.9 Selective quantization and sustained pressure

Twenty QASPER examples isolate which tensors can tolerate int8 COLD storage. Quantizing every key and value is sequence-exact in 4/20 cases. Keys alone are 5/20; values alone are 10/20; the late half is 11/20; and the late quarter is 14/20, with first-token parity in all 20 and a mean 18.85-token common prefix. Aggregate F1 changes remain small, but autoregressive divergence makes even the late-quarter profile a calibration candidate rather than a default.

A separate sustained-session study issues 1,125 requests over QASPER, HotpotQA, and 2Wiki: five seeds, eight resources, three complete rounds and a final revisit, under HOT capacities two, four, and eight. Capacities two and four average 17 reloads per seed and reload the final revisit in every case; capacity eight has no reload or eviction. F1 and answer log-probability are identical across budgets within each dataset. The intervention changes physical reuse, not selected evidence or model semantics.

Table 33: Long-session pressure. Reloads and evictions are means over five seeds.

Dataset	HOT resources	Requests	Reloads	Evictions	Final reload	F1	p95 ms
qasper	2	125	17.0	23.0	100%	0.172	493
qasper	4	125	17.0	21.0	100%	0.172	493
qasper	8	125	0.0	0.0	0%	0.172	495
hotpotqa	2	125	17.0	23.0	100%	0.131	497
hotpotqa	4	125	17.0	21.0	100%	0.131	500
hotpotqa	8	125	0.0	0.0	0%	0.131	501
2wikimultihopqa	2	125	17.0	23.0	100%	0.064	497
2wikimultihopqa	4	125	17.0	21.0	100%	0.064	508
2wikimultihopqa	8	125	0.0	0.0	0%	0.064	502




Figure 14: The eight-resource working-set threshold appears consistently across all three datasets.

The final pressure experiment crosses physical HOT and WARM budgets with local rotating windows of 64 and 256 tokens over 1,000 QASPER requests. With HOT/WARM = 2/2, 75.2% of requests reconstruct from SOURCE; raising WARM to eight removes SOURCE starts. With HOT/WARM = 8/2, starts are 68% HOT, 8% WARM, and 24% SOURCE; at 8/8, the latter two become 32% WARM and zero SOURCE. F1 is 0.123 in every cell, and changing the local window does not materially move completion p95. Disk-backed WARM also fails to improve resolve p95 over this mixture of small-source reconstruction: retaining all eight objects reaches 400–429 ms, while the tighter WARM budget reaches 347–398 ms. The product implication is cost-aware admission, not maximal persistence.

The deterministic admission calculation makes that implication concrete. Across this pressure trace, SOURCE reconstruction averages 335.7 ms and WARM restore averages 346.0 ms; creating the WARM representation costs another 396.8 ms on average. Thus

$$r(C_{\text{source}} - C_{\text{warm}}) > C_{\text{persist}}$$

has no positive break-even reuse count on this cohort. The correct default for these small reconstructable QASPER records is SOURCE, not WARM. This is a hardware- and representation-specific decision, not a rejection of

Table 31: Complete finite-wave storage-concurrency results. Times are ms; exact/HOT is the transport-fidelity gate.

Workload	Tier	c	req/s	p50	p95	p99	queue p95	exact/HOT
shared_resource	hot	1	3.18	313	313	313	0	100.0%
shared_resource	hot	2	3.19	314	621	621	308	100.0%
shared_resource	hot	4	3.12	626	1274	1274	944	100.0%
shared_resource	hot	8	3.22	1226	2484	2484	2165	100.0%
shared_resource	hot	16	3.20	2470	4996	4996	4688	100.0%
independent_resources	hot	1	2.98	335	335	335	0	100.0%
independent_resources	hot	2	3.02	331	657	657	328	100.0%
independent_resources	hot	4	3.02	666	1313	1313	997	100.0%
independent_resources	hot	8	3.06	1328	2607	2607	2274	100.0%
independent_resources	hot	16	2.96	2781	5384	5384	5082	100.0%
shared_resource	warm	1	1.14	877	877	877	0	100.0%
shared_resource	warm	2	2.00	710	1001	1001	710	100.0%
shared_resource	warm	4	2.44	999	1638	1638	1318	100.0%
shared_resource	warm	8	2.75	1648	2909	2909	2593	100.0%
shared_resource	warm	16	2.91	2950	5491	5491	5176	100.0%
independent_resources	warm	1	1.44	695	695	695	0	100.0%
independent_resources	warm	2	1.51	723	1327	1327	723	100.0%
independent_resources	warm	4	1.48	1377	2710	2710	1963	100.0%
independent_resources	warm	8	1.33	2938	5999	5999	5246	100.0%
independent_resources	warm	16	1.22	6673	13086	13086	12310	100.0%
shared_resource	cold_int8	1	0.89	1123	1123	1123	0	0.0%
shared_resource	cold_int8	2	1.95	708	1024	1024	708	0.0%
shared_resource	cold_int8	4	2.40	1037	1668	1668	1358	0.0%
shared_resource	cold_int8	8	2.71	1686	2945	2945	2640	0.0%
shared_resource	cold_int8	16	2.94	2901	5446	5446	5125	0.0%
independent_resources	cold_int8	1	1.50	665	665	665	0	0.0%
independent_resources	cold_int8	2	1.44	715	1389	1389	715	0.0%
independent_resources	cold_int8	4	1.39	1403	2867	2867	2092	0.0%
independent_resources	cold_int8	8	1.30	3173	6135	6135	5373	0.0%
independent_resources	cold_int8	16	1.32	6179	12119	12119	11329	12.5%

Table 32: Selective int8 COLD calibration on 20 QASPER examples.

Profile	Exact	First	Prefix	$\Delta F1$	$\Delta \log p$
lossless	100.0%	100.0%	24.0	+0.0000	+0.0000
all_kv	20.0%	95.0%	8.8	-0.0163	+0.2862
all_keys	25.0%	85.0%	9.8	-0.0039	+0.0905
all_values	50.0%	100.0%	14.6	+0.0121	+0.1798
early_half_kv	15.0%	85.0%	8.6	+0.0012	+0.2602
late_half_kv	55.0%	95.0%	16.6	-0.0028	+0.0771
late_quarter_kv	70.0%	100.0%	18.9	-0.0031	+0.0698

persistent K/V for expensive encoders or remote sources.

The requested 2K, 8K, 32K, and full local-history frontier is not inferred from the 64/256-token cache-capacity control. Those cells require prompts that actually occupy each window. We therefore label the current joint frontier partial: resource budgets 2/8, concurrency 1–16, and consumer-layer profiles are measured, while long local-window occupancy remains a separate staged run. The last-half consumer profile preserves 5/5 controlled recoveries at half the active selected K/V bytes; last-quarter and last-four do not. The larger natural ladder also rejects last-half consumption, so none is a BALANCED or ECONOMY recommendation.

Table 34: Crossed physical-tier and local-window pressure. Start percentages name the representation available before each request. Times are ms.

HOT	WARM	Window	Requests	HOT start	WARM start	SOURCE start	F1	Resolve p95	Total p95
2	2	64	125	0%	25%	75%	0.123	398.2	354
2	2	256	125	0%	25%	75%	0.123	377.9	353
2	8	64	125	0%	100%	0%	0.123	428.6	353
2	8	256	125	0%	100%	0%	0.123	428.4	354
8	2	64	125	68%	8%	24%	0.123	351.9	345
8	2	256	125	68%	8%	24%	0.123	347.2	348
8	8	64	125	68%	32%	0%	0.123	402.8	346
8	8	256	125	68%	32%	0%	0.123	400.0	349




Figure 15: Local-window size has no quality effect in this bounded decode; HOT admission dominates median resolution, while broad WARM retention does not improve the tail on this host.

14 Results Synthesis and Precision Scaling

Across the matched and scaling studies, the stable result is semantic rather than uniformly economic. Native K/V preserves selected-text behavior across the tested Qwen, Llama, and Gemma mechanisms, and its relative cost approaches parity as Qwen model size grows. Reused execution is still slower in the expanded cohort, reduced consumer suffixes fail the held-out quality gate, and lossy COLD storage fails exactness. The next optimization target is therefore fused selected-cache consumption under a full correctness profile.

A companion M5 experiment compares Qwen3-4B BF16, MLX 8-bit, and MLX 4-bit deployed pipelines on the same ten MultiHop-RAG questions, seed, candidate counts, and 2,048-token budget. The pinned BF16 load contains 398 BF16 parameter arrays and no quantized layers. Table 35 shows that BF16 and 8-bit preserve task score, while 4-bit loses one or two of ten answers. The native pipeline’s latency premium also increases from 5.8–6.1% at BF16 to 10.0–10.7% at 8-bit and 14.8–16.1% at 4-bit.

Table 35: Matched Qwen3-4B three-rung deployed-pipeline comparison ($n = 10$, seed 11). Base uses standard selected text; PRA uses hybrid selection and detached native K/V. No learned document-RAG adaptor is qualified.

Encoding	Candidates	No-PRA F1	Native F1	Δ_{NM-NP} F1	Task Δ_{NM-NP}	NM/No-PRA latency
MLX-bfloat16	20	0.1233	0.1214	-0.0019	+0.00	1.058
MLX-bfloat16	50	0.1177	0.1249	+0.0072	+0.00	1.061
MLX-8bit	20	0.0916	0.0881	-0.0035	+0.00	1.107
MLX-8bit	50	0.0903	0.0939	+0.0036	+0.00	1.100
MLX-4bit	20	0.0944	0.0768	-0.0176	-0.20	1.148
MLX-4bit	50	0.0906	0.0815	-0.0091	-0.10	1.161




Figure 16: Quality and latency changes across BF16, INT8, and INT4. Because the selector and transport differ together, this is directional pipeline evidence rather than a selector-frozen precision-invariance result.

14.1 Selector-frozen document-RAG realization

Paper 4.5 defines the retrieval protocol; here we report only its MLX realization result. Fifty held-out MultiHop-RAG questions use Qwen3-4B 4-bit, $N = 20$, a 2K selected-token budget, deterministic generation, and one immutable selection receipt per Selected Context/Native Memory pair. Both the generic PRA selector and a pinned cross-encoder reranker preserve every generated output when selected text is replaced by detached native K/V. Thus the native path does not explain the selector’s quality deficit in this cohort. All primary-arm outputs pass the fixed format check; normalization diagnostics show token overlap rather than exact short-answer matches, so no generation style change is used to manufacture EM parity.

Table 36: MLX selector-frozen realization ($n = 50$). TTFT and total latency are means in milliseconds. Cold includes native encoding; warm native rows reuse retained K/V, while Selected Context is re-prefilled.

Selector	Regime	SC F1	NM F1	SC TTFT	NM TTFT	SC total	NM total
Generic	cold	.071	.071	2633	2648	3024	3121
Generic	warm	.071	.071	2640	166	3031	638
Reranker	cold	.081	.081	2637	2663	3030	3135
Reranker	warm	.081	.081	2643	166	3034	639

Native realization removes about 2,026 source tokens from the visible prompt, leaving 79 visible tokens and retaining approximately 299 MB of active detail. Cold execution costs 96–105 ms more end-to-end than Selected Context. Warm reuse reverses the comparison: TTFT falls by about 2.48 s and total latency by about 2.40 s, with unchanged generated output. This qualifies native realization for repeated identical selections on the measured MLX path; it does not qualify the upstream selector, and it does not establish equivalence to a selected-text path backed by an ordinary prefix-cache hit.

The same protocol at $N = 50$ and 4K preserves every selector-frozen output. Selected Context and Native Memory obtain 0.080 token F1 with generic PRA and 0.081 with the strong reranker; all compared selectors obtain 0.460 official score. Cold native total latency is 6.313 s versus 6.109 s for Selected Context. Retaining the exact 4,084-token native selection reduces total latency to 0.799 s versus 6.121 s for re-prefill, with approximately 602 MB active detail. This extends contiguous realization, not retrieval qualification.

A separate physical run caches independently encoded chunks across 50 changing questions. It reuses 22.0% of chunk lookups, reduces cumulative wall time from 151.19 to 131.78 s, and retains 6.38 GB of detail. It does not preserve output: official score changes from 0.440 for Selected Context to 0.380 for Native Memory. Independently encoded chunks retain source-local post-RoPE positions, whereas visible concatenation assigns global positions. This policy is therefore reported as lifecycle evidence with decomposition gates marked not applicable; it is not included in MLX’s exact contiguous-realization claim.

Qwen3-8B repeats the contiguous experiment on the same 50 questions, 2K budget, and frozen receipts. Generic Selected Context and Native Memory both obtain 0.440 official score and 0.081 F1; the strong-reranker pair obtains 0.480 and 0.096. Every paired output matches. Cold native total latency is 5.636 s versus 5.535 s for Selected Context, while retained native latency is 0.990 s versus 5.541 s for re-prefill. Thus generator scale does not repair selection, but the native realization and repeated-selection result transfer from 4B to 8B on this MLX/Qwen family.

An independent FP32 checkpoint, cross-precision adaptor transfer, peak-memory capture, and a non-Qwen matched replication remain unmeasured. Casting the published BF16 source to float32 would not create an independent FP32 weight reference.

15 Promotion Gates

Gate	Status	Evidence
environment	closed	MLX Metal generation and server repeat reliably
prompt-cache coexistence	smoke closed	364/365 selected tokens reused warm
rotating-cache control	controlled closed	five seeds, four capacities
native semantic parity	controlled closed	five seeds, zero logit error, 5/5 recovery
rotating plus native	controlled closed	5/5 with 64-token local cache
lazy native region	prototype closed	selected resources encode on first selection
consumer profile	open	last 14/28 retains 5/5 on the controlled task; no reduced 4B–32B natural-QA profile passes the smoke gate
persistence	controlled closed	strict fingerprint; zero-error reload in 7.2 ms mean
shared storage lifecycle	controlled closed	segmented WARM 90/90 exact across Qwen/Llama/Gemma; restart recovery; int8 remains gated
concurrency/sharing	controlled closed	five seeds at 1–8 requests; 108 MiB avoided at eight
natural transport	controlled closed	80/80 native; disabled/shuffled near chance
natural QA transport	controlled closed	original QASPER/Hotpot/2Wiki answers with causal controls
routed natural QA	controlled partial	60/60 consumption parity; recall@4 0.348–0.775; oracle gap open
document-RAG contiguous	controlled closed	Qwen3-4B/8B, $n = 50$ each; selector-frozen outputs exact; repeated-selection E2 faster than re-prefill
document-RAG chunk residency	open	22.0% chunk hits and 12.8% lower wall time, but source-local composition changes generation
matched E0/E2 economics	controlled closed	2,086/2,086 expanded outputs; all four E2/E0 cost ratios exceed one
quantized PRA detail	controlled candidate	per-head int8 halves residency; native dequantization before attention
memory pressure	controlled closed	3,145 requests across compact, long-session, and crossed tier/window pressure
SOURCE/WARM admission	controlled closed	WARM is slower than SOURCE on the measured small-record cohort; no positive break-even reuse
large local-window Pareto	open	2K/8K/32K/full require occupied long histories and are not inferred from cache limits

15.1 Observability contract

MLX does not require a container-native telemetry stack. PRA wraps Metal/MLX timing and residency observations in standard spans and normalized Prometheus metrics while the engine remains on the macOS host; Collector, Prometheus, and Grafana may run in containers. The dashboard distinguishes selected tokens, native K/V bytes, HOT/WARM promotion, rotating-cache behavior, TTFT, and ITL. All reported performance qualification keeps detailed tracing disabled.

16 Negative Results, Falsification, and Limitations

Native PRA-MLX is unnecessary if selected visible text plus ordinary MLX caching matches its quality, memory, and latency on realistic long-memory workloads. The server smoke makes that baseline competitive. Conversely, the rotating control shows that bounded sequential KV is not automatically a PRA archive.

The original profile and natural controls center on one small 4-bit Qwen model; Llama/Gemma lifecycle rows

remain smaller mechanism replications. The initial server experiment has only five repeats per condition, so its p95/p99 are not reported; the later gateway load curve reports finite-wave tails but serializes one model runner. The natural controls contain real dataset text but an inserted fixed answer-code mechanism. The first end-task study uses original answers but oracle-scoped evidence. The M4 scaling calibrations reach dense 32B and a 30B-A3B MoE checkpoint, but use only 30 examples per Qwen model. Llama 3.1 8B adds a 60-example cross-family concat-transport replication, but not a cross-family segmented-layer curve. The corrected cost ladder has five seeds but two examples per dataset and seed; its bootstrap intervals are smoke uncertainty, not a production noninferiority analysis. It establishes cost convergence toward parity and rejects the tested reduced-layer profiles, not broad family generalization. The long-context extension likewise has 15 examples per 8K model point and only three at 32K; 64K exceeds the pinned model limit. It measures the MLX allocator rather than whole-system wired memory and does not exercise swap as a benchmark condition. The routed extension now has 50 sampled examples per condition and dataset (49 unique on QASPER), a parameter-free lexical/hash index, and four selected candidates. Strict EM is confounded by verbose continuation. Int8 detail is restored before attention, so it saves retained bytes rather than active attention bytes. The learned-router cohort has only 16 held-out examples per dataset/model and does not include 2Wiki; no learned native Q/K index, continuous batching, or OOM frontier is measured. The residency sweeps use deterministic cyclic schedules rather than a measured production session distribution. The original matched benchmark reuses only 20 examples per dataset; the expanded selector-frozen confirmation increases that cohort but does not establish end-task quality beyond proving that E2 preserves its E0 input. The rotating intervention also changes effective local context and can perturb chat-template semantics.

17 Related Work

MLX targets efficient and flexible machine learning on Apple Silicon [1]. Prompt caching and rotating K/V optimize sequential state; cache quantization reduces its physical size. Retrieval and context compression reduce visible input. PRA instead assigns stable identity to retained non-prefix memory and separates cheap addressing from active native detail. The mechanisms can be combined, but none is a semantic substitute for the others.

18 Conclusion

MLX prompt caching and selected context compose cleanly: the fixed answer is preserved with 76.9% fewer visible tokens, and repeated selected prompts reuse 364 of 365 tokens. Rotating K/V provides large layer-cache savings but fails as an archive in the five-seed control; late selected text does not reliably repair it. The in-process native executor passes the missing mechanism gate: selected non-prefix K/V is bit-exact to ordinary split-cache execution and preserves 5/5 recovery with a rotating local cache. PRA-MLX is therefore a working native prototype. Full-precision native transport also matches ordinary split-cache quality on original-answer QA, while wrong and absent memory fail causally. Per-head int8 halves retained detail without measured degradation in the small cohort. The M4 Pro replication extends this result to Qwen3-4B: ordinary and native aggregate F1 remain equal in all six cohorts and shuffled memory degrades sharply. The earlier 44.5–54.5% completion ratio is retired because it mixed lazy cold E0 with synchronized warm E2. The corrected 4B–32B ladder finds exact concatenated E2 quality, warm cost converging from 1.054 to 1.011 times E0, and a cold advantage contracting from 0.932 to 0.976. Full segmented overhead falls from 1.165 to 1.047, but no reduced consumer suffix passes the natural-QA quality gate. Model scale alone therefore supports cost convergence, not a native speed advantage or a smaller BALANCED consumer fraction. A separate 30-example-per-model M5 replication reaches a 0.980 concat point estimate and 1.031 segmented ratio at 14B, but the concat interval includes parity and 32B exceeds the 16 GiB host’s capacity gate. Its 1,500-request pressure extension reproduces the eight-resource residency threshold and separates repeat-load savings from cold p95. Across a 1,020-request hard-budget sweep, sub-working-set budgets cap compact residency but reload every repeated access; fitting all eight resources removes reloads and halves resolution time without changing F1. With SDK routing, native and ordinary paths remain identical on 60/60 paired generations, but recall@4 ranges from 0.348 to 0.775 and answer likelihood remains below oracle evidence. A five-seed router qualification improves QASPER recall across four MLX identities, but its HotpotQA transfer is negative, so generic routing remains the product default. The next quality improvement still belongs in discovery, not the MLX cache transport. Queued server pressure, an OOM frontier, broader learned-routing objectives, and larger quality cohorts remain necessary before production claims. Matched E2 execution removes roughly 91–93% of visible tokens and preserves every output, but its unfused decode is slower. Last-half consumption remains a narrow controlled result and int8 residency remains a candidate; neither is promoted by the natural scaling ladder.

A segmented one-softmax primitive now replaces the live Qwen attention path and removes K/V concatenation. It retains a 5–17% completion-time overhead and small numerical drift because it is composed from unfused MLX operations; it is not a fused-kernel speed claim.

References

- [1] A. Hannun et al. MLX: Efficient and Flexible Machine Learning on Apple Silicon. 2023.
- [2] PRA Research Program. From Sparse Selection to Physical Savings: A Unified Runtime and Measurement Contract for Progressive Retrieval Attention. Companion manuscript, 2026.