

Persistent Agent Context with Progressive Retrieval Attention: Typed Resource Discovery Before Neural Materialization

Jorge Simao

EInnovator R&D – Center for the Sciences of Computation, Cognition, and Complexity

August 26, 2026

Abstract

Agent runtimes often place every tool definition in every prompt before the model has decided which capability it needs. We separate this process into typed discovery, bounded disclosure, model-side call construction, and host-authorized execution. M0 resolves opaque catalogs up to 8,192 resources; M1 shows causal use of one selected native-K/V definition in a toy decoder. With frozen Qwen3-0.6B, M2 obtains 20/20 exact calls from selected or oracle schemas, versus 0/20 from shuffled, irrelevant, or empty disclosure and 17/20 from all 18 tools. M3 safely executes 20/20 selected calls in an in-memory harness and preserves every result as a typed observation; 17/20 continuations include every returned value. M4 completes 10/15 three-to-five-step workflows with reactive just-in-time disclosure, versus 3/15 with all required tools visible and 0/15 without refresh. M5 then falsifies the proposed speculative extension: a metadata/schema graph raises required-tool recall to .933 but obtains 0/15 task success under static disclosure. M6 also stops negatively: indexed lexical discovery reaches 1.000 single-step Top-1 and token routing reaches .889 multi-step required-tool recall, while native mean K, full token QK, and zero-shot Paper-2.8 projections are worse. M6.5 then tests 306 authored canonical, paraphrastic, colloquial, implicit, contextual, and multilingual requests. On 144 frozen test rows, BM25 falls to .347 Top-1; typed tags reach .743 and a dictionary/embedding hybrid reaches .729. A validation-calibrated staged resolver selects on .597 of requests at .907 selective accuracy. M7 reevaluates frozen native geometry on the same rows: native mean K, full token QK, and two transferred compressors obtain only .049–.063 Top-1. The supported path is therefore model-independent semantic discovery followed by reactive disclosure. M8 removes the manual-registration assumption: ordinary Python signatures and docstrings reach 0.826 Top-1, compared with 0.806 for the rich manual hybrid. Bounded candidate palettes and authoritative typed slices preserve full selected schemas. A controlled zero-configuration ablation attributes the largest isolated gain to a generic synonym dictionary: Top-1 rises from 0.306 for extracted keywords to 0.750. Validation-frozen automatic fusion reaches 0.812, statistically unresolved from the 0.806 manual ceiling on this cohort. Automatic union does not beat fused Top-K at matched budgets, so its generative JIT gate remains closed. A five-seed follow-up scales callable-derived, semantically confusable catalogs from 32 to 8,192 tools. Frozen hybrid Top-1 declines from 0.787 to 0.338. At 8,192 tools and $K = 10$, raw union recovers 0.662 of required tools, but useful precision is only 0.078. Graph speculation, union-default disclosure, and native-Q/K discovery therefore remain bounded experimental modes. M9 adds deterministic typed views for tools and declarative skills. At $K = 8$, selection-to-full skill execution matches full-all strict success at 0.625 while exposing 0.341 of its capability tokens. The two-pass mechanism uses neither callbacks nor dynamic expansion. M10 closes the large-catalog downstream gate. On 40 five-seed 8,192-tool choices at $K = 10$, diversity union obtains 0.375 target recall, 0.667 conditional choice, and 0.250 end-to-end choice, compared with 0.300, 0.083, and 0.025 for fused retrieval. The paired end-to-end effect is exact $p = 0.0117$. At $K = 8$, diversity selection-to-full exposes 0.556 of full-palette capability tokens and obtains 5/20 strict accepted calls. All

ten two-step workflows still fail because per-step target recall is only 0.250. Diversity union is therefore a useful experimental candidate generator, not the SDK default.

1 The Agent Catalog Problem

An append-only agent prompt couples the size of the logical environment to the transformer’s active context. Tool schemas, skills, system instructions, artifacts, and old observations are repeatedly exposed even when a turn uses only one of them. Compaction can reduce prompt length, but it may erase provenance or stale-versus-current distinctions.

PRA suggests a different decomposition:

agent state = bounded native workspace + persistent addressable resources.

This decomposition does not make storage free. Source text, indexes, cached representations, resident K/V, transferred K/V, and attention-visible K/V are different costs. It instead introduces a concrete question before materialization: which resource identity should be resolved, by which discovery mechanism, and with enough confidence to act?

This paper reports a staged sequence of mechanism gates. M0 asks whether a typed resolver can preserve selection quality as a catalog grows, whether policy hints reduce unnecessary search, and whether confidence prevents an accurate ranker from becoming an unsafe selector. M1 asks whether a small decoder causally uses one selected definition after it is encoded as native K/V. M2–M4 add a frozen pretrained model, typed execution, and multi-step reactive discovery. M5 asks whether capability neighborhoods should be disclosed before planning. M6 asks whether model-native query/key state improves resource discovery enough to justify a tighter deployment boundary. M6.5 asks how far an external semantic resolver can go when canonical lexical overlap disappears, and M7 repeats the native comparison on exactly that frozen semantic-hard test set. Persistent-session recall, production speed at large catalog sizes, and external side effects remain outside these gates. M8 asks whether ordinary callable metadata is sufficient for registration, whether bounded candidate unions improve the recall–context frontier, and whether PRA should honor externally selected tool records as atomic authoritative units. Its final ablation isolates raw text, automatic keywords, generic synonyms, inferred concepts, automatic tags, compact embeddings, and bounded channel union without changing the frozen catalog or test identities. Its final scale test preserves those identities while adding seeded callable-derived distractors through 8,192 tools. M9 finally separates the information needed to choose a capability from the information needed to use it. It compares full-all, selection-only, selection-to-full, and oracle-full views for tools and declarative skills.

2 Typed, Versioned Resources

A resource has a stable typed identity, versioned source, aliases, schema, side-effect class, tenant, and revocation state. Representative identifiers are

```
!!ref:tool:calendar:create_event:v3!!  
!!ref:skill:paper-review:v2!!  
!!ref:artifact:experiment-17!!  
!!ref:session:parent/decision-14!!
```

Identity is distinct from text, index entries, semantic summaries, cached K/V, and current residency. A cache fingerprint therefore includes source/version, model, tokenizer, routing configuration, and any position-sensitive encoding state.

The implementation in `src/prahf/agent_resources.py` enforces tenant, revocation, and allowed-URI filters before ranking. Discovery returns a trace, not only a winner: the executed channel path, per-candidate provenance, confidence, fallback count, and final select/ask/abstain decision remain available to the caller. The resolver never authorizes or executes a tool.

3 Discovery Channels and Policy

The SDK exposes

$$\pi \in \{\text{auto}, \text{explicit}, \text{token}, \text{index}, \text{semantic}, \text{hybrid}, \text{adaptive}\}.$$

Hints can be attached to a request or inherited from a resource collection. Strict hints execute one channel. Non-strict hints specify a starting point and permit calibrated fallback.

Explicit. A typed URI is resolved by exact map lookup. This is deterministic after access checks.

Token. Names, aliases, descriptions, and token n-grams are compared lexically. The M0 implementation is deliberately exhaustive and therefore a correctness control, not a production latency path.

Index. Persistent exact, token-posting, n-gram, and BM25 structures narrow the candidate set before scoring. Index fingerprints make freshness auditable.

Semantic. M0 uses a declared 96-dimensional signed-hash control over the resource’s action, object, and family concepts. It contains no resource ID. This idealized control isolates the policy mechanism; it is not evidence about a learned embedding model. M6.5 replaces that control with two deployable external families: typed operation/object concepts with source provenance, and compact local encoders whose tool vectors are precomputed at registration. Neither family calls the generation model or reads its hidden state.

Hybrid and adaptive. Hybrid fuses all available evidence in one expensive stage. Adaptive begins cheaply and escalates only when calibrated confidence is insufficient. A representative ladder is

$$\text{explicit} \rightarrow \text{token} \rightarrow \text{index} \rightarrow \text{semantic} \rightarrow \text{hybrid} \rightarrow \text{ask/abstain}.$$

The concrete path is recorded per request.

3.1 Confidence is part of correctness

For candidate r_i , the calibrator estimates

$$c_i \approx P(r_i \text{ is intended} \mid q, \mathbf{s}_i, p_i),$$

where $\mathbf{s}_i$ denotes channel scores and p_i their provenance. High confidence permits selection; intermediate confidence asks or searches more; low confidence abstains. For a positive request, the actionable outcome is correct only when the right resource is ranked first *and* the decision is **select**. For an ambiguous or nonexistent request, it is correct only when the decision is **ask** or **abstain**. This definition prevents a low-confidence ranker from receiving operational credit merely because its latent top candidate was right.

3.2 Discovery and disclosure are different controls

Discovery answers how candidates are found: explicit URI, token matching, persistent index, semantic score, or a calibrated combination. Disclosure answers which bounded capability set becomes visible to the model. A request can therefore pair indexed discovery with minimal disclosure, or explicit discovery with a broader planning neighborhood. The SDK represents this as

$$\text{AgentResourcePolicy} = (\text{DiscoveryPolicy}, \text{ToolDisclosurePolicy}),$$

rather than overloading one `top_k`. Resolve mode targets one identity $q \mapsto r^*$. Explore mode targets a set $(q, \text{state}) \mapsto R_K$ and requires recall, precision, unsafe exposure, plan coverage, and task-success metrics. Neither form authorizes execution.

The M5 implementation builds a typed capability graph from category, object, API and operation families, curated tags, deterministic description keywords, and directional producer-consumer schema edges. Every disclosed identity retains its root, edge type, depth, weight, and inclusion source. Destructive neighbors are suppressed by default. Fixed `minimal`, `local`, `planning`, `broad`, and `adaptive` profiles bound each signal independently.

4 M0 Experimental Design

4.1 Opaque catalogs

The generator in `src/data/agent_resources.py` creates identifiers such as `tool_zaf_193`. Each resource combines held-out action, object, and family concepts with aliases, a schema, a version, and a side-effect class. The opaque names prevent pretrained knowledge from solving the benchmark. Catalogs contain 8, 32, 128, 512, 2,048, and 8,192 resources.

Each held-out identity produces six positive requests: explicit URI, exact name, alias, typo, semantic paraphrase, and descriptive reference. Each split also contains one ambiguous and one nonexistent request. Validation identities calibrate thresholds; test identities are disjoint. At sizes 32–8,192, each policy is evaluated on 540 positive and 10 negative test requests across seeds {11, 23, 37, 53, 71}. The 8-resource condition uses 180 positive and 10 negative requests. The complete artifact contains 31,680 validation and test policy traces.

4.2 Policies and metrics

We compare five fixed policies (explicit, token, index, semantic, and hybrid), unguided `auto`, a request-level hint, and adaptive fallback. The oracle is the least-cost fixed policy that ranks the target first on that query. Reported metrics include top-1 accuracy, selection coverage, selective and actionable outcome accuracy, safe nonaction and false-action rates, retrieval stages, fallback count, Brier score, expected calibration error (ECE), quality regret, and policy-cost regret. Confidence thresholds are fitted only on the validation split. Intervals are percentile bootstrap intervals over five seed means; deterministic outcomes can consequently have zero-width intervals.

4.3 Systems accounting

Cold index construction and a full immutable rebuild after mutating one percent of resource versions are timed separately. Warm component latency is the mean of repeated CPU Python calls. Policy-path latency is reconstructed by summing the measured components actually executed by each trace. These timings exclude tokenization, a model forward pass, native-K/V encoding or transfer,

attention, and tool execution. They characterize this reference implementation, not an optimized serving system.

5 Results

5.1 Selection scales, but policies differ operationally

Table 1 reports the largest catalog. Hinted and adaptive policies select every positive target and safely reject every negative request. Their equal quality hides a meaningful systems difference: hints reduce the mean path from 2.39 stages to 1.08 and the measured component estimate from 1,084 to 691 ms per request. The fixed hybrid ranker also attains 100% top-1, but calibration leaves 1.7% of positive requests unselected. Its exhaustive 1,405 ms component cost and 4.17 cost-regret units make it a poor default.

Table 1: Held-out M0 results at 8,192 resources. Top-1 is computed on 540 positive requests; outcome includes ten ambiguous/nonexistent controls. Time is reconstructed CPU Python discovery-component time.

Policy	Top-1	Coverage	Outcome	Safe neg.	Stages	ms/query
Explicit	.167	.167	.182	1.000	1.00	5
Token	.713	.537	.545	1.000	1.00	1,174
Index	.833	.000	.018	1.000	1.00	27
Semantic	.333	.000	.018	1.000	1.00	339
Hybrid	1.000	.983	.984	1.000	1.00	1,405
Auto	.833	.000	.018	1.000	1.00	24
User hint	1.000	1.000	1.000	1.000	1.08	691
Adaptive	1.000	1.000	1.000	1.000	2.39	1,084

The unguided automatic policy exposes why ranking is insufficient. It ranks explicit, name, alias, typo, and description requests correctly, misses the semantic-paraphrase stratum, and obtains .833 top-1. Calibration nevertheless causes it to ask on every positive request, so its actionable outcome accuracy is only $2/110 = .018$ per seed: the two safe negative decisions. Fixed index has the same operational pattern. Conversely, fixed token acts selectively and is always right when it acts, but covers only .537 of positive requests.




Figure 1: Positive-request top-1 accuracy as catalog size increases. Ranking alone does not encode whether confidence permits selection.

Figure 1 shows that hints and adaptive fallback maintain 1.0 top-1 at every tested size. The result is not uniformly safe at the smallest size: in the 8-resource condition each policy attains 1.0 positive top-1, but hinted and adaptive routing each falsely select the ambiguous held-out case in all five seeds. Their overall outcome accuracy is .974 and safe-nonaction rate is .5. From 32 resources onward, both negative cases are handled correctly in all seeds. This inversion is a calibration artifact of the controlled score distribution and cautions against treating larger catalogs as automatically harder along every axis.

5.2 No fixed channel is the oracle everywhere

At 8,192 resources, the least-cost correct oracle chooses explicit lookup for 16.7% of positive requests, indexed lookup for 66.7%, and semantic lookup for 16.7%. Neither token nor hybrid is ever the least-cost correct choice. User hints have zero quality regret and 0.083 cost-regret units, compared with zero quality regret and 0.5 cost regret for adaptive fallback. This supports a narrow claim: when callers know the reference form, passing that information avoids discovery work; adaptive escalation remains useful when they do not.




Figure 2: Largest-catalog ranking quality versus measured discovery-component cost. The log axis is needed because explicit and exhaustive hybrid paths differ by more than two orders of magnitude.

5.3 Logical catalog size is not active materialization

Every generated definition contains 18 accounting tokens. The largest catalog therefore contains 147,456 logical definition tokens, whereas resolving one resource requests 18 tokens, or 0.0122% of the catalog. Figure 3 makes that distinction visible. The number is an accounting upper bound on selected source text; M0 does not encode definitions into native K/V and therefore does not establish attention-memory or latency savings.




Figure 3: Logical catalog definitions grow with catalog size while the one-item selected-definition accounting remains at 18 tokens. The axes intentionally separate logical storage from active materialization.

5.4 Persistent indexing helps, but is not yet production-ready

At 8,192 resources, index construction takes 1.727 s, a full rebuild after a one-percent version mutation takes 1.753 s, and the estimated Python index payload is 5.27 MiB. Warm indexed lookup takes 27.4 ms and scores 966 candidates on average (11.8% of the catalog). By comparison, exhaustive semantic, token, and hybrid controls take 339, 1,174, and 1,405 ms. The index build amortizes to 44.6 ms/query after 100 uses and 29.1 ms after 1,000 uses.




Figure 4: Warm CPU Python component latency. Exhaustive channels scale poorly; the indexed path narrows candidates but still touches about 11.8% at 8K.

These numbers establish an optimization target rather than a speed claim. Production token postings should avoid Python object traversal, semantic search should use an approximate vector index, unions should be deduplicated before scoring, and mutation should update affected postings rather than rebuild the entire immutable index. Cold/warm serving measurements must also include tokenization and downstream native-K/V work.

6 What the M0 Gate Establishes

The mechanism satisfies four narrow requirements. First, typed resources can be resolved without exposing the full catalog to a language model. Second, different reference forms favor different channels. Third, request hints can match adaptive quality with fewer stages. Fourth, action-aware calibration can distinguish a correct ranking from a decision safe enough to use.

M0 alone does *not* satisfy the paper’s broader context gate. Its 18-token active count is not a measured model working set. There is no model-side argument generation, observation-conditioned continuation, persistent-history recall, skill composition, or subagent inheritance. Its negative controls are small and its semantic control is idealized. We therefore treat M0 as resolver evidence rather than evidence that PRA improves agents.

7 M1: A Causal Toy-Model Gate

7.1 Protocol separates discovery, generation, and execution

The synthetic language gives each tool an opaque identity and a compact definition. Semantics appear only in that definition. A held-out query states an action, object, family, and raw argument, but contains neither the resource identity nor one of four formatting schemas. The host first discovers a stable URI and binds it to a temporary slot such as @0. The model must emit `slot|schema|argument`; the host resolves the slot and accepts the call only when the selected URI, schema, and formatted argument are all exact. A deterministic observation is then supplied and the model must continue with the expected answer. No external side effect is performed.

This boundary matters. Discovery quality is measured by stable URI identity. The model is responsible for schema-conditioned call construction. The host is responsible for binding and validation. A correct retrieval cannot receive end-to-end credit for an incorrect call, and a memorized call string cannot receive selection credit for an incorrect resource.

7.2 Model, training, and causal conditions

The character-level decoder has 506,400 parameters, four layers, hidden width 96, feed-forward width 384, RoPE, and PRA native-K/V at layers 1 and 3. Each of five seeds trains for 3,000 steps with a mixture of native-memory examples (50%), directly presented selected definitions (30%), and eager eight-item catalogs (20%). Loss is applied only to the call and post-observation continuation. Evaluation uses eight examples per seed at catalog sizes 8, 32, and 128.

We compare oracle memory, idealized semantically discovered memory, shuffled memory, disabled memory, the selected definition directly in the native prompt, and the eager catalog. All conditions share model weights. Shuffling preserves memory presence and size but selects the wrong definition; disabling removes memory. Eager conditions that exceed the 512-token native window are recorded as context overflow and are not truncated into an easier task.

7.3 Selected native K/V is used causally

Table 2 reports the strict end-to-end metric: correct resource, exact call, and exact continuation. Discovered memory matches oracle memory at all three sizes. Shuffled and disabled conditions obtain zero, so performance cannot be explained by the query alone or by generic memory activation.

Table 2: M1 end-to-end success, averaged over five seeds and eight held-out examples per seed. A dash denotes native-window overflow rather than failure.

Catalog	Discovered	Oracle	Direct	Shuffled	Disabled	Eager
8	1.000	1.000	0.925	0.000	0.000	0.025
32	0.950	0.950	0.825	0.000	0.000	–
128	0.925	0.925	0.800	0.000	0.000	–




Figure 5: M1 end-to-end success. Oracle and discovered curves overlap; shuffled and disabled controls remain at zero. Eager evaluation is defined only where the complete catalog fits.

Every seed favors oracle memory over shuffled and disabled memory at every catalog size. With five paired seeds, however, the minimum attainable exact two-sided sign-flip value is $p = .0625$. Oracle and discovered memory have zero paired difference. Direct presentation is lower by .075 at size 8 and .125 at sizes 32 and 128, but those paired differences are not statistically decisive. We report the direction without claiming that PRA presentation is intrinsically better than direct presentation.

7.4 Logical growth is separated from active K/V

Figure 6 measures the neural working set that M0 could only account for. Logical catalog text grows with catalog size and crosses the native window before size 32. The selected native-K/V payload remains near one definition per PRA layer. Direct selected-definition prompts also remain bounded,

making them the relevant non-PRA control. This is evidence of bounded active materialization in the toy implementation, not a production memory or latency claim.




Figure 6: M1 logical catalog text, materialized native K/V per PRA layer, direct selected prompt length, and the 512-token native window. Logical storage and attention-visible working memory are distinct quantities.

The causal result remains narrow. The semantic discoverer is an idealized control, not a learned production index. The grammar is synthetic, the executor is pure and deterministic, and the decoder is trained specifically for this task. The experiment establishes that the selected native K/V carries necessary information through call construction and observation-conditioned continuation; it does not establish general tool use.

8 M2–M4: From Selection to Typed Execution

8.1 Frozen pretrained bridge

M2 freezes Qwen3-0.6B at revision `c1899de289a04d12100db370d81485cdf75e47ca`. Four realistic single-tool tasks cross five deterministic prompt presentations. The model receives OpenAI-compatible function schemas and must emit one parseable `<tool_call>` with the exact function and arguments. Conditions expose the selected tool, the identical oracle tool, a same-family shuffled tool, an irrelevant tool, no tools, or all 18 catalog tools. The five presentations vary instruction wording and eager order; they are not independent model-training seeds.

Selected and oracle disclosure both obtain 1.000 strict success (20/20). Shuffled, irrelevant, and empty conditions obtain 0/20. Eager disclosure obtains 0.850 (17/20). Thus the selected schema causes the output and a small model can use it without task-specific training; the eager catalog is not an oracle and already introduces interference.

8.2 Safe execution and typed observations

M3 inserts a host boundary after generation. The parser treats model JSON as untrusted. The executor checks that the function was disclosed, the stable URI was authorized, required arguments and types match the schema, and write or destructive authority was granted separately. Only registered pure in-memory handlers can run. A result becomes

```
!!ref:observation:<namespace>:<call-id>:v1!!
```

with producer URI, result schema, call identity, tenant, and provenance. Tool outputs are therefore retrievable resources rather than irreversible transcript text. All 20 selected calls execute and create typed observations. In 17/20 continuations (0.850), the model includes every returned value. The remaining three are consumption failures, despite correct execution.

8.3 Reactive multi-tool composition

M4 uses three held-out workflows of length three, four, and five. Reactive JIT shows only the next discovered schema after each typed result; eager-required shows the complete evaluator-required set at every step; no-refresh leaves only the initial tool visible. The same model and prompt presentations are used. Strict task success requires every expected identity, exact arguments, accepted execution, and typed observation in order.




Figure 7: M2–M4 strict gates. Selected definitions cause exact pretrained calls; typed execution is reliable, while observation consumption remains imperfect. Reactive disclosure substantially outperforms both a static required set and a no-refresh control on the small workflow cohort.

Reactive JIT completes 0.667 (10/15), eager-required 0.200 (3/15), and no-refresh 0/15. The three- and four-step families account for the reactive successes; the five-step cross-domain repository workflow fails after its first result in every presentation. M4 is therefore a positive reactive-composition gate, not general long-horizon agent competence.

9 M5: Speculative Capability Disclosure

M5 asks whether exposing a useful capability neighborhood before the first call improves planning relative to direct Top-1 or reactive retrieval. The 18-node graph has 171 typed edges. We compare all-tools eager (P0), direct Top-1 and Top- K (P1–P2), family/category, tag/keyword, and schema expansion (P3–P5), their bounded union (P6), reactive JIT (P7), horizon-matched speculative disclosure (P8), and the exact required set (P9). Labels distinguish required, useful, related-but-unnecessary, irrelevant, and unsafe tools.

Table 3: M5 capability-set and execution results over three workflows. Recall and initial tool counts are averaged across tasks. Model success has 15 trials for P1, P6–P9; a dash denotes a set-only policy.

Policy	Required recall	Initial tools	Unsafe	Model success
P0 all eager	1.000	18.00	1.00	–
P1 direct Top-1	.261	1.00	0	.000
P2 direct Top- K	.822	4.00	0	–
P5 schema graph	.628	3.67	0	–
P6 combined graph	0.933	6.67	0	0.000
P7 reactive JIT	.261	1.00	0	.667
P8 speculative matched	.628	3.67	0	.000
P9 oracle capabilities	1.000	4.00	0	.200




Figure 8: Initial required-tool recall is not task success. The combined graph recovers most required capabilities without unsafe exposure, yet the frozen model does not execute successfully from the larger static set.

The M5 stop/go gate is negative. P6 fully covers two workflows and reaches .933 mean required recall, but static P6 and P8 disclosure obtain 0/15 model success. Even P9 obtains only 3/15, whereas P7 reactive JIT obtains 10/15. The experiment does not support training an adaptive

disclosure controller. The fixed graph remains useful for auditable candidate construction, but a useful capability set expands the action space; it does not solve action sequencing.

10 M6: Optional PRA-Native Resource Discovery

Native Q/K is not treated as another ordinary SDK function. In a generic deployment, the agent and resolver sit outside an opaque model API and cannot read hidden states, native query vectors, key vectors, or PRA cache entries. M6 therefore defines four explicit arrangements: co-located execution; a same-host shared-memory index that exports only $z_q = W_q q$; a model-server resolver endpoint that returns typed URIs and scores; and replicated query computation requiring the same model, tokenizer, layer, projection, and position treatment. Only the co-located mechanism is executed here. The SDK also implements the projected-query payload and identity-only server reply; raw Q/K is not persisted in artifacts.

At frozen Qwen layer 27, 1,625 tokens from 18 structured definitions are encoded with explicit name, description, parameter, return, side-effect, and tag boundaries. Five prompt presentations cover four single-tool and three multi-tool tasks. Baselines include token and indexed lexical routing, a signed-hash control, input-embedding means, native mean K, full token QK, and a lexical/native fusion. The rank-16 and rank-8/eight-centroid projections are the five Paper-2.8 HotpotQA/QASPER checkpoints applied zero-shot; they receive no tool supervision.

Table 4: M6 co-located discovery. Multi-step recall uses a budget equal to the workflow horizon. Bytes are declared persistent routing-index payloads, not Python/process overhead or backing native K/V.

Mode	Single Top-1	Multi recall	Successor recall	Index bytes
Token	.800	.889	1.000	13,427
Index	1.000	.862	1.000	13,427
Signed hash	.800	.670	.906	9,216
Input-embedding mean	.650	.619	.572	36,864
Lexical + rank-16	.800	.690	.689	65,427
Native mean K	0.000	.439	.611	36,864
Native token QK	.000	.439	.611	3,328,000
Paper-2.8 rank-16	.000	0.222	.167	52,000
Paper-2.8 rank-8/8-centroid	.000	.372	.167	2,304




Figure 9: M6 multi-step retrieval quality against routing-index payload. The compact transferred projection is inexpensive, but lexical channels dominate quality on this catalog. Full native-token scoring is both worse and larger.

M6 also stops negatively. Native query/key state does not improve this tool catalog, and the QA-trained low-rank map does not transfer to capability language. The lexical/native fusion is below either lexical channel alone. Consequently the supported generic architecture remains external explicit/token/index discovery. Model-server native routing is an optional research interface whose tighter coupling is not justified by these results.

11 M6.5: External Semantic-Hard Discovery

11.1 Frozen benchmark and split discipline

M6 answers a canonical-language question; it does not show that lexical search is sufficient for sloppy intent. We therefore freeze 306 single-tool requests over the same 18-tool catalog. Each identity has one held-out test and one validation realization in H0 (canonical), H1 (synonym), H2 (colloquial), H3 (implicit goal), H4 (bounded context plus shorthand), and H5 (Portuguese, Spanish, or French against English tool definitions). Eighteen additional audit rows cover the catalog. Validation and test contain 144 identities each; the complete query fingerprint is `e348334ad2362f284a217f1f852f175917a17540dd403f79953d7d2a7ebd7adb`. For H2 and later, tests reject exact tool names and aliases. H4 gives every channel the same context. No test wording selects an encoder, representation, fusion weight, or confidence threshold.

This is a controlled, project-authored fixture rather than a naturally sampled user corpus. Its multilingual mappings and tool tags share the benchmark’s typed concept inventory. Their accuracy measures the value of registration-time semantic metadata under those declared assumptions, not open-domain translation or multilingual generalization. The benchmark evaluates discovery only; the multi-step planning and execution boundaries remain those measured in M4–M5.

11.2 External policies

The concept map resolves domain-bounded operation/object phrases and retains language, source, and weight for every expansion. Tool registration emits four representations: description; name plus description; a deterministic card with purpose, operation, object, input/output types, and side effect; and three separate vectors for those fields. Three compact encoders are selected on validation only: all-MiniLM-L6-v2, BGE-small-en-v1.5, and multilingual MiniLM-L12-v2, all pinned by revision. Validation selects BGE with description cards and context-plus-query for English, and multilingual MiniLM with structured cards and query-only encoding for H5. Tool vectors are persisted; third-party weights are not copied into the repository.

The policy ladder compares token overlap (P0), BM25 (P1), typed dictionary expansion (P2), registration-time tags (P3), their lexical combinations, English and multilingual embeddings, a validation-frozen hybrid, and an identity oracle. A staged P10 resolver tries lexical, dictionary/tags, and the embedding hybrid before asking. Its four thresholds are fitted on validation.

Table 5: M6.5 results on 144 frozen test requests. Latency is the median measured end-to-end resolver time for this 18-candidate Python implementation; embedding registration is amortized.

Policy	Top-1	MRR	Recall@3	Median ms
Token	.257	.398	.403	8.6
BM25	.347	.485	.514	8.6
Typed dictionary	.729	.812	.868	8.6
Registration tags	.743	.823	.875	8.6
English embedding	.563	.682	.750	12.3
Multilingual embedding	.556	.690	.785	10.8
Dictionary/embedding hybrid	.729	.809	.875	12.3
Identity oracle	1.000	1.000	1.000	8.6

Tags improve Top-1 over BM25 by .396 (paired identity bootstrap 95% interval [.299, .500]; exact discordance $p < 10^{-10}$). The gain is not uniform. English BGE is strongest on H3 implicit goals (.722), tags on H4 context (.556), and curated tags are perfect on H5 because all three languages are represented in the frozen concept inventory. The hybrid improves H2 to .667 and H5 to .963 but does not dominate each specialized channel. This is why the deployment policy is a calibrated ladder, not a claim that one embedding replaces lexical metadata.




Figure 10: External discovery by hardness. Lexical quality declines as canonical wording disappears; typed metadata and embeddings recover different strata.




Figure 11: Required quality-versus-distance analysis. Bins contain 21, 37, 11, and 75 test identities from lower to higher lexical distance. Dictionary and hybrid channels retain substantial accuracy when token overlap is zero.

11.3 Selective action and deployment cost

On test, staged P10 selects 86 of 144 requests and asks on 58: coverage is .597, selective accuracy is .907, actionable accuracy is .542, and the false-action rate is .056. No selected top tool is in

the row’s unsafe set. Its Brier score is .135 and 10-bin ECE is .098. These held-out numbers are weaker than validation (.949 selective accuracy), which is a reason to retain ASK rather than lower thresholds after seeing test outcomes.

The 384-dimensional persistent tool vectors occupy 27.6 kB. The selected BGE encoder has 133.4 MB of parameters, loads into the process in 3.75 s, and encodes the full 306-query cohort in 1.13 s on the recorded CUDA host; its median per-query resolver path is 12.3 ms versus 8.6 ms for dictionary/tags. The multilingual encoder is larger (470.6 MB) and is not selected for the generic fusion.




Figure 12: H5 retrieval against English tool definitions. The curated dictionary and tag channels encode the three tested languages by construction; the multilingual encoder is an independent model-backed control.




Figure 13: Validation-calibrated P10 stage outcomes on test. More semantic-hard rows reach the hybrid or ASK stages instead of forcing a low-confidence call.

12 M7: Native Geometry on the Frozen Hard Set

M7 reuses all 144 test identities without changing wording, tool cards, or external policies. Frozen Qwen3-0.6B layer-27 pre-RoPE state supplies native mean K and full token QK. The Paper-2.8 rank-16 ensemble and rank-8/eight-centroid ensemble are transferred from five HotpotQA/QASPER checkpoints without tool supervision. Query encoding is charged online; raw native state is not persisted.

Table 6: M7 matched semantic-hard comparison. Model bytes include the frozen Qwen required to reproduce native queries; index bytes exclude backing model weights. The rank-8 timing includes the current per-query centroid construction and is not a production index latency.

Mode	Top-1	MRR	Recall@3	Median ms
BM25	.347	.485	.514	8.6
Tags	.743	.823	.875	8.6
English embedding	.563	.682	.750	12.3
External hybrid	.729	.809	.875	12.3
Native mean K	.056	.233	.313	116.7
Native token QK	.062	.220	.181	115.9
Paper-2.8 rank-16	.056	.194	.167	122.9
Paper-2.8 rank-8/8-centroid	.049	.190	.167	1,658.4

The native Top-1 range, .049–.063, contains the uniform catalog chance rate $1/18 = .056$. Tags exceed native mean K by .688 (paired bootstrap 95% interval [.604,.764]) and full token QK by .681 ([.597,.764]). Thus the harder wording does not reveal a frozen native advantage hidden by M6’s canonical queries. This result is narrower than “native Q/K cannot encode tool intent”: it says that one frozen layer, two direct reducers, and QA-trained low-rank projections do not expose

it zero-shot for this catalog.

Tool-specific native training does not open. External accuracy leaves oracle headroom, but none of the four frozen native diagnostics supplies the required positive signal, and reproducing native queries requires a 1.19 GB model plus co-location or a model-server extension. The next justified experiment would need tool-supervised evidence independent of this test set, not adaptation to the 144 held-out requests.




Figure 14: M7 quality–latency frontier on matched queries. External typed semantics dominates the tested frozen native modes in both quality and online latency on this small catalog.

13 M8: Automatic Tool Records and Record-Aware Materialization

13.1 Ordinary callables become typed resources

M6.5’s curated registration metadata raises an adoption question: must every function author learn a PRA-specific vocabulary? M8 adds `tool_record_from_callable`, which inspects but never invokes a Python callable. It preserves module and qualified name, sync/async status, signature, parameter kinds and defaults, input and return annotations, and descriptions parsed from Google, NumPy, or Sphinx-style docstrings. A provider-independent `ToolSchema` exports to provider formats only at the boundary. Its type cache covers primitives, generic containers, optional and union types, `Literal`, enumerations, dataclasses, `TypedDict`, annotated classes, and Pydantic when installed; no optional dependency is required.

Semantic enrichment remains auditable. Function-name terms, parameter and return types, docstring terms, canonical operation/object concepts, and dictionary expansions retain source and deterministic weight. Curated and automatic tags occupy separate fields. Generic verbs such as `get` and `read` can establish operation type without receiving the same lexical weight as `delete` or `update`. Name-inferred destructive operations are conservatively hidden from ordinary disclosure, but this hint does not grant or replace host execution authority.

E1 aligns 18 ordinary annotated callables with the existing rich catalog and uses the same 144 frozen test requests. Fusion weights are selected only on the 144 validation requests. The

automatic hybrid obtains 0.826 Top-1 and 0.951 Recall@3, compared with 0.806 Top-1 for the rich manual hybrid. Thus automatic metadata retains all measured gain over its lexical baseline in this controlled catalog. This is evidence for a low-annotation registration path, not evidence that docstrings are always complete or that inferred risk metadata is authorization.




Figure 15: E1 frozen-test quality. Ordinary signatures and docstrings recover the rich catalog result after validation-only fusion selection.

13.2 A zero-configuration ladder identifies the useful signals

The combined E1 result does not show which automatically derived fields matter. We therefore freeze one ToolRecord per callable and score every policy against the same 144 test requests. The record contains only function and qualified names, module, signature, parameter and return annotations, descriptions parsed from the docstring, and deterministic fields derived from them. Every derived term records whether it came from a name, docstring, parameter, return, type schema, dictionary expansion, inferred operation or object, automatic tag, or embedding field. No condition receives a per-tool alias, tag, operation, object, or capability edge.

Table 7 gives the complete ladder. Raw BM25 reaches 0.403 Top-1. Reweighting automatically extracted keywords alone raises Recall@3 but lowers Top-1 to 0.306; extraction is therefore not itself a semantic resolver. The global, tool-independent synonym dictionary supplies the largest isolated gain, reaching 0.750 Top-1. Adding inferred concepts and tags by elementwise score maxima does not improve that ranker in isolation. The selected name-plus-description embedding is complementary rather than sufficient. A validation-selected mixture of synonym, automatic-tag, and embedding scores reaches 0.812 Top-1 and 0.938 Recall@3, versus 0.806 and 0.910 for manually enriched registration. The one-row Top-1 difference is not evidence that automatic metadata is generally superior; it is a parity result on one controlled catalog.

Policy	Manual	Syn.	Concepts	Embed.	Top-1	R@3	R@5
A0 Raw BM25	no	no	no	no	0.403	0.556	0.639
A1 Auto keywords	no	no	no	no	0.306	0.597	0.681
A2 Keywords + synonyms	no	yes	no	no	0.750	0.882	0.910
A3 Inferred concepts	no	yes	yes	no	0.632	0.882	0.903
A4 Automatic tags	no	yes	yes	no	0.618	0.875	0.903
A5 Compact embedding	no	no	no	yes	0.556	0.764	0.840
A6 Automatic hybrid	no	yes	yes	yes	0.812	0.938	0.958
A7 raw union ($K = 10$)	no	yes	yes	yes	–	0.882	0.910
Manual-rich ceiling	yes	yes	yes	yes	0.806	0.910	0.958

Table 7: Frozen-test automatic discovery ladder. A7 has no unique Top-1: R@3 and R@5 evaluate caps three and five, while validation selects raw union at $K = 10$ for its maximum set recall. “Manual” denotes per-tool authored semantic metadata; all A0–A7 fields come from the same callable record.




Figure 16: A0–A6 contribution ladder and the manual-rich control. A5 is an independent embedding condition rather than an incremental addition to A4; A6 is the validation-frozen fusion.

The hardness decomposition in Figure 17 clarifies the complementarity. BM25 is already strong on direct H0 requests but degrades on implicit and contextual requests. The generic dictionary dominates the multilingual H5 fixtures, whereas the English compact embedding is strongest on H2–H4 paraphrastic and implicit English requests. Their fusion is therefore more stable than either channel alone. These multilingual rows test the frozen dictionary coverage, not open-domain translation.




Figure 17: Top-1 by frozen semantic-hardness stratum. The automatic hybrid tracks the manual ceiling while using different channels for lexical, implicit, and multilingual requests.

One-at-a-time source controls locate practical SDK requirements. Removing docstring terms lowers A1 Recall@3 from .597 to 0.521; type/schema terms add 0.042 Recall@3, although they do not improve A1 Top-1 and do not alter A4. With no docstring, embedding Top-1 falls from .556 to 0.375. Replacing descriptive names with opaque identifiers while retaining the docstring and signature yields 0.507. Thus registration remains useful with weak names, but documentation quality matters, especially to the embedding fallback. Figure 18 reports these controlled variants; they do not modify the frozen request text.




Figure 18: Sensitivity to callable documentation and naming. Each intervention changes only the named source while preserving signatures, target identities, and requests.

13.3 Candidate palettes do not reduce to one winner

The resolver can now return a bounded `CandidateSet` rather than only a Top-1 URI. Every candidate retains all channel hits and its admission source; deduplication uses stable URI. E2 sweeps budgets $K \in \{1, 2, 3, 4, 5, 6, 8, 10\}$ and compares the best single channel, normalized score fusion, raw union, and a diversity-aware union that admits one unseen candidate from each enabled chan-

nel before filling remaining slots. Metrics include required/useful recall, useful precision, unsafe exposure, schema tokens, and provenance diversity.

At $K = 4$, fused Top-K reaches 0.882 required-tool recall versus 0.847 for diversity union. Recall rises with larger palettes, but useful precision falls and labeled unsafe exposure increases. The requested default-union gate is therefore negative: candidate sets are first-class and union remains configurable, while validation-selected fusion remains the supported semantic palette policy for this cohort.




Figure 19: E2 required-tool recall against candidate count and serialized schema cost. Diversity preserves channel representation but does not dominate fused Top-K at matched budgets.

The isolated A7 rerun reaches the same decision with five strictly automatic channels: BM25, extracted keywords, keyword-plus-synonym scores, automatic-tag/concept scores, and the selected embedding. At $K = 4$, fused Top-K obtains 0.938 required-tool recall, versus 0.910 for raw union and 0.861 for diversity union. Validation selects raw union at $K = 10$ if recall is optimized alone; on test it reaches 0.993 recall but exposes at least one labeled unsafe candidate on 0.868 of requests and serializes 811 mean schema tokens. Channel complementarity is real but sparse: 8 of 144 targets are recovered only by the embedding, 3 only by an automatic lexical/concept channel, 129 by multiple channels, and 4 by none at per-channel $K = 3$. Larger set size, rather than broad unique recovery, explains most of the raw-union frontier.

The preregistered automatic-union JIT gate therefore remains closed. No new generative run is performed: on validation, neither union variant materially exceeds matched-budget fusion while holding unsafe exposure within one percentage point. This is an intentional stopped experiment, not missing data. Candidate provenance and union remain exposed by the SDK for diagnostics and future cohorts, but validation-frozen fusion remains the default.




Figure 20: A7 frozen-test frontier. Union does not dominate fused Top-K in required-tool recall and raises unsafe-candidate exposure rapidly as the palette grows. Discovery exposure is diagnostic; host authorization remains a separate execution boundary.

13.4 Semantic confusion changes the large-catalog result

M0 established that a typed index can locate opaque resources in catalogs up to 8,192 entries. That result did not test whether zero-configuration semantic channels preserve ranking quality when thousands of plausible callable descriptions compete for the same intent. The scale follow-up therefore keeps the 18 target callables and 144 frozen H0–H5 test requests, then constructs catalogs of 32, 128, 512, 2,048, and 8,192 ordinary Python callables under five catalog seeds. Each added callable shares exactly one canonical facet with a target—its operation or its object—but never both. Names, docstrings, signatures, annotations, and modules are inspectable by the same registration path as the targets. Every fifth eligible distractor uses a destructive operation on the anchor object so unsafe exposure remains measurable. No generated distractor receives a manual tag or target label at ranking time.

The evaluation reruns A2 lexical-plus-synonym scoring, A5 compact embeddings, A6 with the representation and weights frozen on the original 18-tool validation set, and A7 equal fusion, raw union, and diversity union. Candidate budgets $K \in \{2, 3, 4, 5, 6, 8, 10, 20\}$ are matched exactly. Table 8 reports seed means at $K = 10$; “Union P” is useful precision and “Unsafe” is the fraction of requests whose candidate palette contains at least one labeled destructive non-target. Exposure is diagnostic and never grants execution authority.

Catalog	A2 Top-1	A5 Top-1	A6 Top-1	A6 R@10	Union R	Union P	Unsafe	Tokens
32	0.729	0.529	0.787	0.967	0.944	0.153	0.414	916
128	0.629	0.394	0.675	0.914	0.879	0.120	0.399	1026
512	0.564	0.272	0.546	0.850	0.801	0.101	0.208	1050
2,048	0.472	0.229	0.421	0.692	0.715	0.085	0.125	1057
8,192	0.451	0.206	0.338	0.553	0.662	0.078	0.110	1059

Table 8: Five-seed semantically confusable callable-catalog scaling. A2–A6 columns report ranking quality. Union columns report raw A7 union at matched $K = 10$; context tokens serialize complete selected schemas.

The original frozen hybrid remains strongest at 32 and 128 tools. Its Top-1 advantage disappears

by 512 and reverses at 2,048: at 8,192, A2 lexical evidence reaches 0.451 Top-1 while A6 reaches 0.338. The compact embedding alone falls to .206. Semantic similarity accumulates plausible neighbors as the catalog grows, whereas exact lexical evidence retains a sharper identity signal. Figure 21 shows that this is a gradual ranking failure, not a threshold artifact.




Figure 21: Top-1 and Recall@10 under five seeded callable catalogs. The A6 mixture is frozen rather than retuned at each scale.

Channel diversity nevertheless becomes useful for bounded recovery. At 8,192 tools and $K = 10$, raw union reaches 0.662 required-tool recall, compared with 0.553 for A6. The gain is not free: useful precision is 0.078, the palette exposes a labeled unsafe candidate on 0.110 of requests, and full schemas consume 1059 tokens on average. The declining unsafe-exposure rate with catalog size does not imply safer catalogs; under fixed K , increasingly many non-destructive confusable candidates displace destructive ones from the visible palette. Figure 22 therefore reports recall and exposure together.




Figure 22: Matched- $K = 10$ required-tool recall and unsafe-candidate exposure. Raw and diversity union preserve more large-catalog recall, but neither is a safe or precise default disclosure policy.

Figure 23 exposes the 8K budget tradeoff rather than selecting $K = 10$ after the fact. Union’s recall advantage grows with palette size while useful precision falls; $K = 20$ is a diagnostic upper

bound, not a recommended disclosure budget.




Figure 23: The 8,192-tool matched-budget frontier. Candidate-set recall and useful precision must be read together.

The Python reference implementation scales approximately linearly. At 8,192 tools, exhaustive BM25 scoring takes 123.6 ms/query and tensorized A7 candidate construction takes 25.9 ms/query after channel scores are available. Discovery components occupy 50.7 MiB, while serializing the entire logical catalog would require 895909 tokens. These measurements establish the cost shape, not a production latency claim: an inverted lexical index and approximate semantic index should narrow candidates before exact reranking.




Figure 24: Reference implementation latency and measured discovery-component memory. Both axes are logarithmic; query encoding is batched and reported separately in the machine-readable cost rows.

The large-catalog result reopens A7 only as a retrieval research direction. It does not reopen the default-union or generative JIT gate because recall does not dominate useful precision, safety exposure, and schema context jointly. A production follow-up should use indexed candidate narrowing, calibrate abstention at each catalog scale, and test whether a second-stage reranker can retain the union’s recall before any schema becomes model-visible.

13.5 Selection authority and atomic tool records

M8 separates selection from materialization. The agent resolver owns catalog semantics, execution state, tenant and safety filters, and the candidate budget. It emits an authoritative `TOOL_CATALOG_SLICE` whose children are atomic `TOOL_DEFINITION` records. PRA preserves child identity, version, fingerprint, parent relation, field boundaries, and channel provenance, then materializes every admitted child in full. It does not rerank or prune chunks inside a selected schema.

If the slice exceeds the declared budget, the runtime must request a narrower slice, return an error, use an explicitly allowed temporary expansion, or drop whole records in candidate-priority order. Silent partial-schema fallback is invalid. The same substrate can construct tool responses, logs, terminal outputs, database results, and RAG chunks. M9 adds only predefined selection and full views for capability records; dynamic field reduction and detail expansion remain future work rather than hidden behavior in this paper.

E4 reruns the three frozen workflows across five prompt presentations with Top-1 JIT, diversity-union JIT at $K = 2, 4, 6, 8$, static required-set and graph sets, and all safe tools. The frozen Qwen3-0.6B Q4 model obtains 0.000 strict task success at Top-1 and 0.333, 0.333, 0.400, and 0.600 at the four union budgets. All-tools reaches only 0.133 despite complete candidate recall. The prompt presentations are robustness conditions, not training seeds; no model parameter changes. Figure 25 and Table 9 report execution success against disclosure cost rather than treating set recall as execution success.

Policy	Success	Required recall	Schema tokens	Disclosed tools
Top-1 JIT	0.000	0.857	191	2.3
Union JIT, $K = 2$	0.333	1.000	495	5.3
Union JIT, $K = 4$	0.333	1.000	964	6.9
Union JIT, $K = 6$	0.400	1.000	1371	9.2
Union JIT, $K = 8$	0.600	1.000	1922	11.8
Static required-set oracle	0.267	1.000	839	4.0
Static graph set	0.133	1.000	1337	8.0
All tools	0.133	1.000	2411	17.0

Table 9: E4 aggregate over 15 workflows per policy. Schema tokens are cumulative over attempted JIT steps; disclosed tools count unique definitions per workflow.




Figure 25: E4 strict workflow success against cumulative visible schema tokens. The frozen model, host authorization, and pure handlers are unchanged across candidate policies.

E5–E6 compare provider-structured full records, flat textual serialization, parameters without a description, descriptions without parameters, and a schema with one required field deliberately removed. Across 20 calls per condition, full atomic records obtain 1.000 accepted execution. Parameter-bearing structured records also reach 1.000 on these simple single-step calls, while description-only records reach 0.150. An explicit content-only internal-chunk router, which selects either the description or parameter-schema segment, likewise reaches only 0.150. Dropping one required field lowers acceptance to 0.000 and causes required-argument omission at rate 1.000. These controls support full selected- tool materialization as the default. M9 retains that atomic full view after selection; it does not introduce progressive detail expansion inside a schema.




Figure 26: E5–E6 exact calls and required-argument omissions. The generic router selects one internal schema segment; all partial-schema conditions are experimental controls, not runtime fallback policies.

14 M9: Typed Progressive Disclosure

Tool and skill records require different detail during choice and use. A model can often choose between capabilities from identity, signature or applicability, and a short description. It needs the complete parameter schema only when constructing a call, and the complete instruction body only when applying a skill. M9 represents this distinction as two deterministic views of one typed, versioned record:

$$R_{\text{select}} \longrightarrow R_{\text{full}}.$$

This transition is not token truncation or learned summarization. Field selection is defined by record type, and both views remain semantic record units. A full view is admitted atomically or not at all.

14.1 Records and explicit two-phase execution

A tool selection view contains stable name, call signature, short description, and side-effect class. Its full view contains the complete nested input schema, required and optional fields, enums and defaults, return schema, identity, version, and authorization-relevant metadata. A declarative skill selection view contains name, description, and when-to-use text. Its full view additionally contains ordered instructions, constraints, examples, dependencies, references, namespace, tenant, and version metadata. M9 supports text-only skills; it rejects script or executable payload metadata.

The runtime performs two ordinary model invocations. Phase A materializes the selection view of every externally bounded candidate and requests one structured capability identity. A validated `CapabilityTransition` permits only `selection`→`full`. Phase B removes non-selected candidates, materializes the selected full record, and asks the model to produce arguments or follow the selected instructions. Discovery and materialization do not grant execution authority; the existing host validator still checks selected identity, schema, side-effect policy, and request-scoped authorization.

The implementation is deliberately simple. Named views and policy validation live in `src/prahf/context_records.py`; the transition and token/KV accounting live in `src/prahf/progressive_disclosure.py`; and declarative skill records live in `src/prahf/skill_records.py`. There is no generation callback, decoder interruption, confidence-triggered expansion, arbitrary compression, or mid-token materialization. Those dynamic mechanisms belong to a later large-record study.

For candidate full-view tokens F_K , Phase-A selection tokens A_K , and the selected full-view tokens B , we report

$$\text{DisclosureRatio} = \frac{A_K}{F_K}, \quad \text{TotalDisclosureRatio} = \frac{A_K + B}{F_K}.$$

The latter counts both model phases and therefore does not hide the second invocation.

14.2 Protocol

The tool experiment freezes 8 semantic-hard requests over the 18 M8 callables. The skill benchmark contains 25 realistic, semantically confusable declarative skills and 50 name-blind requests split evenly between validation and test. Validation selects the same external BM25, automatic semantic, compact-embedding, and fused discovery stack used for tools. Generation uses 8 held-out skill requests. Both experiments compare full-all (T0/S0), selection-only (T1/S1), selection-to-full (T2/S2), and oracle-full (T3/S3) at $K \in \{2, 4, 6, 8\}$, followed by all-18-tool and all-25-skill stress controls.

All generation rows use the exact frozen Qwen3-0.6B revision in FP16 on an NVIDIA GeForce GTX 950M, deterministic decoding, and no parameter updates. The small eight-request generation cohorts are mechanism checks, not population estimates. Tool evaluation separates required Recall@ K , capability choice, schema validity, argument semantics, host acceptance, and strict success. Skill evaluation uses deterministic identity markers, required output headings, and forbidden external-action claims. Timing includes CUDA synchronization, serialization, both TTFTs, and both generations.

14.3 Discovery and view size

The shared discovery stack transfers without a second skill-specific retrieval system. On the 25-query skill test split, BM25 obtains 0.720 Top-1, compact embeddings obtain 0.800, and validation-selected fusion obtains 0.880. The corresponding tool bars in Figure 27 reuse M8’s 144 test requests and should not be read as a paired cross-type comparison.




Figure 27: The same external discovery families applied to tools and declarative skills. Cohort sizes differ; the figure tests substrate reuse, not which resource type is intrinsically easier.

Full tool schemas contain a median 233.5 tokens (p95 244), and their selection views retain a median 0.368 of that count. Full skill instructions contain a median 468 tokens (p95 480), while skill selection views retain 0.215. These bodies are intentionally nontrivial: full-all instruction cost therefore grows linearly with the candidate palette.




Figure 28: Median full-view tokens and median selection/full token ratio for the 18 tool schemas and 25 declarative skills. Ratios are percentages in the red bars; no view is formed by first- N -token truncation.

14.4 Tools: choice improves at moderate budgets, not at catalog scale

Table 10 reports the full tool matrix. At $K = 4$, T2 raises choice accuracy from 0.500 to 0.875 and strict success from 0.250 to 0.375, while exposing 0.589 of T0’s capability tokens. At $K = 8$, T0 and T2 both obtain 0.625 choice accuracy; T2 uses 0.479 of full-all tokens but adds 0.492 seconds mean wall time. At all 18 tools, T2 choice falls to 0.375 versus 0.625 for T0. Compact views do not remove the need for bounded discovery.

Selection-only obtains no strict successes: compact signatures are sufficient for choosing in some rows, but are not a reliable replacement for full provider schemas during argument construction. T3 obtains 1.000 schema validity and host acceptance at every budget, yet only 0.375 strict success because five of eight requests elicit non-exact argument values even when the correct tool is supplied alone. This oracle ceiling is why M9 does not interpret low strict success as a routing-only failure.

14.5 Skills: context economy without loss at $K = 8$

Skill full views consistently recover the deterministic instruction contract. At $K = 8$, S2 and S0 both obtain 0.625 strict success. S2 exposes 0.341 of S0’s capability tokens and reduces mean wall time from 63.870 to 15.900 seconds despite two model invocations. This is a measured property of quadratic SDPA on the stated 4-GB GPU, not a general latency claim. S1 fails the required identity and output contract, confirming that compact applicability text cannot replace the full instructions.

K	n	Condition	Recall	Choice	Valid/follow	Success	Tokens	Wall (s)
2	8	T0 full-all	0.88	0.75	0.88	0.50	472.00	4.36
2	8	T1 selection-only	0.88	0.00	0.00	0.00	173.88	2.93
2	8	T2 select→full	0.88	0.75	0.88	0.25	381.88	5.00
2	8	T3 oracle-full	0.88	1.00	1.00	0.38	236.88	4.04
4	8	T0 full-all	1.00	0.50	0.88	0.25	945.50	4.63
4	8	T1 selection-only	1.00	0.00	0.00	0.00	349.12	2.88
4	8	T2 select→full	1.00	0.88	0.88	0.38	557.12	5.43
4	8	T3 oracle-full	1.00	1.00	1.00	0.38	236.88	4.06
6	8	T0 full-all	1.00	0.62	1.00	0.25	1417.00	4.96
6	8	T1 selection-only	1.00	0.00	0.00	0.00	523.62	3.31
6	8	T2 select→full	1.00	0.75	1.00	0.25	759.88	6.42
6	8	T3 oracle-full	1.00	1.00	1.00	0.38	236.88	4.08
8	8	T0 full-all	1.00	0.62	0.88	0.38	1889.25	5.75
8	8	T1 selection-only	1.00	0.00	0.00	0.00	697.88	3.44
8	8	T2 select→full	1.00	0.62	0.88	0.12	904.50	6.25
8	8	T3 oracle-full	1.00	1.00	1.00	0.38	236.88	3.92
18	8	T0 full-all	1.00	0.62	0.88	0.38	4235.00	8.49
18	8	T1 selection-only	1.00	0.00	0.00	0.00	1567.00	8.55
18	8	T2 select→full	1.00	0.38	0.62	0.00	1715.12	8.85
18	8	T3 oracle-full	1.00	1.00	1.00	0.38	236.88	4.06

Table 10: Tool progressive-disclosure matrix. n is the number of context-fitting generation cases. “Valid/follow” is schema-valid call rate. Strict success also requires the exact benchmark arguments and accepted host execution.

K	n	Condition	Recall	Choice	Valid/follow	Success	Tokens	Wall (s)
2	8	S0 full-all	0.88	0.75	0.88	0.62	921.38	13.11
2	8	S1 selection-only	0.88	0.00	0.00	0.00	199.88	6.78
2	8	S2 select→full	0.88	0.62	0.88	0.62	602.62	11.33
2	8	S3 oracle-full	0.88	1.00	1.00	1.00	460.62	11.27
4	8	S0 full-all	0.88	0.75	1.00	0.75	1852.62	24.79
4	8	S1 selection-only	0.88	0.00	0.00	0.00	402.38	9.36
4	8	S2 select→full	0.88	0.62	0.88	0.62	805.00	12.74
4	8	S3 oracle-full	0.88	1.00	1.00	1.00	460.62	11.06
6	8	S0 full-all	0.88	0.75	0.88	0.62	2794.00	34.74
6	8	S1 selection-only	0.88	0.00	0.00	0.00	606.62	9.60
6	8	S2 select→full	0.88	0.50	0.75	0.50	1067.12	15.91
6	8	S3 oracle-full	0.88	1.00	1.00	1.00	460.62	11.01
8	8	S0 full-all	1.00	0.62	1.00	0.62	3728.75	63.87
8	8	S1 selection-only	1.00	0.00	0.00	0.00	808.25	10.40
8	8	S2 select→full	1.00	0.62	0.88	0.62	1269.88	15.90
8	8	S3 oracle-full	1.00	1.00	1.00	1.00	460.62	11.15
25	0	S0 full-all	1.00	—	—	—	—	—
25	8	S1 selection-only	1.00	0.00	0.00	0.00	2527.00	33.70
25	8	S2 select→full	1.00	0.50	0.75	0.50	2928.50	23.98
25	8	S3 oracle-full	1.00	1.00	1.00	1.00	460.62	11.16

Table 11: Declarative-skill progressive-disclosure matrix. “Valid/follow” is deterministic instruction-contract success. The $K = 25$ S0 row has $n = 0$ because all eight 11,765–11,771-token prompts exceeded the declared 8,192-token runtime limit before generation.

At $K = 25$, the 11,702 capability tokens of S0 cannot run under the measured GPU limit, whereas S2 fits at 2,929 mean capability tokens and obtains 0.500 strict success. This stress result establishes feasibility under one hardware envelope; it does not imply that larger-memory runtimes cannot execute full-all.




Figure 29: Strict task success against total capability tokens exposed at $K = 8$. Totals include both phases for T2/S2. The oracle points isolate the call or instruction-following ceiling after externally correct selection.

14.6 Gate decision

The architectural gate opens: typed selection/full views, explicit transitions, atomic selected records, and shared tool/skill discovery work without callback machinery. The skill inclusion gate also opens because fused discovery exceeds BM25 by 0.160 Top-1 and selected-full execution cuts instruction context while remaining deterministically evaluable. The deployment gate is conditional. Selection-to-full is the SDK policy for an already bounded capability palette, but the resolver must retain its calibrated candidate limit and host safety checks. This small cohort does not establish statistical non-inferiority, and the all-tool result rejects using compact views as a substitute for discovery.

15 M10: Large-Catalog Palette Use

The 8,192-tool retrieval result leaves a downstream question open. A larger candidate set is useful only when the model can identify the target among its additional confusable tools. M10 freezes the same callable generator and eight M9 tool requests, then separates three probabilities:

$$P(t \in C), \quad P(\hat{t} = t \mid t \in C), \quad P(\hat{t} = t).$$

The first is resolver recall, the second is model discrimination conditional on a successful resolver, and the third is their end-to-end product. Every raw row retains the chosen identity, candidate count, channel sources and ranks, unsafe and useful candidates, view tokens, and wall time.

15.1 Protocol and model boundary

The full retrieval grid covers $N \in \{32, 128, 512, 2048, 8192\}$, $K \in \{2, 4, 6, 8, 10\}$, and BM25, frozen fusion, raw union, and diversity union. Seed 11 covers the Cartesian grid. The decisive $N = 8192, K = 10$ comparison uses all five catalog seeds and 40 query-seed pairs; a fixed agreement bonus is an additional control. Candidate identities are bound to one-digit labels and the official Ollama `qwen3:0.6b` Q4 artifact generates one label at temperature zero. This low-memory artifact is kept separate from the exact FP16 CUDA model used below for call construction. Catalog seeds vary generated distractors, not model weights.

The 40-row recall values in Table 12 differ from the 144-query values in Table 8 because M10 uses the smaller argument-annotated generation cohort. They answer different questions: M8 estimates resolver behavior on all frozen test requests; M10 tests whether a model can use selected palettes on requests with executable expected calls.




Figure 30: Five-seed large-catalog Recall@10. Union increasingly helps candidate generation as the logical catalog grows, before any candidate is shown to the generation model.

15.2 Candidate recall is not candidate use

Policy	n	Recall	Cond. choice	End-to-end	Unsafe choice	Tokens
BM25	40	0.375	0.600	0.225	0.000	1465
Fused	40	0.300	0.083	0.025	0.075	1514
Raw union	40	0.375	0.000	0.000	0.000	1464
Diversity union	40	0.375	0.667	0.250	0.000	1465
Agreement union	40	0.375	0.467	0.175	0.050	1464

Table 12: Generated capability choice at $N = 8192$, $K = 10$ across eight requests and five catalog seeds. Conditional choice excludes target-absent rows; end-to-end choice counts them as failures. Tokens expose compact selection views only.

Raw union raises recall from 0.300 for frozen fusion to 0.375, but Qwen chooses none of its 15 available targets. The model cannot use raw union’s additional hypotheses. Diversity union retains the same 0.375 recall and chooses 10/15 available targets, yielding 0.250 end-to-end choice. Fusion obtains 1/12 and 0.025 respectively. On paired query–seed identities, diversity wins ten rows that fusion loses and loses one row that fusion wins (exact two-sided $p = 0.0117$). The fixed agreement bonus obtains 7/15 conditional successes; its paired effect versus fusion is unresolved ($p = 0.0703$). This is evidence for candidate ordering and channel diversity, not for broad schema disclosure.




Figure 31: Conditional model choice across candidate budgets. A resolver gain becomes operational only when the model can distinguish target from admitted distractors.

The compact/full view control uses 128 matched decisions over four stratified requests. Compact views average 0.453 of full-view tokens and obtain 10/128 correct choices, versus 7/128 for full views. Both rates are low and this study does not claim non-inferiority. It shows that exposing every nested schema field does not repair large-catalog selection for this small model. Three executable large-schema controls contain nested dataclasses, lists, optional objects, enums, literals, and multiple required fields; their full views contain a median 721 tokens (maximum 736), while selection views retain 0.219–0.279.

15.3 Selection to atomic execution

The causal execution matrix returns to the exact frozen Qwen3-0.6B FP16 revision. At $N = 8192$ and $K \in \{4, 8\}$, all five catalog seeds and four requests compare full-all, compact-only, compact-to-full, and oracle-to-full. The model first selects from complete compact records. The runtime then removes losing records, materializes one full schema atomically, asks for arguments, and submits the result to the same host validator. No resolver score grants execution authority.

Policy	View	Recall	Choice	Accepted	Strict	Tokens
Fused	Full-all	0.300	0.100	0.750	0.100	2615
Fused	Select→full	0.300	0.100	0.650	0.100	1495
Fused	Oracle full	0.300	1.000	1.000	0.750	278
Raw union	Full-all	0.250	0.150	0.700	0.150	2561
Raw union	Select→full	0.250	0.050	0.600	0.050	1441
Raw union	Oracle full	0.250	1.000	1.000	0.750	278
Diversity union	Full-all	0.250	0.000	0.500	0.000	2558
Diversity union	Select→full	0.250	0.250	0.550	0.250	1419
Diversity union	Oracle full	0.250	1.000	1.000	0.750	278

Table 13: Exact-FP16 large-catalog tool execution at $N = 8192, K = 8$ ($n = 20$ per row). Strict success requires target choice, exact benchmark arguments, and accepted host execution. Oracle full isolates post-selection call quality.

Diversity compact-to-full recovers every target-present row (5/5), obtains 5/20 strict accepted calls, and exposes 1,419 capability tokens on average, 0.556 of its 2,558-token full-all palette. It materializes one full schema and avoids 2,294 full-schema tokens. Measured native-K/V accounting falls from 293.4 MB-equivalent bytes for full-all to 162.7 MB-equivalent bytes. Raw union obtains 1/20 strict calls and fused obtains 2/20. Diversity versus fusion has the same positive direction in 4 versus 1 discordant rows, but the 20-pair effect is unresolved (exact $p = 0.375$). Oracle-to-full obtains 15/20 strict calls, locating the remaining ceiling in discovery and model choice rather than schema execution alone.




Figure 32: Large-catalog strict accepted calls against total capability tokens at $N = 8192, K = 8$. Two-phase totals include compact selection views and the one selected full schema.

15.4 Large-catalog JIT remains recall-bound

Two two-step workflows pair account retrieval with update and document search with export. For each step the resolver searches the full 8,192-tool logical catalog, exposes $K \in \{4, 8\}$ compact records, materializes one selected full schema, validates the call, records a typed observation, and advances. The five catalog seeds yield ten workflows per policy.

Policy	Steps	Recall	Cond. choice	Accepted	Workflow	Tokens
Fused	20	0.300	0.333	0.650	0.000	1495
Raw union	20	0.250	0.200	0.600	0.000	1441
Diversity union	20	0.250	1.000	0.550	0.000	1419

Table 14: Large-catalog two-step JIT at $K = 8$. Conditional choice can be high while workflow success remains zero when a required step is absent from the palette.

Diversity union again uses its successful palettes well: it chooses 5/5 target-present steps. But per-step recall is only 0.250, so no policy completes either required step in all ten workflows. Raw union and fusion also obtain zero workflow success. No generated proposal is an unsafe tool, but host acceptance remains 0.55–0.65 because wrong identities and malformed or incomplete arguments are still rejected. The default-union gate therefore remains closed. The supported policy is frozen fusion by default, with diversity union exposed as an experimental high-recall mode and explicit host authorization after every selection.

16 Next Experimental Gates

The next gates return to the broader persistent-context thesis: hierarchical skills, mutation and revocation, bounded `#_head` history, superseded facts, and permission-scoped child sessions

$$M_{\text{child}} = R(M_{\text{parent}}) \cup \Delta M_{\text{child}}.$$

Cross-model replication is also required because M2–M10 use one Qwen model family for pretrained mechanisms and downstream capability choice. The generated large-catalog study does not replace natural catalogs or a tool-identity-disjoint benchmark; both should test whether the observed lexical advantage and union recovery survive distribution shift. The OpenAI-compatible endpoint belongs after these mechanism checks, so harness behavior does not hide resolver, disclosure, or execution failures.

17 Reproducibility

M8’s E1–E3 driver is `run_auto_union_records.py`. Candidate sets for E4 are frozen by `run_union_jit_records.py` with `--prepare-only`. The five-presentation E4–E6 driver is `run_union_jit_ollama.py`. It uses the official `qwen3:0.6b Q4_K_M` artifact on CPU with temperature zero, no GPU offload, per-call unloading, and deterministic host-side schema validation. This low-memory runtime is separate from the exact-FP16 M2–M7 rows. The runner checkpoints at each completed condition. The companion summarizer regenerates tables, macros, findings, and all four M8 figures.

The zero-configuration source, type, quality, and A0–A7 channel ablations use `run_auto_discovery_ablation.py`; `summarize_auto_discovery_ablation.py` regenerates its table,

macros, four figures, and machine-readable summary. The driver constructs one frozen record for each of the same 18 callables, selects embedding representation and fusion weights only on validation, and evaluates all reported policies on 144 test requests. It used the compact BGE encoder on an NVIDIA GeForce GTX 950M and completed in 103 seconds. The manifest records every automatic field’s source. The A7 JIT status is a checked-in gate-closed artifact because no union condition met the validation recall-and-safety criterion.

The semantically confusable scale follow-up uses `run_scaled_auto_discovery.py`; it checkpoints one gzip shard per catalog size and seed and resumes only when the protocol hash matches. The run covers five sizes, seeds {11, 23, 37, 53, 71}, 144 frozen test requests, three ranking policies, six matched candidate strategies, and eight budgets on CUDA. `summarize_scaled_auto_discovery.py` audits row completeness and budget compliance, aggregates catalog-seed means and standard deviations, and regenerates the table, macros, four figures, and machine-readable findings. The raw rows distinguish index construction, query encoding, channel scoring, candidate construction, component memory, logical schema tokens, and selected schema tokens.

M9 is prepared by `prepare_progressive_disclosure.py`, executed by `run_progressive_disclosure.py`, and summarized by `summarize_progressive_disclosure.py`. The preparation stage freezes selection/full views, the 25-skill catalog, 50 semantic-hard skill queries, validation-selected fusion weights, candidate sets, and disclosure costs. The generation runner checkpoints each condition and records prompt/context fit, native-K/V bytes, both TTFTs, transition overhead, wall time, raw output, and deterministic evaluator fields. The summarizer regenerates two full condition tables, three figures, macros, and machine-readable findings. No model or encoder weights are committed.

M10 is prepared by `prepare_missing_experiments.py`, executed by `run_missing_experiments.py`, and summarized by `summarize_missing_experiments.py`. Preparation regenerates the same five callable catalogs, freezes 5,000 policy/budget palettes, serializes 3,122 unique candidate views, and records three realistic nested-schema stress tools. The choice phase uses the official Ollama `qwen3:0.6b` Q4 artifact with persistent CPU loading and checkpoints every bounded block. The execution phase resumes the same checkpoint with the exact FP16 local revision on CUDA. Raw artifacts contain 968 compact choices, 160 matched full-view choices, 704 execution rows, 120 JIT step rows, per-seed identities, paired exact effects, timing, token/KV accounting, and generated figures and tables. The split precision protocol is explicit in the manifests; Q4 and FP16 rows are never pooled as repeated seeds of one estimator.

The M0 benchmark entry points are `experiments/paper6.5_tools/run_m0_policy_study.py` and `summarize_m0_policy_study.py`. Machine-readable outputs include the manifest, one row per policy/query trace, index costs, seed summaries, stratum summaries, oracle distributions, and plots under `docs/papers/shared/results/paper6.5_tools`. The recorded environment is Python 3.10.11 on Windows, Intel64 Family 6 Model 94. All reported M0 results are deterministic conditional on the five catalog seeds.

M1 is reproduced by `run_m1_toy_model.py` followed by `summarize_m1_toy_model.py`. Checked-in artifacts include all 720 evaluation rows, per-step training histories, seed and aggregate summaries, paired effects, a manifest, and generated figures. The run used PyTorch on an NVIDIA GeForce GTX 950M and seeds {11, 23, 37, 53, 71}. Exact configuration and runtime metadata are recorded in `paper6.5_tools/m1/m1_manifest.json`.

M2–M4 are reproduced by `run_m2_m4_pretrained.py`. M5 uses `run_m5_disclosure.py`; its direct, reactive, and oracle model rows reuse the frozen M4 summaries, while P6/P8 are newly executed static disclosures. M6 uses `run_m6_native_discovery.py`. The final tables, macros, and figures are generated by `summarize_m2_m6.py`. Raw rows contain one call, workflow step, policy/task, or routing-query record as appropriate. Manifests record the frozen model revision,

CUDA device, presentation seeds, index accounting, checkpoint provenance, and the fact that all execution handlers are pure in-memory fixtures. M6.5 uses `run_m6_5_external_semantics.py`; M7 uses `run_m7_semantic_native.py`; and `summarize_semantic_hard.py` regenerates paired effects, tables, macros, and figures. Their manifests pin query fingerprints, split discipline, encoder/model revisions, representation selection, thresholds, and CUDA runtime. Tool vectors are committed, but third-party encoder weights and raw native Q/K are not.

18 Conclusion

A large agent catalog should not be mistaken for the context needed on every turn. M0 resolves typed identities without exposing the catalog, and M1 shows that one selected native-K/V definition is used causally. M2–M4 carry the mechanism into a frozen pretrained model: selected schemas support exact calls, host authorization remains separate from generation, typed observations retain provenance, and reactive JIT supports bounded multi-step execution better than static disclosure in this cohort.

The extensions also locate two boundaries. M5 capability graphs improve set recall but not task success; more useful tools visible at once are still more actions for a small model to sequence. M6’s canonical queries favor lexical indexing, whereas M6.5 shows that authored tags, dictionaries, and compact embeddings recover semantic-hard intent that BM25 misses. M7 then finds no zero-shot native benefit on those same hard rows. These are positive external- resolver results and bounded native stop results, not evidence that all tool semantics are lexical or that trained native geometry cannot work.

The supported architecture combines explicit identity, indexed lexical search, provenance-bearing dictionaries, compact embedding fallback, calibrated ASK, reactive disclosure, host-authorized execution, and typed observations. It works with an opaque generation API; graph neighborhoods and native Q/K remain optional research interfaces. M8 adds ordinary callables as typed records and atomic PRA transport of the selected palette. On this catalog, generic synonyms supply most of the isolated gain, and fusion with automatic tags and embeddings matches the rich manual result. Union remains optional because it does not dominate fusion on the original cohort. Under five semantically confusable catalog seeds, frozen hybrid ranking degrades faster than lexical evidence; bounded union restores part of the 8K recall but leaves low useful precision and nonzero unsafe exposure. The result motivates indexed high-recall retrieval followed by calibrated reranking, not default broad disclosure.

M9 narrows the cost after discovery. Deterministic selection views preserve a bounded choice palette; the runtime then removes losing candidates and exposes one complete authoritative tool schema or declarative skill. The skill result shows equal $K = 8$ strict success at roughly one third of full-all capability tokens, while the all-skill full prompt does not fit the measured runtime envelope. The tool result is less uniform and degrades at all-catalog scale, so progressive disclosure complements rather than replaces calibrated discovery. Basic declarative skills are now covered; hierarchical skills, persistent history, mutation, revocation, and session inheritance remain untested.

M10 closes the downstream large-catalog question without erasing its boundary. Raw union’s recall gain is unusable by the Q4 chooser, whereas diversity union turns 10/15 target-present palettes into correct choices and significantly improves paired end-to-end choice over frozen fusion. In exact-FP16 execution, diversity selection-to-full converts all 5/5 target-present rows into strict accepted calls while using 0.556 of full-all capability tokens. The two-step JIT result nevertheless remains zero because only one quarter of required steps contain the target. Typed progressive disclosure therefore solves the cost of using a bounded palette; it does not solve large-catalog

discovery. Fusion remains the SDK default, diversity union is a measured experimental high-recall mode, and callbacks or dynamic detail generation remain outside this paper.